

■ Python Project Code Export

- [PY] class_count.py
- [PY] clean_pngs.py
- [PY] config.py
- [PY] convert_polygons.py
- [PY] dataset.py
- [PY] demo.py
- [PY] evaluate.py
- [PY] evaluation\metrics.py
- [PY] infer.py
- [PY] inference\engine.py
- [PY] models\dvx.py
- [PY] models\encoder.py
- [PY] models\extrusion.py
- [PY] models\heads.py
- [PY] models\model.py
- [PY] param.py
- [PY] setup.py
- [PY] train.py
- [PY] training\losses.py
- [PY] training\trainer.py
- [PY] utils\visualization.py
- [PY] vortex.py

■ File: class_count.py

```
=====
1: import cv2, numpy as np, glob
2: from collections import Counter
3:
4: all_classes = set()
5: class_counts = Counter()
6:
7: # Go through all mask images
8: for mask_file in glob.glob("./data/floorplans/train/*/mask.png"):
9:     mask = cv2.imread(mask_file, cv2.IMREAD_GRAYSCALE)
10:    unique, counts = np.unique(mask, return_counts=True)
11:    for u, c in zip(unique, counts):
12:        class_counts[u] += c
13:    all_classes.update(unique)
14:
15: # Total pixels
16: total_pixels = sum(class_counts.values())
17:
18: print("Classes found in dataset:", sorted(all_classes))
19: print("\nPixel distribution per class:")
20: for cls in sorted(class_counts.keys()):
21:     percentage = (class_counts[cls] / total_pixels) * 100
22:     print(f"Class {cls}: {class_counts[cls]} pixels ({percentage:.2f}%")
-----
```

■ File: clean_pngs.py

```
=====
1: from PIL import Image
2: import os
3: import shutil
4:
5: data_dir = r"data/floorplans"
6:
7: def safe_clean_image(path):
8:     """Safely clean image by only removing ICC profiles, preserving all pixel data"""
9:     try:
10:         # Create backup first (optional safety measure)
11:         backup_path = path + ".backup"
12:
13:         with Image.open(path) as img:
14:             # Check if image is already clean
15:             if 'icc_profile' not in img.info:
16:                 print(f"Already clean: {path}")
17:                 return True
18:
19:             # Create backup
20:             shutil.copy2(path, backup_path)
21:
22:             # Method 1: Just strip ICC profile while preserving everything else
23:             img_data = img.copy()
24:
25:             # Remove only the problematic ICC profile
26:             if 'icc_profile' in img_data.info:
27:                 del img_data.info['icc_profile']
28:
29:             # Save with same format and quality, just without ICC profile
30:             img_data.save(path, format="PNG", optimize=False) # No optimization to preserve exact pixels
31:
32:             # Remove backup if successful
33:             os.remove(backup_path)
34:
35:         print(f"Cleaned ICC profile from: {path}")
36:         return True
37:
38:     except Exception as e:
39:         # Restore backup if it exists
40:         backup_path = path + ".backup"
41:         if os.path.exists(backup_path):
42:             shutil.move(backup_path, path)
43:             print(f"Restored backup for: {path}")
44:
-----
```

```

45:         print(f"Failed {path}: {e}")
46:         return False
47:
48: def verify_image_integrity(path):
49:     """Verify image can still be loaded properly after cleaning"""
50:     try:
51:         with Image.open(path) as img:
52:             # Try to access pixel data to ensure image is valid
53:             _ = img.size
54:             _ = img.mode
55:             # Try to load a small sample of pixel data
56:             _ = img.getpixel((0, 0))
57:         return True
58:     except Exception as e:
59:         print(f"WARNING: Image integrity check failed for {path}: {e}")
60:         return False
61:
62: # Process only image.png files
63: processed_files = []
64: cleaned_count = 0
65: failed_count = 0
66: already_clean = 0
67:
68: print("Starting safe ICC profile removal for dataset...")
69: print("This preserves all pixel data and only removes problematic metadata.")
70:
71: for root, _, files in os.walk(data_dir):
72:     for f in files:
73:         if f == "image.png": # Only process image.png files
74:             path = os.path.join(root, f)
75:             processed_files.append(path)
76:
77:             # Check if already clean
78:             try:
79:                 with Image.open(path) as img:
80:                     if 'icc_profile' not in img.info:
81:                         already_clean += 1
82:                         continue
83:             except:
84:                 pass
85:
86:             if safe_clean_image(path):
87:                 # Verify integrity after cleaning
88:                 if verify_image_integrity(path):
89:                     cleaned_count += 1
90:                 else:
91:                     failed_count += 1
92:             else:
93:                 failed_count += 1
94:
95: print(f"\n" + "="*50)
96: print(f"DATASET CLEANING SUMMARY")
97: print(f"="*50)
98: print(f"Total image.png files found: {len(processed_files)}")
99: print(f"Already clean (no ICC profile): {already_clean}")
100: print(f"Successfully cleaned: {cleaned_count}")
101: print(f"Failed to clean: {failed_count}")
102: print(f"Total files processed: {already_clean + cleaned_count + failed_count}")
103:
104: if failed_count > 0:
105:     print(f"\nWARNING: {failed_count} files couldn't be cleaned.")
106:     print(f"Check these files manually - they may be corrupted.")
107:
108: print(f"\nDataset should now be ready for training without libpng warnings!")
109:
110: # Optional: Test load a few random images to verify dataset integrity
111: print(f"\nTesting random samples for integrity...")
112: import random
113: test_files = random.sample(processed_files, min(5, len(processed_files)))
114: for test_path in test_files:
115:     if verify_image_integrity(test_path):
116:         print(f"? {test_path}")
117:     else:

```

```
118:         print(f"? {test_path} - POTENTIAL ISSUE")
```

■ File: config.py

```
=====
1: """
2: Configuration settings for the Neural-Geometric 3D Model Generator
3: Enhanced with dynamic curriculum and adaptive training strategies
4: """
5: from dataclasses import dataclass
6: from typing import Tuple, Dict, Any, Optional, List
7: import torch
8:
9:
10: @dataclass
11: class DataConfig:
12:     """Data-related configuration"""
13:     data_dir: str = "./data/floorplans"
14:     image_size: Tuple[int, int] = (256, 256)    # keep full resolution for accuracy
15:     voxel_size: int = 64
16:     batch_size: int = 4                        # balance speed & memory
17:     num_workers: int = 8                      # faster dataloader (tune per CPU)
18:     augment: bool = True
19:
20:
21: @dataclass
22: class ModelConfig:
23:     """Model architecture configuration optimized for high accuracy"""
24:     input_channels: int = 3
25:     num_classes: int = 5
26:     feature_dim: int = 512    # reduced from 768 ? faster while keeping strong accuracy
27:     num_attributes: int = 6
28:     voxel_size: int = 64
29:     max_polygons: int = 20    # enough for complex layouts
30:     max_points: int = 50      # good detail without huge cost
31:     dropout: float = 0.05
32:     use_attention: bool = True
33:     use_deep_supervision: bool = True
34:
35:     # Auxiliary heads for novel training strategies
36:     use_latent_consistency: bool = True
37:     use_graph_constraints: bool = True
38:     latent_embedding_dim: int = 256
39:
40:
41: @dataclass
42: class CurriculumConfig:
43:     """Dynamic curriculum learning configuration"""
44:     # Adaptive stage transitioning
45:     use_dynamic_curriculum: bool = True
46:     stage_switch_patience: int = 5
47:     min_improvement_threshold: float = 0.001
48:     plateau_detection_window: int = 3
49:
50:     # GradNorm / gradient tracking
51:     gradient_norm_window: int = 100
52:
53:     # Objectives for multi-objective optimization
54:     objectives: Optional[List[str]] = None
55:
56:     # Topology-aware scheduling
57:     topology_schedule: str = "progressive"    # "progressive", "linear_ramp", "exponential"
58:     topology_start_weight: float = 0.1
59:     topology_end_weight: float = 1.0
60:     topology_ramp_epochs: int = 20
61:
62:     # config.py (snippet ? add into the existing config class/dict)
63:     # Mixed precision and training conveniences
64:     use_mixed_precision = True                # enable AMP
65:     cache_in_memory = False                  # set True if host RAM can hold dataset
66:     accumulation_steps = 1                   # effective batch size multiplier
67:     dvx_step_freq = 1                        # run DVX refinement every N steps (1 = every step)
```

```

68:     persistent_workers = True                # DataLoader persistent workers
69:     prefetch_factor = 4                      # DataLoader prefetch
70:     num_workers = 8                          # default num workers for DataLoader (tune by CPU)
71:     # Progressive resolution settings (example)
72:     voxel_size_stage = { "stage1": 32, "stage2": 32, "stage3": 64 } # voxel sizes per stage
73:     image_size_stage = { "stage1": (128,128), "stage2": (192,192), "stage3": (256,256)}
74:
75:
76:     # Loss component scheduling
77:     loss_schedule: Dict[str, str] = None
78:
79:     # Multi-objective optimization (GradNorm)
80:     use_gradnorm: bool = True
81:     gradnorm_alpha: float = 0.12
82:     gradnorm_update_freq: int = 5
83:
84:     # Graph constraint scheduling
85:     graph_weight_schedule: str = "delayed_ramp"
86:     graph_start_epoch: int = 15
87:     graph_end_weight: float = 0.5
88:
89:     def __post_init__(self):
90:         # Provide default loss schedule if not set
91:         if self.loss_schedule is None:
92:             self.loss_schedule = {
93:                 "segmentation": "static",
94:                 "dice": "static",
95:                 "sdf": "early_decay",
96:                 "attributes": "static",
97:                 "polygon": "staged_ramp",
98:                 "voxel": "late_ramp",
99:                 "topology": "progressive",
100:                 "latent_consistency": "mid_ramp",
101:                 "graph": "delayed_ramp",
102:             }
103:
104:         # Default objectives used by GradNorm / trainer monitoring
105:         if self.objectives is None:
106:             self.objectives = [
107:                 "segmentation",
108:                 "dice",
109:                 "sdf",
110:                 "attributes",
111:                 "polygon",
112:                 "voxel",
113:                 "topology",
114:                 "latent_consistency",
115:                 "graph",
116:             ]
117:
118:
119: @dataclass
120: class TrainingConfig:
121:     """Training configuration with adaptive strategies"""
122:     device: str = "cuda" if torch.cuda.is_available() else "cpu"
123:
124:     # Dynamic epoch limits (maxima; curriculum may switch earlier)
125:     max_stage1_epochs: int = 40
126:     max_stage2_epochs: int = 25
127:     max_stage3_epochs: int = 60
128:
129:     # Minimum epochs per stage (avoid switching too early)
130:     min_stage1_epochs: int = 8
131:     min_stage2_epochs: int = 5
132:     min_stage3_epochs: int = 12
133:
134:     # Learning rates (per stage)
135:     stage1_lr: float = 1e-5 # was 3e-4
136:     stage1_weight_decay: float = 1e-5
137:
138:     stage2_lr: float = 5e-6 # was 1e-4
139:     stage2_weight_decay: float = 1e-5
140:

```

```

141:     stage3_lr: float = 1e-6 # was 5e-5
142:     stage3_weight_decay: float = 1e-5
143:
144:     # Advanced training techniques
145:     use_mixed_precision: bool = True
146:     use_cosine_restarts: bool = True
147:     warmup_epochs: int = 5
148:     grad_clip_norm: float = 0.5
149:
150:     # Gradient monitoring for dynamic weighting
151:     track_gradient_norms: bool = True
152:     gradient_norm_window: int = 10 # rolling window for gradient tracking
153:
154:     # Checkpointing
155:     checkpoint_freq: int = 1
156:
157:     # Curriculum configuration
158:     curriculum: CurriculumConfig = None
159:
160:     def __post_init__(self):
161:         if self.curriculum is None:
162:             self.curriculum = CurriculumConfig()
163:
164:
165: @dataclass
166: class LossConfig:
167:     """Loss function weights (will be dynamically adjusted during training)"""
168:     # Base weights (starting values)
169:     seg_weight: float = 1.0
170:     dice_weight: float = 1.0
171:     sdf_weight: float = 0.5
172:     attr_weight: float = 1.0
173:     polygon_weight: float = 1.0
174:     voxel_weight: float = 1.0
175:     topology_weight: float = 0.1 # start low, ramp up
176:
177:     # New loss components
178:     latent_consistency_weight: float = 0.5
179:     graph_constraint_weight: float = 0.3
180:
181:     # Dynamic weighting parameters
182:     enable_dynamic_weighting: bool = True
183:     weight_update_freq: int = 10
184:     weight_momentum: float = 0.9
185:
186:
187: @dataclass
188: class InferenceConfig:
189:     """Inference configuration"""
190:     model_path: str = "final_model.pth"
191:     test_images_dir: str = "./data/test_images"
192:     output_dir: str = "./outputs"
193:     export_intermediate: bool = True
194:     polygon_threshold: float = 0.5
195:
196:
197: # Curriculum stage transition logic
198: class StageTransitionCriteria:
199:     """Defines criteria for automatic stage transitions"""
200:
201:     @staticmethod
202:     def should_transition_from_stage1(train_losses, val_losses, config: CurriculumConfig) -> bool:
203:         """Check if should transition from Stage 1 to Stage 2"""
204:         if len(val_losses) < config.plateau_detection_window:
205:             return False
206:
207:         # Check for plateau in segmentation + dice losses
208:         recent_losses = val_losses[-config.plateau_detection_window:]
209:         if len(recent_losses) < 2:
210:             return False
211:
212:         # Calculate improvement rate
213:         old_avg = sum(recent_losses[:len(recent_losses)//2]) / (len(recent_losses)//2)

```

```

214:         new_avg = sum(recent_losses[len(recent_losses)//2:]) / (len(recent_losses) - len(recent_losses)//2)
215:
216:         improvement_rate = (old_avg - new_avg) / (old_avg + 1e-8)
217:         return improvement_rate < config.min_improvement_threshold
218:
219:     @staticmethod
220:     def should_transition_from_stage2(polygon_losses, config: CurriculumConfig) -> bool:
221:         """Check if should transition from Stage 2 to Stage 3"""
222:         if len(polygon_losses) < config.plateau_detection_window:
223:             return False
224:
225:         # Check polygon loss plateau
226:         recent_losses = polygon_losses[-config.plateau_detection_window:]
227:         if len(recent_losses) < 2:
228:             return False
229:
230:         old_avg = sum(recent_losses[:len(recent_losses)//2]) / (len(recent_losses)//2)
231:         new_avg = sum(recent_losses[len(recent_losses)//2:]) / (len(recent_losses) - len(recent_losses)//2)
232:
233:         improvement_rate = (old_avg - new_avg) / (old_avg + 1e-8)
234:         return improvement_rate < config.min_improvement_threshold
235:
236:
237: # Default configurations (import these in your trainer)
238: DEFAULT_DATA_CONFIG = DataConfig()
239: DEFAULT_MODEL_CONFIG = ModelConfig()
240: DEFAULT_TRAINING_CONFIG = TrainingConfig()
241: DEFAULT_LOSS_CONFIG = LossConfig()
242: DEFAULT_INFERENCE_CONFIG = InferenceConfig()

```

■ File: convert_polygons.py

```

=====
1: import json
2: from pathlib import Path
3:
4: def flatten_polygon_file(file_path: Path):
5:     with open(file_path, "r") as f:
6:         data = json.load(f)
7:
8:         flat_list = []
9:
10:        # If file already flat, skip
11:        if isinstance(data, list):
12:            print(f"[SKIP] Already flat: {file_path}")
13:            return
14:
15:        # Otherwise, flatten categories
16:        for category in ["walls", "doors", "windows", "floors", "fixtures"]:
17:            if category in data and isinstance(data[category], list):
18:                for idx, poly in enumerate(data[category]):
19:                    flat_list.append({
20:                        "id": idx,
21:                        "type": category[:-1], # "walls" -> "wall"
22:                        "points": poly.get("points", []),
23:                        "area": poly.get("area", None)
24:                    })
25:
26:        # Save back in flat format
27:        with open(file_path, "w") as f:
28:            json.dump(flat_list, f, indent=2)
29:
30:        print(f"[OK] Converted: {file_path}")
31:
32:
33: def batch_convert(root_dir="data/floorplans"):
34:     root = Path(root_dir)
35:     for polygon_file in root.rglob("polygon.json"):
36:         flatten_polygon_file(polygon_file)
37:
38:
39: if __name__ == "__main__":

```

```

40:     # Change path if needed
41:     batch_convert("data/floorplans")

```

File: dataset.py

```

=====
1: """
2: Dataset classes for the Neural-Geometric 3D Model Generator
3: Enhanced with in-memory caching for faster training
4: """
5:
6: import cv2
7: import json
8: import numpy as np
9: import torch
10: from torch.utils.data import Dataset
11: from pathlib import Path
12: from typing import Dict, List, Tuple, Optional, Union
13: import time
14:
15: from config import DEFAULT_DATA_CONFIG
16:
17:
18: class AdvancedFloorPlanDataset(Dataset):
19:     """
20:     Research-grade dataset with complete ground truth:
21:     - Floorplan image + segmentation mask
22:     - Attribute dictionary (geometric parameters)
23:     - Ground-truth mesh + voxelized occupancy
24:     - Polygon outlines for vectorization supervision
25:
26:     Enhanced with optional in-memory caching for performance
27:     """
28:
29:     def __init__(
30:         self,
31:         data_dir: str = None,
32:         split: str = "train",
33:         image_size: Tuple[int, int] = None,
34:         voxel_size: int = None,
35:         augment: bool = None,
36:         config=None,
37:     ):
38:         # Use config if provided, otherwise defaults from config.py
39:         if config is None:
40:             config = DEFAULT_DATA_CONFIG
41:
42:         self.data_dir = Path(data_dir or config.data_dir)
43:         self.split = split
44:         self.image_size = image_size or config.image_size
45:         self.voxel_size = voxel_size or config.voxel_size
46:         self.augment = (
47:             augment if augment is not None else config.augment
48:         ) and split == "train"
49:
50:         # Collect all samples that contain every required file
51:         self.samples = self._find_complete_samples()
52:         print(f"Found {len(self.samples)} complete samples for {self.split}")
53:
54:         # NEW: In-memory caching for performance
55:         self.cache_in_memory = getattr(config, "cache_in_memory", False)
56:         self._cache = None
57:
58:         if self.cache_in_memory and len(self.samples) > 0:
59:             print(f"[DATA] Preloading {len(self.samples)} samples into RAM (cache_in_memory=True).")
60:             print("[DATA] This may take significant memory but will speed up training...")
61:
62:             # Estimate memory usage
63:             estimated_mb = self._estimate_memory_usage()
64:             print(f"[DATA] Estimated memory usage: {estimated_mb:.1f} MB")
65:
66:             start_time = time.time()

```



```

67:         self._preload_cache()
68:         load_time = time.time() - start_time
69:         print(f"[DATA] Cache preloading completed in {load_time:.2f}s")
70:
71:     def _estimate_memory_usage(self):
72:         """Estimate memory usage for caching"""
73:         if not self.samples:
74:             return 0.0
75:
76:         H, W = self.image_size
77:         n_samples = len(self.samples)
78:
79:         # Rough estimates in bytes
80:         image_bytes = H * W * 3 # RGB uint8
81:         mask_bytes = H * W # grayscale uint8
82:         voxel_bytes = self.voxel_size ** 3 * 4 # float32
83:         json_bytes = 1024 # rough estimate for params + polygons
84:
85:         total_per_sample = image_bytes + mask_bytes + voxel_bytes + json_bytes
86:         total_mb = (total_per_sample * n_samples) / (1024 * 1024)
87:
88:         return total_mb
89:
90:     def _preload_cache(self):
91:         """Preload all samples into memory"""
92:         self._cache = []
93:
94:         for i, sample in enumerate(self.samples):
95:             if i % 100 == 0:
96:                 print(f"[DATA] Loading sample {i+1}/{len(self.samples)}")
97:
98:                 try:
99:                     # Load image
100:                    img = cv2.imread(str(sample["image"]))
101:                    if img is None:
102:                        print(f"Warning: Could not load image {sample['image']}")
103:                        continue
104:                    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
105:                    img = cv2.resize(img, self.image_size) # (W, H) format for cv2.resize
106:
107:                    # Load mask
108:                    mask = cv2.imread(str(sample["mask"]), cv2.IMREAD_GRAYSCALE)
109:                    if mask is None:
110:                        print(f"Warning: Could not load mask {sample['mask']}")
111:                        continue
112:                    mask = cv2.resize(mask, self.image_size, interpolation=cv2.INTER_NEAREST)
113:
114:                    # Load voxel data
115:                    try:
116:                        voxel_data = np.load(sample["voxel"])
117:                        vox = voxel_data["voxels"] # Keep as numpy array
118:                    except Exception as e:
119:                        print(f"Warning: Could not load voxel data {sample['voxel']}: {e}")
120:                        # Create dummy voxel data
121:                        vox = np.zeros((self.voxel_size, self.voxel_size, self.voxel_size), dtype=np.float32)
122:
123:                    # Load parameters
124:                    try:
125:                        with open(sample["params"], "r") as f:
126:                            params = json.load(f)
127:                    except Exception as e:
128:                        print(f"Warning: Could not load params {sample['params']}: {e}")
129:                        params = self._get_default_attributes()
130:
131:                    # Load polygons
132:                    try:
133:                        with open(sample["polygon"], "r") as f:
134:                            polygons = json.load(f)
135:                    except Exception as e:
136:                        print(f"Warning: Could not load polygons {sample['polygon']}: {e}")
137:                        polygons = {"walls": []}
138:
139:                    self._cache.append({

```

```

140:         "image": img,
141:         "mask": mask,
142:         "vox": vox,
143:         "params": params,
144:         "polygons": polygons,
145:         "sample_id": sample["image"].parent.name,
146:     })
147:
148:     except Exception as e:
149:         print(f"Error loading sample {i}: {e}")
150:         continue
151:
152: def _get_default_attributes(self):
153:     """Return default attributes for missing param files"""
154:     return {
155:         "wall_height": 2.6,
156:         "wall_thickness": 0.15,
157:         "window_base_height": 0.7,
158:         "window_height": 0.95,
159:         "door_height": 2.6,
160:         "pixel_scale": 0.02,
161:     }
162:
163: # -----
164: def _find_complete_samples(self):
165:     """Locate samples that contain all the expected files."""
166:     samples = []
167:     split_dir = self.data_dir / self.split
168:
169:     if not split_dir.exists():
170:         print(f"Warning: Split directory {split_dir} does not exist")
171:         return samples
172:
173:     for sample_dir in split_dir.iterdir():
174:         if not sample_dir.is_dir():
175:             continue
176:
177:         required_files = {
178:             "image": sample_dir / "image.png",
179:             "mask": sample_dir / "mask.png",
180:             "params": sample_dir / "params.json",
181:             "mesh": sample_dir / "model.obj",
182:             "voxel": sample_dir / "voxel_GT.npz",
183:             "polygon": sample_dir / "polygon.json",
184:         }
185:
186:         if all(f.exists() for f in required_files.values()):
187:             samples.append(required_files)
188:
189:     return samples
190:
191: # -----
192: def __len__(self):
193:     return len(self._cache) if self._cache is not None else len(self.samples)
194:
195: # -----
196: def __getitem__(self, idx):
197:     # Use cached data if available
198:     if self._cache is not None:
199:         cached_sample = self._cache[idx]
200:         image = cached_sample['image']
201:         mask = cached_sample['mask']
202:         vox = cached_sample['vox']
203:         attributes = cached_sample['params']
204:         polygons_gt = cached_sample['polygons']
205:         sample_id = cached_sample['sample_id']
206:     else:
207:         # Fallback: load from disk on-the-fly
208:         sample = self.samples[idx]
209:
210:         # Load image and mask
211:         image = cv2.imread(str(sample["image"]))
212:         image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

```

```

213:         image = cv2.resize(image, self.image_size)
214:
215:         mask = cv2.imread(str(sample["mask"]), cv2.IMREAD_GRAYSCALE)
216:         mask = cv2.resize(mask, self.image_size, interpolation=cv2.INTER_NEAREST)
217:
218:         # Load attributes
219:         with open(sample["params"], "r") as f:
220:             attributes = json.load(f)
221:
222:         # Load voxel ground truth
223:         voxel_data = np.load(sample["voxel"])
224:         vox = voxel_data["voxels"]
225:
226:         # Load polygon ground truth
227:         with open(sample["polygon"], "r") as f:
228:             polygons_gt = json.load(f)
229:
230:         sample_id = sample["image"].parent.name
231:
232:         # Normalize image to [0,1]
233:         image = image.astype(np.float32) / 255.0
234:
235:         # Clean mask (remove class 5 if present)
236:         mask[mask == 5] = 0
237:
238:         # Convert to tensors
239:         image_tensor = torch.from_numpy(image).float().permute(2, 0, 1)
240:         mask_tensor = torch.from_numpy(mask).long()
241:         voxels_tensor = torch.from_numpy(vox.astype(np.float32))
242:
243:         attr_tensor = self._process_attributes(attributes)
244:         polygon_tensor = self._process_polygons(polygons_gt)
245:
246:         # Apply augmentation if enabled
247:         if self.augment:
248:             image_tensor, mask_tensor = self._augment(image_tensor, mask_tensor)
249:
250:         # Add validation before returning
251:         self._validate_sample_data(idx, image_tensor, mask_tensor, attr_tensor, voxels_tensor, polygon_tensor)
252:
253:         return {
254:             "image": image_tensor,
255:             "mask": mask_tensor,
256:             "attributes": attr_tensor,
257:             "voxels_gt": voxels_tensor,
258:             "polygons_gt": polygon_tensor,
259:             "sample_id": sample_id,
260:         }
261:
262:         # -----
263:         def _validate_sample_data(self, idx, image, mask, attributes, voxels, polygons):
264:             """Validate sample data for NaN/Inf values"""
265:             tensors_to_check = [
266:                 ("image", image),
267:                 ("mask", mask),
268:                 ("attributes", attributes),
269:                 ("voxels", voxels),
270:                 ("polygons", polygons["polygons"])
271:             ]
272:
273:             corrupted_data = False
274:
275:             for name, tensor in tensors_to_check:
276:                 if torch.isnan(tensor).any():
277:                     print(f"ERROR: {name} contains NaN values at sample {idx}")
278:                     corrupted_data = True
279:                 if torch.isinf(tensor).any():
280:                     print(f"ERROR: {name} contains Inf values at sample {idx}")
281:                     corrupted_data = True
282:
283:             if corrupted_data:
284:                 print(f"WARNING: Corrupted data detected in sample {idx}, replacing with safe fallback values")
285:

```

```

286:         # Replace corrupted tensors with safe fallback values
287:     for name, tensor in tensors_to_check:
288:         if torch.isnan(tensor).any() or torch.isinf(tensor).any():
289:             if name == "image":
290:                 # Replace with zeros (black image)
291:                 image.data = torch.zeros_like(image)
292:             elif name == "mask":
293:                 # Replace with zeros (background class)
294:                 mask.data = torch.zeros_like(mask).long()
295:             elif name == "attributes":
296:                 # Replace with reasonable default values (0.5 normalized)
297:                 attributes.data = torch.ones_like(attributes) * 0.5
298:             elif name == "voxels":
299:                 # Replace with empty voxel grid
300:                 voxels.data = torch.zeros_like(voxels)
301:             elif name == "polygons":
302:                 # Replace polygons with zeros
303:                 polygons["polygons"].data = torch.zeros_like(polygons["polygons"])
304:
305: # -----
306: def _process_attributes(self, attributes):
307:     """Convert attribute dictionary to a normalized tensor."""
308:     # Normalize common architectural parameters into [0,1]
309:     attr_list = [
310:         attributes.get("wall_height", 2.6) / 5.0,
311:         attributes.get("wall_thickness", 0.15) / 0.5,
312:         attributes.get("window_base_height", 0.7) / 3.0,
313:         attributes.get("window_height", 0.95) / 2.0,
314:         attributes.get("door_height", 2.6) / 5.0,
315:         attributes.get("pixel_scale", 0.01) / 0.02,
316:     ]
317:
318:     # Ensure no NaN/Inf values in attribute processing
319:     safe_attr_list = []
320:     for val in attr_list:
321:         if np.isnan(val) or np.isinf(val):
322:             safe_attr_list.append(0.5) # Default normalized value
323:         else:
324:             safe_attr_list.append(max(0.0, min(1.0, val))) # Clamp to [0,1]
325:
326:     return torch.tensor(safe_attr_list, dtype=torch.float32)
327:
328: # -----
329: def _process_polygons(self, polygons_gt):
330:     """Convert polygon ground truth into a fixed tensor representation.
331:     Handles both formats:
332:     1. Nested dict: { "walls": [...], "doors": [...], ... }
333:     2. Flat list: [ {"type": "wall", "points": [...]}, ... ]
334:     """
335:     max_polygons = 30 # number of polygons per sample
336:     max_points = 100 # max points per polygon
337:
338:     processed = torch.zeros(max_polygons, max_points, 2)
339:     valid_mask = torch.zeros(max_polygons, dtype=torch.bool)
340:
341:     poly_idx = 0
342:
343:     try:
344:         # --- Case 1: dict format ---
345:         if isinstance(polygons_gt, dict):
346:             for class_name, polygon_list in polygons_gt.items():
347:                 if not isinstance(polygon_list, list):
348:                     continue
349:                 for polygon in polygon_list:
350:                     if poly_idx >= max_polygons:
351:                         break
352:                     if "points" not in polygon:
353:                         continue
354:
355:                     points = np.array(polygon["points"])
356:                     if len(points) > max_points:
357:                         # Subsample evenly if too many points
358:                         indices = np.linspace(0, len(points) - 1, max_points, dtype=int)

```

```

359:             points = points[indices]
360:
361:             # Check for NaN/Inf in points
362:             if np.any(np.isnan(points)) or np.any(np.isinf(points)):
363:                 print(f"Warning: Invalid polygon points detected, skipping polygon")
364:                 continue
365:
366:             # Normalize to [0,1] relative to image size
367:             points = points / np.array(self.image_size)
368:             # Clamp to valid range
369:             points = np.clip(points, 0.0, 1.0)
370:
371:             processed[poly_idx, : len(points)] = torch.from_numpy(points).float()
372:             valid_mask[poly_idx] = True
373:             poly_idx += 1
374:
375:         # --- Case 2: list format ---
376:         elif isinstance(polygons_gt, list):
377:             for polygon in polygons_gt:
378:                 if poly_idx >= max_polygons:
379:                     break
380:                 if "points" not in polygon:
381:                     continue
382:
383:                 points = np.array(polygon["points"])
384:                 if len(points) > max_points:
385:                     indices = np.linspace(0, len(points) - 1, max_points, dtype=int)
386:                     points = points[indices]
387:
388:                 # Check for NaN/Inf in points
389:                 if np.any(np.isnan(points)) or np.any(np.isinf(points)):
390:                     print(f"Warning: Invalid polygon points detected, skipping polygon")
391:                     continue
392:
393:                 points = points / np.array(self.image_size)
394:                 points = np.clip(points, 0.0, 1.0)
395:
396:                 processed[poly_idx, : len(points)] = torch.from_numpy(points).float()
397:                 valid_mask[poly_idx] = True
398:                 poly_idx += 1
399:
400:         except Exception as e:
401:             print(f"Warning: Error processing polygons: {e}")
402:             # Return safe empty polygon data
403:             processed = torch.zeros(max_polygons, max_points, 2)
404:             valid_mask = torch.zeros(max_polygons, dtype=torch.bool)
405:
406:         return {"polygons": processed, "valid_mask": valid_mask}
407:
408:     # -----
409:     def _augment(self, image, mask):
410:         """Enhanced data augmentation with rotations, flips, and intensity changes."""
411:         # Random rotation (multiples of 90° only for architectural data)
412:         if torch.rand(1) < 0.5:
413:             k = torch.randint(1, 4, (1,)).item()
414:             image = torch.rot90(image, k, dims=[1, 2])
415:             mask = torch.rot90(mask, k, dims=[0, 1])
416:
417:         # Random horizontal flip
418:         if torch.rand(1) < 0.5:
419:             image = torch.flip(image, dims=[2])
420:             mask = torch.flip(mask, dims=[1])
421:
422:         # Random vertical flip
423:         if torch.rand(1) < 0.5:
424:             image = torch.flip(image, dims=[1])
425:             mask = torch.flip(mask, dims=[0])
426:
427:         # Slight brightness/contrast adjustment with safety checks
428:         if torch.rand(1) < 0.3:
429:             brightness = torch.rand(1) * 0.2 - 0.1 # ±0.1
430:             contrast = torch.rand(1) * 0.2 + 0.9 # 0.9-1.1
431:             image = torch.clamp(image * contrast + brightness, 0, 1)

```

```

432:
433:         # Additional safety check for augmented image
434:         if torch.isnan(image).any() or torch.isinf(image).any():
435:             print("Warning: Augmentation produced invalid values, reverting to original")
436:             # Revert to safe values
437:             image = torch.clamp(image, 0, 1)
438:             image = torch.where(torch.isnan(image) | torch.isinf(image),
439:                               torch.zeros_like(image), image)
440:
441:         return image, mask
442:
443:     # -----
444:     def get_cache_info(self):
445:         """Return information about caching status"""
446:         return {
447:             "cache_enabled": self.cache_in_memory,
448:             "cache_loaded": self._cache is not None,
449:             "cached_samples": len(self._cache) if self._cache else 0,
450:             "total_samples": len(self.samples),
451:             "estimated_memory_mb": self._estimate_memory_usage() if self.cache_in_memory else 0
452:         }
453:
454:     def disable_cache(self):
455:         """Disable caching and free memory"""
456:         if self._cache is not None:
457:             print(f"[DATA] Disabling cache and freeing memory for {len(self._cache)} samples")
458:             self._cache = None
459:             self.cache_in_memory = False
460:
461:     def enable_cache(self):
462:         """Enable caching if not already enabled"""
463:         if not self.cache_in_memory and self.samples:
464:             self.cache_in_memory = True
465:             print("[DATA] Enabling cache...")
466:             self._preload_cache()
467:
468:
469:     # =====
470:     # Synthetic sample generator for testing without dataset
471:     # =====
472:     def create_synthetic_data_sample():
473:         """Generate a synthetic floorplan with attributes, voxels, and polygons."""
474:         image = np.ones((256, 256, 3), dtype=np.uint8) * 255
475:         mask = np.zeros((256, 256), dtype=np.uint8)
476:
477:         # Simple square room
478:         room_points = np.array([[50, 50], [200, 50], [200, 200], [50, 200]])
479:         cv2.fillPoly(mask, [room_points], 1) # Room = class 1
480:         cv2.polylines(image, [room_points], True, (0, 0, 0), 3)
481:
482:         # Add door
483:         cv2.rectangle(mask, (90, 50), (110, 70), 2) # Door = class 2
484:         cv2.rectangle(image, (90, 50), (110, 70), (255, 0, 0), -1)
485:
486:         # Attributes
487:         attributes = {
488:             "wall_height": 2.6,
489:             "wall_thickness": 0.15,
490:             "window_base_height": 0.7,
491:             "window_height": 0.95,
492:             "door_height": 2.6,
493:             "pixel_scale": 0.02,
494:         }
495:
496:         # Simple voxel GT
497:         voxels = np.zeros((64, 64, 64), dtype=bool)
498:         voxels[:20, 10:50, 10:50] = True
499:
500:         # Polygon GT
501:         polygons = {"walls": [{"points": room_points.tolist()}]}
502:
503:         return image, mask, attributes, voxels, polygons
504:

```

```

505:
506: class SyntheticFloorPlanDataset(Dataset):
507:     """
508:     Synthetic dataset for testing and development when real data is not available
509:     """
510:
511:     def __init__(self, num_samples=1000, image_size=(256, 256), voxel_size=64):
512:         self.num_samples = num_samples
513:         self.image_size = image_size
514:         self.voxel_size = voxel_size
515:
516:     def __len__(self):
517:         return self.num_samples
518:
519:     def __getitem__(self, idx):
520:         # Generate deterministic synthetic data based on index
521:         np.random.seed(idx)
522:         torch.manual_seed(idx)
523:
524:         image, mask, attributes, voxels, polygons_gt = create_synthetic_data_sample()
525:
526:         # Convert to tensors
527:         image_tensor = torch.from_numpy(image.astype(np.float32) / 255.0).permute(2, 0, 1)
528:         mask_tensor = torch.from_numpy(mask).long()
529:         voxels_tensor = torch.from_numpy(voxels.astype(np.float32))
530:
531:         # Process attributes and polygons using same methods as main dataset
532:         dataset = AdvancedFloorPlanDataset.__new__(AdvancedFloorPlanDataset)
533:         dataset.image_size = self.image_size
534:
535:         attr_tensor = dataset._process_attributes(attributes)
536:         polygon_tensor = dataset._process_polygons(polygons_gt)
537:
538:         return {
539:             "image": image_tensor,
540:             "mask": mask_tensor,
541:             "attributes": attr_tensor,
542:             "voxels_gt": voxels_tensor,
543:             "polygons_gt": polygon_tensor,
544:             "sample_id": f"synthetic_{idx:06d}",
545:         }

```

File: demo.py

```

=====
1: """
2: Demo script for the Neural-Geometric 3D Model Generator
3: """
4:
5: import torch
6: import cv2
7: import numpy as np
8: from pathlib import Path
9:
10: from models.model import NeuralGeometric3DGenerator
11: from dataset import create_synthetic_data_sample
12: from utils.visualization import visualize_predictions, create_model_summary_report
13:
14:
15: def demo_pipeline():
16:     """Demonstrate the complete pipeline with synthetic data"""
17:     print("Neural-Geometric 3D Model Generator Demo")
18:     print("=" * 50)
19:
20:     # Create output directory
21:     demo_dir = Path("./demo_outputs")
22:     demo_dir.mkdir(exist_ok=True)
23:
24:     # Create synthetic sample
25:     print("Creating synthetic data sample...")
26:     image, mask, attributes, voxels, polygons = create_synthetic_data_sample()
27:

```

```

28:     # Save synthetic data
29:     cv2.imwrite(str(demo_dir / "demo_input.png"), image)
30:     cv2.imwrite(str(demo_dir / "demo_mask.png"), mask * 50)
31:
32:     # Create model (random weights for demo)
33:     print("Initializing model...")
34:     model = NeuralGeometric3DGenerator()
35:     model.eval()
36:
37:     # Convert to tensors
38:     image_tensor = torch.from_numpy(image / 255.0).float().permute(2, 0, 1).unsqueeze(0)
39:
40:     # Create model summary
41:     create_model_summary_report(model, image_tensor, str(demo_dir / "model_summary.txt"))
42:
43:     # Forward pass
44:     print("Running forward pass...")
45:     with torch.no_grad():
46:         predictions = model(image_tensor)
47:
48:     print("Forward pass completed")
49:     print(f"Segmentation shape: {predictions['segmentation'].shape}")
50:     print(f"Attributes shape: {predictions['attributes'].shape}")
51:     print(f"Polygons shape: {predictions['polygons'].shape}")
52:     print(f"Voxels shape: {predictions['voxels_pred'].shape}")
53:
54:     # Extract and save results
55:     seg_pred = torch.argmax(predictions["segmentation"], dim=1).squeeze().numpy()
56:     attr_pred = predictions["attributes"].squeeze().numpy()
57:
58:     cv2.imwrite(str(demo_dir / "demo_seg_pred.png"), seg_pred * 50)
59:
60:     print(f"Predicted attributes: {attr_pred}")
61:
62:     # Create visualization
63:     print("Creating visualizations...")
64:
65:     # Create targets for visualization
66:     targets = {
67:         "mask": torch.from_numpy(mask).unsqueeze(0),
68:         "attributes": torch.from_numpy(np.array([
69:             attributes["wall_height"] / 5.0,
70:             attributes["wall_thickness"] / 0.5,
71:             attributes["window_base_height"] / 3.0,
72:             attributes["window_height"] / 2.0,
73:             attributes["door_height"] / 5.0,
74:             attributes["pixel_scale"] / 0.02,
75:         ])).float().unsqueeze(0)
76:     }
77:
78:     visualize_predictions(
79:         image_tensor,
80:         predictions,
81:         targets,
82:         save_path=str(demo_dir / "demo_predictions.png")
83:     )
84:
85:     print(f"Demo completed successfully! Results saved to {demo_dir}")
86:
87:
88: def demo_with_pretrained(model_path, input_image_path=None):
89:     """Demo with a pretrained model"""
90:     if not Path(model_path).exists():
91:         print(f"Model file {model_path} not found!")
92:         return
93:
94:     print(f"Running demo with pretrained model: {model_path}")
95:
96:     # Load model
97:     model = NeuralGeometric3DGenerator()
98:     checkpoint = torch.load(model_path, map_location="cuda")
99:     model.load_state_dict(checkpoint["model_state_dict"])
100:    model.cuda()

```



```

101:     model.eval()
102:
103:     # Use provided image or create synthetic
104:     if input_image_path and Path(input_image_path).exists():
105:         image = cv2.imread(input_image_path)
106:         image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
107:         image = cv2.resize(image, (256, 256))
108:         image_tensor = torch.from_numpy(image / 255.0).float().permute(2, 0, 1).unsqueeze(0).cuda()
109:         print(f"Using input image: {input_image_path}")
110:     else:
111:         print("Using synthetic data...")
112:         image, _, _, _, _ = create_synthetic_data_sample()
113:         image_tensor = torch.from_numpy(image / 255.0).float().permute(2, 0, 1).unsqueeze(0).cuda()
114:
115:     # Run inference
116:     with torch.no_grad():
117:         predictions = model(image_tensor)
118:
119:     # Visualize results
120:     demo_dir = Path("./demo_outputs")
121:     demo_dir.mkdir(exist_ok=True)
122:
123:     visualize_predictions(
124:         image_tensor,
125:         predictions,
126:         save_path=str(demo_dir / "pretrained_demo.png")
127:     )
128:
129:     print(f"Pretrained demo completed! Results saved to {demo_dir}")
130:
131:
132: if __name__ == "__main__":
133:     import argparse
134:
135:     parser = argparse.ArgumentParser(description="Demo Neural-Geometric 3D Model Generator")
136:     parser.add_argument("--model_path", type=str, default=None,
137:                         help="Path to pretrained model (optional)")
138:     parser.add_argument("--input_image", type=str, default=None,
139:                         help="Input image path (optional)")
140:
141:     args = parser.parse_args()
142:
143:     if args.model_path:
144:         demo_with_pretrained(args.model_path, args.input_image)
145:     else:
146:         demo_pipeline()

```

■ File: evaluate.py

```

=====
1: """
2: evaluate.py
3: Comprehensive evaluation CLI for Neural-Geometric 3D Model Generator.
4:
5: Usage examples:
6:     python evaluate.py --model_path checkpoints/final_model.pth --data_dir ./data/floorplans
7:     python evaluate.py --model_path checkpoints/final_model.pth --data_dir ./data/floorplans --visualize --sa
8: """
9:
10: import argparse
11: import json
12: from pathlib import Path
13: from typing import List, Dict
14:
15: import torch
16: from torch.utils.data import DataLoader
17:
18: # Project imports (match your repo layout)
19: from dataset import AdvancedFloorPlanDataset
20: from evaluation.metrics import ModelEvaluator
21: from inference.engine import ResearchInferenceEngine
22: from utils.visualization import visualize_predictions

```

```

23:
24:
25: def save_json(obj, path: Path):
26:     path.parent.mkdir(parents=True, exist_ok=True)
27:     with open(path, "w") as f:
28:         json.dump(obj, f, indent=2)
29:     print(f"[?] Saved JSON -> {path}")
30:
31:
32: def gather_per_sample_metrics(
33:     evaluator: ModelEvaluator,
34:     dataset: AdvancedFloorPlanDataset,
35:     device: str,
36:     max_samples: int = None,
37: ) -> List[Dict]:
38:     """
39:     Re-run evaluation loop sample-by-sample and collect per-sample metrics.
40:     We use evaluator._evaluate_* helper methods (present in evaluation/metrics.py)
41:     so metrics match the overall evaluation.
42:     """
43:     loader = DataLoader(dataset, batch_size=1, shuffle=False)
44:     per_sample = []
45:
46:     with torch.no_grad():
47:         for idx, batch in enumerate(loader):
48:             if max_samples is not None and idx >= max_samples:
49:                 break
50:
51:             # Move tensors to device where applicable
52:             batch_for_model = {}
53:             for k, v in batch.items():
54:                 if torch.is_tensor(v):
55:                     batch_for_model[k] = v.to(device)
56:                 else:
57:                     batch_for_model[k] = v # dicts/strings stay as-is
58:
59:             # Forward
60:             preds = evaluator.model(batch_for_model["image"])
61:
62:             # segmentation
63:             seg_res = evaluator._evaluate_segmentation(preds["segmentation"], batch_for_model["mask"])
64:             # attributes
65:             attr_res = evaluator._evaluate_attributes(preds["attributes"], batch_for_model["attributes"].to(device))
66:             # voxels
67:             voxel_res = evaluator._evaluate_voxels(preds["voxels_pred"], batch_for_model["voxels_gt"].to(device))
68:             # polygons ? evaluator._evaluate_polygons expects format used in metrics.py
69:             # batch["polygons_gt"] is a dict with "polygons" and "valid_mask"
70:             poly_res = evaluator._evaluate_polygons(preds["polygons"], preds.get("polygon_validity", preds["valid_mask"]))
71:
72:             sample_id = batch["sample_id"][0] if isinstance(batch["sample_id"], (list, tuple)) else batch["sample_id"]
73:             sample_metrics = {
74:                 "sample_id": str(sample_id),
75:                 "segmentation": seg_res,
76:                 "attributes": attr_res,
77:                 "voxels": voxel_res,
78:                 "polygons": poly_res,
79:             }
80:             per_sample.append(sample_metrics)
81:
82:             if (idx + 1) % 10 == 0:
83:                 print(f"[INFO] Collected per-sample metrics for {idx+1}/{len(loader)} samples")
84:
85:     return per_sample
86:
87:
88: def run_visualization_and_exports(
89:     engine: ResearchInferenceEngine,
90:     dataset: AdvancedFloorPlanDataset,
91:     output_dir: Path,
92:     device: str,
93:     num_viz: int = 10,
94:     max_export: int = 5,
95: ):

```

```

96:     """
97:     For the first `num_viz` samples, create visualizations using the model and optionally
98:     run deterministic 3D export to save intermediate results and a .obj.
99:     """
100:    output_dir.mkdir(parents=True, exist_ok=True)
101:    loader = DataLoader(dataset, batch_size=1, shuffle=False)
102:
103:    viz_count = 0
104:    export_count = 0
105:
106:    with torch.no_grad():
107:        for idx, batch in enumerate(loader):
108:            sample_id = batch["sample_id"][0] if isinstance(batch["sample_id"], (list, tuple)) else batch["sample_id"]
109:            sample_dir = Path(dataset.data_dir) / "test" / str(sample_id)
110:
111:            # Prepare tensors
112:            image_tensor = batch["image"].to(device)
113:            target_mask = batch["mask"].unsqueeze(0) if torch.is_tensor(batch["mask"]) else None
114:            target_attrs = batch["attributes"].unsqueeze(0) if torch.is_tensor(batch["attributes"]) else None
115:
116:            # Model predictions using engine.model (same underlying model)
117:            preds = engine.model(image_tensor)
118:
119:            # Visualization
120:            if viz_count < num_viz:
121:                vis_path = output_dir / f"viz_{sample_id}.png"
122:                try:
123:                    visualize_predictions(
124:                        image_tensor,
125:                        preds,
126:                        {"mask": target_mask, "attributes": target_attrs},
127:                        save_path=str(vis_path),
128:                    )
129:                    print(f"[?] Saved visualization for sample {sample_id} -> {vis_path}")
130:                except Exception as e:
131:                    print(f"[!] Visualization failed for {sample_id}: {e}")
132:                viz_count += 1
133:
134:            # Export deterministic 3D (uses the image file path)
135:            if export_count < max_export:
136:                image_file = sample_dir / "image.png"
137:                out_obj = output_dir / f"{sample_id}_predicted_model.obj"
138:                try:
139:                    success = engine.generate_3d_model(str(image_file), str(out_obj), export_intermediate=True)
140:                    if success:
141:                        print(f"[?] Exported deterministic 3D model for {sample_id} -> {out_obj}")
142:                    else:
143:                        print(f"[!] 3D export returned False for {sample_id}")
144:                except Exception as e:
145:                    print(f"[!] 3D export failed for {sample_id}: {e}")
146:                export_count += 1
147:
148:            if viz_count >= num_viz and export_count >= max_export:
149:                break
150:
151:
152: def main():
153:     parser = argparse.ArgumentParser(description="Evaluate Neural-Geometric 3D model")
154:     parser.add_argument("--model_path", "-m", required=True, help="Path to model checkpoint (checkpoint.pth)")
155:     parser.add_argument("--data_dir", "-d", default="./data/floorplans", help="Dataset root with train/val/test splits")
156:     parser.add_argument("--device", default=None, help="Device to use (cuda or cpu). Auto-detect if omitted")
157:     parser.add_argument("--visualize", action="store_true", help="Save visual comparison images (pred vs GT)")
158:     parser.add_argument("--save_outputs", action="store_true", help="Run deterministic 3D export for some samples")
159:     parser.add_argument("--output_dir", default="./evaluation_outputs", help="Where to save reports/visuals")
160:     parser.add_argument("--num_viz", type=int, default=10, help="How many visualizations to produce (default=10)")
161:     parser.add_argument("--max_exports", type=int, default=3, help="How many deterministic 3D exports to run")
162:     parser.add_argument("--per_sample_json", action="store_true", help="Save per-sample metrics JSON (may be useful for debugging)")
163:     parser.add_argument("--limit_samples", type=int, default=None, help="If set, limit evaluation to first N samples")
164:
165:     args = parser.parse_args()
166:
167:     # Determine device
168:     device = args.device or ("cuda" if torch.cuda.is_available() else "cpu")

```

```

169:     print(f"[INFO] Using device: {device}")
170:
171:     model_path = Path(args.model_path)
172:     if not model_path.exists():
173:         print(f"[ERROR] Model not found at: {model_path}")
174:         return
175:
176:     # Load test dataset
177:     dataset = AdvancedFloorPlanDataset(data_dir=args.data_dir, split="test")
178:     if len(dataset) == 0:
179:         print("[ERROR] No test samples found (dataset may be empty or data_dir incorrect).")
180:         return
181:
182:     # If user asked for a limited quick run, slice dataset.samples accordingly.
183:     if args.limit_samples is not None:
184:         # Create a shallow copy dataset pointing to first N samples
185:         dataset.samples = dataset.samples[: args.limit_samples]
186:         print(f"[INFO] Limiting evaluation to first {len(dataset)} samples")
187:
188:     # Create evaluator and run full evaluation
189:     evaluator = ModelEvaluator(str(model_path), device=device)
190:     summary = evaluator.evaluate_dataset(dataset)
191:     evaluator.print_evaluation_results(summary)
192:
193:     # Save summary JSON
194:     out_dir = Path(args.output_dir)
195:     out_dir.mkdir(parents=True, exist_ok=True)
196:     save_json(summary, out_dir / f"{model_path.stem}_summary.json")
197:
198:     # Optionally collect detailed per-sample metrics
199:     if args.per_sample_json:
200:         print("[INFO] Collecting per-sample metrics (this re-runs model inference sample-by-sample)...")
201:         per_sample = gather_per_sample_metrics(evaluator, dataset, device, max_samples=None)
202:         save_json(per_sample, out_dir / f"{model_path.stem}_per_sample_metrics.json")
203:
204:     # Visualization and/or exports
205:     if args.visualize or args.save_outputs:
206:         print("[INFO] Initializing inference engine for visualizations/exports...")
207:         engine = ResearchInferenceEngine(model_path=str(model_path), device=device)
208:         run_visualization_and_exports(
209:             engine,
210:             dataset,
211:             out_dir,
212:             device,
213:             num_viz=args.num_viz,
214:             max_export=args.max_exports,
215:         )
216:
217:     print("[?] Evaluation finished.")
218:
219:
220: if __name__ == "__main__":
221:     main()

```

File: evaluation\metrics.py

```

=====
1: """
2: Evaluation metrics and utilities for the Neural-Geometric 3D Model Generator
3: """
4:
5: import torch
6: import numpy as np
7: from torch.utils.data import DataLoader
8:
9: from models.model import NeuralGeometric3DGenerator
10: from dataset import AdvancedFloorPlanDataset
11:
12:
13: def compute_iou(pred, target):
14:     """Compute IoU for segmentation"""
15:     intersection = (pred & target).float().sum()

```

```

16:     union = (pred | target).float().sum()
17:     return (intersection / (union + 1e-6)).item()
18:
19:
20: def compute_3d_iou(pred, target):
21:     """Compute 3D IoU for voxel grids"""
22:     pred_bool = pred.bool()
23:     target_bool = target.bool()
24:
25:     intersection = (pred_bool & target_bool).float().sum()
26:     union = (pred_bool | target_bool).float().sum()
27:
28:     return (intersection / (union + 1e-6)).item()
29:
30:
31: def compute_polygon_metrics(pred_polygons, gt_polygons, validity_pred, validity_gt):
32:     """Compute metrics for polygon prediction"""
33:     # Chamfer distance between polygon sets
34:     valid_pred = pred_polygons[validity_pred > 0.5]
35:     valid_gt = gt_polygons[validity_gt]
36:
37:     if len(valid_pred) == 0 or len(valid_gt) == 0:
38:         return {"chamfer_distance": float('inf'), "validity_accuracy": 0.0}
39:
40:     # Simplified chamfer distance computation
41:     chamfer_dist = 0.0
42:     for pred_poly in valid_pred:
43:         min_dist = float('inf')
44:         for gt_poly in valid_gt:
45:             dist = torch.norm(pred_poly - gt_poly, dim=-1).min().item()
46:             min_dist = min(min_dist, dist)
47:             chamfer_dist += min_dist
48:
49:     chamfer_dist /= len(valid_pred)
50:
51:     # Validity accuracy
52:     validity_acc = ((validity_pred > 0.5) == validity_gt).float().mean().item()
53:
54:     return {
55:         "chamfer_distance": chamfer_dist,
56:         "validity_accuracy": validity_acc
57:     }
58:
59:
60: class ModelEvaluator:
61:     """Comprehensive model evaluation"""
62:
63:     def __init__(self, model_path, device="cuda"):
64:         self.device = device
65:         self.model = NeuralGeometric3DGenerator()
66:
67:         # Load model
68:         checkpoint = torch.load(model_path, map_location=device)
69:         self.model.load_state_dict(checkpoint["model_state_dict"])
70:         self.model.to(device)
71:         self.model.eval()
72:
73:         print(f"Loaded model from {model_path}")
74:
75:     def evaluate_dataset(self, test_dataset):
76:         """Comprehensive evaluation on test dataset"""
77:         test_loader = DataLoader(test_dataset, batch_size=1, shuffle=False)
78:
79:         # Metrics storage
80:         metrics = {
81:             "segmentation": {"ious": [], "class_ious": []},
82:             "attributes": {"maes": [], "mses": []},
83:             "voxels": {"ious": [], "dice_scores": []},
84:             "polygons": {"chamfer_distances": [], "validity_accs": []},
85:         }
86:
87:         with torch.no_grad():
88:             for batch_idx, batch in enumerate(test_loader):

```

```

89:         batch = {k: v.to(self.device) if torch.is_tensor(v) else v
90:                   for k, v in batch.items()}
91:
92:         predictions = self.model(batch["image"])
93:
94:         # Evaluate segmentation
95:         seg_metrics = self._evaluate_segmentation(
96:             predictions["segmentation"], batch["mask"]
97:         )
98:         metrics["segmentation"]["ious"].append(seg_metrics["iou"])
99:         metrics["segmentation"]["class_ious"].append(seg_metrics["class_ious"])
100:
101:         # Evaluate attributes
102:         attr_metrics = self._evaluate_attributes(
103:             predictions["attributes"], batch["attributes"]
104:         )
105:         metrics["attributes"]["maes"].append(attr_metrics["mae"])
106:         metrics["attributes"]["mses"].append(attr_metrics["mse"])
107:
108:         # Evaluate voxels
109:         voxel_metrics = self._evaluate_voxels(
110:             predictions["voxels_pred"], batch["voxels_gt"]
111:         )
112:         metrics["voxels"]["ious"].append(voxel_metrics["iou"])
113:         metrics["voxels"]["dice_scores"].append(voxel_metrics["dice"])
114:
115:         # Evaluate polygons
116:         poly_metrics = self._evaluate_polygons(
117:             predictions["polygons"],
118:             predictions["polygon_validity"],
119:             batch["polygons_gt"]
120:         )
121:         metrics["polygons"]["chamfer_distances"].append(poly_metrics["chamfer_distance"])
122:         metrics["polygons"]["validity_accs"].append(poly_metrics["validity_accuracy"])
123:
124:         if (batch_idx + 1) % 10 == 0:
125:             print(f"Evaluated {batch_idx + 1}/{len(test_loader)} samples")
126:
127:         return self._compute_summary_metrics(metrics)
128:
129:     def _evaluate_segmentation(self, pred_seg, target_mask):
130:         """Evaluate segmentation performance"""
131:         seg_pred = torch.argmax(pred_seg, dim=1)
132:
133:         # Overall IoU
134:         overall_iou = compute_iou(seg_pred, target_mask)
135:
136:         # Per-class IoU
137:         num_classes = pred_seg.shape[1]
138:         class_ious = []
139:
140:         for c in range(num_classes):
141:             pred_c = (seg_pred == c)
142:             target_c = (target_mask == c)
143:
144:             if target_c.sum() > 0: # Only compute if class exists in ground truth
145:                 iou_c = compute_iou(pred_c, target_c)
146:                 class_ious.append(iou_c)
147:
148:         return {
149:             "iou": overall_iou,
150:             "class_ious": class_ious
151:         }
152:
153:     def _evaluate_attributes(self, pred_attrs, target_attrs):
154:         """Evaluate attribute prediction"""
155:         mae = torch.mean(torch.abs(pred_attrs - target_attrs)).item()
156:         mse = torch.mean((pred_attrs - target_attrs) ** 2).item()
157:
158:         return {"mae": mae, "mse": mse}
159:
160:     def _evaluate_voxels(self, pred_voxels, target_voxels):
161:         """Evaluate 3D voxel prediction"""

```

```

162:         pred_binary = (torch.sigmoid(pred_voxels) > 0.5).float()
163:         target_float = target_voxels.float()
164:
165:         # 3D IoU
166:         iou_3d = compute_3d_iou(pred_binary, target_float)
167:
168:         # 3D Dice score
169:         intersection = (pred_binary * target_float).sum()
170:         dice = (2 * intersection) / (pred_binary.sum() + target_float.sum() + 1e-6)
171:
172:         return {
173:             "iou": iou_3d,
174:             "dice": dice.item()
175:         }
176:
177:     def _evaluate_polygons(self, pred_polygons, pred_validity, gt_polygons):
178:         """Evaluate polygon prediction"""
179:         return compute_polygon_metrics(
180:             pred_polygons[0],
181:             gt_polygons["polygons"][0],
182:             pred_validity[0],
183:             gt_polygons["valid_mask"][0]
184:         )
185:
186:     def _compute_summary_metrics(self, metrics):
187:         """Compute summary statistics"""
188:         summary = {}
189:
190:         # Segmentation
191:         summary["segmentation_mIoU"] = np.mean(metrics["segmentation"]["ious"])
192:         summary["segmentation_std"] = np.std(metrics["segmentation"]["ious"])
193:
194:         # Attributes
195:         summary["attribute_MAE"] = np.mean(metrics["attributes"]["maes"])
196:         summary["attribute_MAE_std"] = np.std(metrics["attributes"]["maes"])
197:
198:         # Voxels
199:         summary["voxel_mIoU"] = np.mean(metrics["voxels"]["ious"])
200:         summary["voxel_mIoU_std"] = np.std(metrics["voxels"]["ious"])
201:         summary["voxel_dice"] = np.mean(metrics["voxels"]["dice_scores"])
202:
203:         # Polygons
204:         valid_chamfer = [d for d in metrics["polygons"]["chamfer_distances"] if d != float('inf')]
205:         if valid_chamfer:
206:             summary["polygon_chamfer"] = np.mean(valid_chamfer)
207:             summary["polygon_chamfer_std"] = np.std(valid_chamfer)
208:         else:
209:             summary["polygon_chamfer"] = float('inf')
210:             summary["polygon_chamfer_std"] = 0.0
211:
212:         summary["polygon_validity_acc"] = np.mean(metrics["polygons"]["validity_accs"])
213:
214:         return summary
215:
216:     def print_evaluation_results(self, summary):
217:         """Print formatted evaluation results"""
218:         print("=" * 60)
219:         print("COMPREHENSIVE EVALUATION RESULTS")
220:         print("=" * 60)
221:
222:         print(f"Segmentation mIoU: {summary['segmentation_mIoU']:.4f} ± {summary['segmentation_std']:.4f}")
223:         print(f"Attribute MAE: {summary['attribute_MAE']:.4f} ± {summary['attribute_MAE_std']:.4f}")
224:         print(f"Voxel 3D mIoU: {summary['voxel_mIoU']:.4f} ± {summary['voxel_mIoU_std']:.4f}")
225:         print(f"Voxel Dice Score: {summary['voxel_dice']:.4f}")
226:
227:         if summary['polygon_chamfer'] != float('inf'):
228:             print(f"Polygon Chamfer Distance: {summary['polygon_chamfer']:.4f} ± {summary['polygon_chamfer_std']:.4f}")
229:         else:
230:             print("Polygon Chamfer Distance: No valid polygons")
231:
232:         print(f"Polygon Validity Accuracy: {summary['polygon_validity_acc']:.4f}")
233:         print("=" * 60)
234:

```

```

235:
236: def evaluate_model(model_path, data_dir="./data/floorplans"):
237:     """Standalone evaluation function"""
238:     # Load test dataset
239:     test_dataset = AdvancedFloorPlanDataset(data_dir, split="test")
240:
241:     if len(test_dataset) == 0:
242:         print("No test samples found!")
243:         return None
244:
245:     # Create evaluator
246:     evaluator = ModelEvaluator(model_path)
247:
248:     # Run evaluation
249:     summary = evaluator.evaluate_dataset(test_dataset)
250:
251:     # Print results
252:     evaluator.print_evaluation_results(summary)
253:
254:     return summary

```

File: infer.py

```

1: """
2: Main inference script for generating 3D models from 2D floorplans
3: """
4:
5: import argparse
6: from pathlib import Path
7:
8: from inference.engine import ResearchInferenceEngine
9: from config import DEFAULT_INFERENCE_CONFIG
10:
11:
12: def main():
13:     parser = argparse.ArgumentParser(description="Generate 3D models from 2D floorplans")
14:     parser.add_argument("--model_path", type=str, default="final_model.pth",
15:                         help="Path to trained model")
16:     parser.add_argument("--input", type=str, required=True,
17:                         help="Input image path or directory")
18:     parser.add_argument("--output", type=str, required=True,
19:                         help="Output path or directory")
20:     parser.add_argument("--device", type=str, default="cuda",
21:                         help="Inference device")
22:     parser.add_argument("--export_intermediate", action="store_true",
23:                         help="Export intermediate results")
24:     parser.add_argument("--polygon_threshold", type=float, default=0.5,
25:                         help="Threshold for polygon validity")
26:
27:     args = parser.parse_args()
28:
29:     # Initialize inference engine
30:     print(f"Initializing inference engine...")
31:     engine = ResearchInferenceEngine(
32:         model_path=args.model_path,
33:         device=args.device
34:     )
35:
36:     input_path = Path(args.input)
37:     output_path = Path(args.output)
38:
39:     if input_path.is_file():
40:         # Single image inference
41:         print(f"Processing single image: {input_path}")
42:
43:         if not output_path.suffix:
44:             output_path = output_path / f"{input_path.stem}_model.obj"
45:
46:         success = engine.generate_3d_model(
47:             str(input_path),
48:             str(output_path),

```



```

49:         export_intermediate=args.export_intermediate
50:     )
51:
52:     if success:
53:         print(f"? Successfully generated: {output_path}")
54:     else:
55:         print(f"? Failed to generate model for: {input_path}")
56:
57: elif input_path.is_dir():
58:     # Batch processing
59:     print(f"Processing directory: {input_path}")
60:
61:     # Find all image files
62:     image_extensions = ['.png', '.jpg', '.jpeg', '.bmp', '.tiff']
63:     image_files = [
64:         f for f in input_path.iterdir()
65:         if f.suffix.lower() in image_extensions
66:     ]
67:
68:     if not image_files:
69:         print("No image files found in input directory!")
70:         return
71:
72:     print(f"Found {len(image_files)} image files")
73:
74:     # Create output directory
75:     output_path.mkdir(exist_ok=True)
76:
77:     # Process batch
78:     results = engine.process_batch(image_files, output_path)
79:
80:     # Print summary
81:     successful = sum(1 for r in results if r["success"])
82:     print(f"\nBatch processing completed:")
83:     print(f"? Successful: {successful}/{len(results)}")
84:     print(f"? Failed: {len(results) - successful}/{len(results)}")
85:
86:     # List failed cases
87:     failed_cases = [r for r in results if not r["success"]]
88:     if failed_cases:
89:         print("\nFailed cases:")
90:         for case in failed_cases:
91:             error_msg = case.get("error", "Unknown error")
92:             print(f" - {Path(case['input']).name}: {error_msg}")
93:
94:     else:
95:         print(f"Error: Input path {input_path} does not exist!")
96:
97:
98: if __name__ == "__main__":
99:     main()

```

File: inference\engine.py

```

=====
1: """
2: Research-grade inference engine for 2D to 3D floorplan generation
3: """
4:
5: import torch
6: import cv2
7: import numpy as np
8: import json
9: import trimesh
10: from pathlib import Path
11:
12: from models.model import NeuralGeometric3DGenerator
13: from config import DEFAULT_INFERENCE_CONFIG
14:
15:
16: class ResearchInferenceEngine:
17:     """

```

```

18: Complete inference system that converts 2D floorplans to 3D models
19: following the deterministic export pipeline
20: """
21:
22: def __init__(self, model_path=None, device="cuda", config=None):
23:     if config is None:
24:         config = DEFAULT_INFERENCE_CONFIG
25:
26:     self.device = device
27:     self.config = config
28:     self.model = NeuralGeometric3DGenerator()
29:
30:     # Load trained model
31:     model_path = model_path or config.model_path
32:     checkpoint = torch.load(model_path, map_location=device)
33:     self.model.load_state_dict(checkpoint["model_state_dict"])
34:     self.model.to(device)
35:     self.model.eval()
36:
37:     print(f"Loaded trained model from {model_path}")
38:
39: def generate_3d_model(
40:     self,
41:     image_path: str,
42:     output_path: str,
43:     export_intermediate: bool = None
44: ):
45:     """
46:     Complete pipeline: Image -> Segmentation -> Polygons -> 3D Model
47:     """
48:     export_intermediate = export_intermediate or self.config.export_intermediate
49:
50:     # Load and preprocess image
51:     image = self._load_image(image_path)
52:
53:     with torch.no_grad():
54:         # Neural network inference
55:         predictions = self.model(image)
56:
57:         # Extract predictions
58:         segmentation = predictions["segmentation"]
59:         attributes = predictions["attributes"]
60:         polygons = predictions["polygons"]
61:         validity = predictions["polygon_validity"]
62:
63:         print("Neural network inference complete")
64:
65:         # Convert to deterministic representations
66:         mask_np = self._extract_mask(segmentation)
67:         attributes_dict = self._extract_attributes(attributes)
68:         polygons_list = self._extract_polygons(polygons, validity)
69:
70:         print(f"Extracted: {len(polygons_list)} valid polygons")
71:
72:         # Export intermediate results if requested
73:         if export_intermediate:
74:             self._export_intermediates(
75:                 mask_np, attributes_dict, polygons_list, Path(output_path).parent
76:             )
77:
78:         # Generate 3D model using deterministic pipeline
79:         success = self._generate_deterministic_3d(
80:             mask_np, attributes_dict, polygons_list, output_path
81:         )
82:
83:         return success
84:
85: def _load_image(self, image_path):
86:     """Load and preprocess input image"""
87:     image = cv2.imread(image_path)
88:     if image is None:
89:         raise ValueError(f"Could not load image from {image_path}")
90:

```

```

91:         image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
92:         image = cv2.resize(image, (256, 256))
93:         image = torch.from_numpy(image / 255.0).float()
94:         image = image.permute(2, 0, 1).unsqueeze(0)
95:         return image.to(self.device)
96:
97:     def _extract_mask(self, segmentation):
98:         """Convert soft segmentation to hard mask"""
99:         seg_pred = torch.argmax(segmentation, dim=1)
100:         mask_np = seg_pred.squeeze().cpu().numpy().astype(np.uint8)
101:         return mask_np
102:
103:     def _extract_attributes(self, attributes):
104:         """Convert normalized attributes back to physical values"""
105:         attr_np = attributes.squeeze().cpu().numpy()
106:
107:         # Denormalize (reverse of normalization in dataset)
108:         attributes_dict = {
109:             "wall_height": float(attr_np[0] * 5.0),
110:             "wall_thickness": float(attr_np[1] * 0.5),
111:             "window_base_height": float(attr_np[2] * 3.0),
112:             "window_height": float(attr_np[3] * 2.0),
113:             "door_height": float(attr_np[4] * 5.0),
114:             "pixel_scale": float(attr_np[5] * 0.02),
115:         }
116:
117:         return attributes_dict
118:
119:     def _extract_polygons(self, polygons, validity, threshold=None):
120:         """Extract valid polygons from network predictions"""
121:         threshold = threshold or self.config.polygon_threshold
122:         batch_size, num_polys, num_points, _ = polygons.shape
123:
124:         polygons_list = []
125:
126:         for poly_idx in range(num_polys):
127:             if validity[0, poly_idx] > threshold: # Only valid polygons
128:                 poly_points = polygons[0, poly_idx].cpu().numpy()
129:
130:                 # Remove zero-padded points
131:                 valid_points = poly_points[poly_points.sum(axis=1) > 0]
132:
133:                 if len(valid_points) >= 3: # Minimum for a polygon
134:                     # Convert to image coordinates (assuming 256x256 input)
135:                     valid_points = valid_points * 256
136:                     polygons_list.append(
137:                         {
138:                             "points": valid_points.tolist(),
139:                             "class": "wall", # Simplified - in practice classify polygon type
140:                         }
141:                     )
142:
143:         return polygons_list
144:
145:     def _export_intermediates(self, mask, attributes, polygons, output_dir):
146:         """Export intermediate results for debugging/analysis"""
147:         output_dir = Path(output_dir)
148:         output_dir.mkdir(exist_ok=True)
149:
150:         # Export mask
151:         cv2.imwrite(str(output_dir / "predicted_mask.png"), mask * 50)
152:
153:         # Export attributes
154:         with open(output_dir / "predicted_attributes.json", "w") as f:
155:             json.dump(attributes, f, indent=2)
156:
157:         # Export polygons
158:         with open(output_dir / "predicted_polygons.json", "w") as f:
159:             json.dump(polygons, f, indent=2)
160:
161:         # Visualize polygons on mask
162:         vis_img = np.zeros((256, 256, 3), dtype=np.uint8)
163:         vis_img[:, :, 0] = mask * 50 # Background

```

```

164:
165:     for poly in polygons:
166:         points = np.array(poly["points"], dtype=np.int32)
167:         cv2.polylines(vis_img, [points], True, (0, 255, 0), 2)
168:
169:     cv2.imwrite(str(output_dir / "polygon_visualization.png"), vis_img)
170:
171:     print(f"Intermediate results exported to {output_dir}")
172:
173: def _generate_deterministic_3d(self, mask, attributes, polygons, output_path):
174:     """Generate 3D model using deterministic geometric operations"""
175:     try:
176:         # Initialize mesh components
177:         vertices = []
178:         faces = []
179:         vertex_count = 0
180:
181:         # Extract geometric parameters
182:         wall_height = attributes.get("wall_height", 2.6)
183:         wall_thickness = attributes.get("wall_thickness", 0.15)
184:         pixel_scale = attributes.get("pixel_scale", 0.01)
185:
186:         print(
187:             f"Generating 3D model with wall_height={wall_height:.2f}m, thickness={wall_thickness:.2f}m"
188:         )
189:
190:         # Process each polygon (walls, rooms, etc.)
191:         for poly_idx, polygon in enumerate(polygons):
192:             poly_vertices, poly_faces = self._extrude_polygon_3d(
193:                 polygon["points"],
194:                 wall_height,
195:                 wall_thickness,
196:                 pixel_scale,
197:                 vertex_count,
198:             )
199:
200:             vertices.extend(poly_vertices)
201:             faces.extend(poly_faces)
202:             vertex_count += len(poly_vertices)
203:
204:         # Add floor and ceiling
205:         floor_verts, floor_faces = self._generate_floor_ceiling(
206:             mask, pixel_scale, wall_height, vertex_count
207:         )
208:         vertices.extend(floor_verts)
209:         faces.extend(floor_faces)
210:
211:         if len(vertices) == 0:
212:             print("No geometry generated")
213:             return False
214:
215:         # Create mesh
216:         mesh = trimesh.Trimesh(vertices=np.array(vertices), faces=np.array(faces))
217:
218:         # Clean up mesh
219:         mesh.remove_duplicate_faces()
220:         mesh.remove_unreferenced_vertices()
221:         mesh.fix_normals()
222:
223:         # Export
224:         mesh.export(output_path)
225:         print(f"3D model exported to {output_path}")
226:         print(
227:             f"Mesh statistics: {len(mesh.vertices)} vertices, {len(mesh.faces)} faces"
228:         )
229:
230:         return True
231:
232:     except Exception as e:
233:         print(f"Error generating 3D model: {str(e)}")
234:         return False
235:
236: def _extrude_polygon_3d(self, points, height, thickness, scale, vertex_offset):

```

```

237:     """Extrude a 2D polygon to create 3D wall geometry"""
238:     vertices = []
239:     faces = []
240:
241:     # Convert points to 3D coordinates
242:     points_3d = []
243:     for point in points:
244:         x = (point[0] - 128) * scale # Center and scale
245:         z = (128 - point[1]) * scale # Flip Y and scale
246:         points_3d.append([x, 0, z])
247:
248:     # Create bottom vertices (y=0)
249:     bottom_outer = points_3d
250:     bottom_inner = self._inset_polygon(points_3d, thickness)
251:
252:     # Create top vertices (y=height)
253:     top_outer = [[p[0], height, p[2]] for p in bottom_outer]
254:     top_inner = [[p[0], height, p[2]] for p in bottom_inner]
255:
256:     # Combine all vertices
257:     all_vertices = bottom_outer + bottom_inner + top_outer + top_inner
258:     vertices.extend(all_vertices)
259:
260:     n_points = len(points_3d)
261:
262:     # Generate faces for walls
263:     for i in range(n_points):
264:         next_i = (i + 1) % n_points
265:
266:         # Outer wall faces
267:         v1 = vertex_offset + i # bottom outer
268:         v2 = vertex_offset + next_i # bottom outer next
269:         v3 = vertex_offset + 2 * n_points + next_i # top outer next
270:         v4 = vertex_offset + 2 * n_points + i # top outer
271:
272:         faces.extend([[v1, v2, v3], [v1, v3, v4]])
273:
274:         # Inner wall faces (reverse winding)
275:         v1 = vertex_offset + n_points + i # bottom inner
276:         v2 = vertex_offset + n_points + next_i # bottom inner next
277:         v3 = vertex_offset + 3 * n_points + next_i # top inner next
278:         v4 = vertex_offset + 3 * n_points + i # top inner
279:
280:         faces.extend([[v1, v3, v2], [v1, v4, v3]])
281:
282:     # Top cap (between outer and inner)
283:     for i in range(n_points):
284:         next_i = (i + 1) % n_points
285:
286:         v1 = vertex_offset + 2 * n_points + i # top outer
287:         v2 = vertex_offset + 2 * n_points + next_i # top outer next
288:         v3 = vertex_offset + 3 * n_points + next_i # top inner next
289:         v4 = vertex_offset + 3 * n_points + i # top inner
290:
291:         faces.extend([[v1, v2, v3], [v1, v3, v4]])
292:
293:     # Bottom cap (between outer and inner)
294:     for i in range(n_points):
295:         next_i = (i + 1) % n_points
296:
297:         v1 = vertex_offset + i # bottom outer
298:         v2 = vertex_offset + next_i # bottom outer next
299:         v3 = vertex_offset + n_points + next_i # bottom inner next
300:         v4 = vertex_offset + n_points + i # bottom inner
301:
302:         faces.extend([[v1, v3, v2], [v1, v4, v3]])
303:
304:     return vertices, faces
305:
306: def _inset_polygon(self, points, inset_distance):
307:     """Create inset polygon for wall thickness"""
308:     if len(points) < 3:
309:         return points

```

```

310:
311:     # Simple inset by moving each point inward along angle bisector
312:     inset_points = []
313:     n = len(points)
314:
315:     for i in range(n):
316:         prev_i = (i - 1) % n
317:         next_i = (i + 1) % n
318:
319:         p_prev = np.array(points[prev_i])
320:         p_curr = np.array(points[i])
321:         p_next = np.array(points[next_i])
322:
323:         # Vectors to adjacent points
324:         v1 = p_curr - p_prev
325:         v2 = p_next - p_curr
326:
327:         # Normalize vectors (in XZ plane, ignore Y)
328:         v1_norm = np.array([v1[0], 0, v1[2]])
329:         v2_norm = np.array([v2[0], 0, v2[2]])
330:
331:         v1_len = np.linalg.norm(v1_norm)
332:         v2_len = np.linalg.norm(v2_norm)
333:
334:         if v1_len > 1e-6:
335:             v1_norm /= v1_len
336:         if v2_len > 1e-6:
337:             v2_norm /= v2_len
338:
339:         # Angle bisector
340:         bisector = v1_norm + v2_norm
341:         bisector_len = np.linalg.norm(bisector)
342:
343:         if bisector_len > 1e-6:
344:             bisector /= bisector_len
345:
346:         # Move point inward
347:         inset_point = p_curr - bisector * inset_distance
348:         inset_points.append([inset_point[0], inset_point[1], inset_point[2]])
349:     else:
350:         inset_points.append(points[i])
351:
352:     return inset_points
353:
354: def _generate_floor_ceiling(self, mask, scale, wall_height, vertex_offset):
355:     """Generate floor and ceiling geometry from segmentation mask"""
356:     vertices = []
357:     faces = []
358:
359:     # Find floor regions (assuming class 0 = floor/room)
360:     floor_mask = (mask == 0).astype(np.uint8)
361:
362:     # Find contours
363:     contours, _ = cv2.findContours(
364:         floor_mask, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE
365:     )
366:
367:     for contour in contours:
368:         if cv2.contourArea(contour) < 100: # Skip small regions
369:             continue
370:
371:         # Simplify contour
372:         epsilon = 0.02 * cv2.arcLength(contour, True)
373:         approx = cv2.approxPolyDP(contour, epsilon, True)
374:
375:         if len(approx) < 3:
376:             continue
377:
378:         # Convert to 3D coordinates
379:         floor_points = []
380:         for point in approx.reshape(-1, 2):
381:             x = (point[0] - 128) * scale
382:             z = (128 - point[1]) * scale

```

```

383:         floor_points.append([x, 0, z]) # Floor at y=0
384:
385:         ceiling_points = []
386:         for point in approx.reshape(-1, 2):
387:             x = (point[0] - 128) * scale
388:             z = (128 - point[1]) * scale
389:             ceiling_points.append([x, wall_height, z]) # Ceiling at y=wall_height
390:
391:         # Add vertices
392:         n_points = len(floor_points)
393:         vertices.extend(floor_points)
394:         vertices.extend(ceiling_points)
395:
396:         # Triangulate floor
397:         if n_points >= 3:
398:             for i in range(1, n_points - 1):
399:                 faces.append(
400:                     [vertex_offset, vertex_offset + i + 1, vertex_offset + i]
401:                 )
402:
403:         # Triangulate ceiling (reverse winding)
404:         for i in range(1, n_points - 1):
405:             faces.append(
406:                 [
407:                     vertex_offset + n_points,
408:                     vertex_offset + n_points + i,
409:                     vertex_offset + n_points + i + 1,
410:                 ]
411:             )
412:
413:         vertex_offset += 2 * n_points
414:
415:     return vertices, faces
416:
417: def process_batch(self, image_paths, output_dir):
418:     """Process multiple images in batch"""
419:     output_dir = Path(output_dir)
420:     output_dir.mkdir(exist_ok=True)
421:
422:     results = []
423:
424:     for img_path in image_paths:
425:         img_path = Path(img_path)
426:         print(f"Processing: {img_path.name}")
427:
428:         output_path = output_dir / f"{img_path.stem}_model.obj"
429:
430:         try:
431:             success = self.generate_3d_model(
432:                 str(img_path), str(output_path), export_intermediate=True
433:             )
434:
435:             results.append({
436:                 "input": str(img_path),
437:                 "output": str(output_path),
438:                 "success": success
439:             })
440:
441:             if success:
442:                 print(f"? Generated: {output_path}")
443:             else:
444:                 print(f"? Failed: {img_path.name}")
445:
446:         except Exception as e:
447:             print(f"? Error processing {img_path.name}: {str(e)}")
448:             results.append({
449:                 "input": str(img_path),
450:                 "output": str(output_path),
451:                 "success": False,
452:                 "error": str(e)
453:             })
454:
455:     return results

```

■ File: models\dvx.py

```
=====
1: """
2: Robust Differentiable Vectorization (DVX) module.
3:
4: Improvements vs naive DVX:
5: - Projects backbone feature maps to `feature_dim` if channels don't match via 1x1 conv.
6: - Multi-step iterative refinement (improves final polygon accuracy).
7: - Safe guards for shapes, device handling, and grid-sampling.
8: - Returns init_polygons, final polygons, per-step displacements, and validity scores.
9:
10: Usage:
11: - features: dict of feature maps (e.g. "p2", "p4"), each tensor (B, C, H, W).
12: - segmentation: (B, 1, H_img, W_img) or similar ? only used for optional initialization logic.
13: """
14:
15: from typing import Dict, Any, Optional, Tuple
16: import torch
17: import torch.nn as nn
18: import torch.nn.functional as F
19:
20:
21: class DifferentiableVectorization(nn.Module):
22:     def __init__(
23:         self,
24:         max_polygons: int = 20,
25:         max_points: int = 50,
26:         feature_dim: int = 256,
27:         displacement_scale: float = 0.12,
28:         num_refinement_steps: int = 3,
29:         align_corners: bool = False,
30:         padding_mode: str = "border", # options for grid_sample
31:         use_proj_conv: bool = True,
32:     ):
33:         """
34:         Args:
35:             max_polygons: maximum polygons to predict per image
36:             max_points: number of control points per polygon
37:             feature_dim: number of channels the DVX expects (will project backbone features to this)
38:             displacement_scale: multiplier for predicted displacement (tanh output)
39:             num_refinement_steps: how many iterative refinement steps to apply (>=1)
40:             align_corners: align_corners for F.grid_sample
41:             padding_mode: padding_mode for F.grid_sample
42:             use_proj_conv: whether to use 1x1 conv to project backbone features to feature_dim (recommended)
43:         """
44:         super().__init__()
45:         assert max_points > 2, "max_points must be > 2"
46:         assert num_refinement_steps >= 1
47:
48:         self.max_polygons = int(max_polygons)
49:         self.max_points = int(max_points)
50:         self.feature_dim = int(feature_dim)
51:         self.displacement_scale = float(displacement_scale)
52:         self.num_refinement_steps = int(num_refinement_steps)
53:         self.align_corners = bool(align_corners)
54:         self.padding_mode = padding_mode
55:         self.use_proj_conv = bool(use_proj_conv)
56:
57:         # init_net: from pooled p4 -> flattened -> produce normalized coords in [0,1]
58:         # AdaptiveAvgPool2d(8) -> (B, C, 8, 8) -> flatten -> Linear(C*8*8 -> hidden)
59:         hidden = max(512, feature_dim * 2)
60:         self.init_pool = nn.AdaptiveAvgPool2d(8)
61:
62:         # we'll create a projector conv for p4/p2 channels if necessary at runtime
63:         # but also create an MLP init_net that assumes feature_dim channels after pooling
64:         self.init_mlp = nn.Sequential(
65:             nn.Flatten(),
66:             nn.Linear(self.feature_dim * 8 * 8, hidden),
67:             nn.ReLU(inplace=True),
68:             nn.Linear(hidden, 1024),
69:             nn.ReLU(inplace=True),
70:             nn.Linear(1024, self.max_polygons * self.max_points * 2),
71:             nn.Sigmoid(),

```



```

72:         )
73:
74:         # refinement network: maps (feature_dim + 2) -> displacement in [-1,1]
75:         self.refine_net = nn.Sequential(
76:             nn.Linear(self.feature_dim + 2, 256),
77:             nn.ReLU(inplace=True),
78:             nn.Linear(256, 128),
79:             nn.ReLU(inplace=True),
80:             nn.Linear(128, 2),
81:             nn.Tanh(),
82:         )
83:
84:         # validity net (reads flattened coords only)
85:         self.validity_net = nn.Sequential(
86:             nn.Linear(self.max_points * 2, 128),
87:             nn.ReLU(inplace=True),
88:             nn.Linear(128, 1),
89:             nn.Sigmoid(),
90:         )
91:
92:         # projector convs (create lazily when first seen a feature channel mismatch)
93:         # stored per-key: e.g., self._proj_convs['p2'] = nn.Conv2d(in_ch, feature_dim, 1)
94:         self._proj_convs = nn.ModuleDict()
95:         self._proj_created = set()
96:
97:     def _ensure_projector(self, key: str, in_channels: int):
98:         """
99:         Ensure a 1x1 conv exists that projects `in_channels` -> self.feature_dim for feature map `key`.
100:         """
101:         if not self.use_proj_conv:
102:             return None
103:         if key in self._proj_created:
104:             return self._proj_convs[key]
105:
106:         if in_channels != self.feature_dim:
107:             conv = nn.Conv2d(in_channels, self.feature_dim, kernel_size=1, stride=1, padding=0)
108:             # initialize conv: kaiming
109:             nn.init.kaiming_normal_(conv.weight, a=0.2)
110:             if conv.bias is not None:
111:                 nn.init.zeros_(conv.bias)
112:             self._proj_convs[key] = conv
113:         else:
114:             # identity mapping using 1x1 conv with weights = identity-like is tricky
115:             # Instead simply keep no conv; we'll pass feature as-is
116:             self._proj_convs[key] = nn.Identity()
117:         self._proj_created.add(key)
118:         return self._proj_convs[key]
119:
120:     def _project_feature(self, key: str, feat: torch.Tensor) -> torch.Tensor:
121:         """
122:         Project or verify feature map to have self.feature_dim channels.
123:         If projector conv wasn't present and channels == feature_dim, returns feat unchanged.
124:         """
125:         in_ch = feat.shape[1]
126:         proj = self._ensure_projector(key, in_ch)
127:         if proj is None:
128:             # projection not desired; assert channels match
129:             if in_ch != self.feature_dim:
130:                 raise RuntimeError(
131:                     f"Feature '{key}' channels ({in_ch}) != feature_dim ({self.feature_dim}) "
132:                     "and projection disabled."
133:                 )
134:             return feat
135:         # if proj is Identity, apply it still (fast path)
136:         return proj(feat)
137:
138:     def forward(
139:         self,
140:         features: Dict[str, torch.Tensor],
141:         segmentation: Optional[torch.Tensor] = None,
142:         return_all_steps: bool = False,
143:     ) -> Dict[str, Any]:
144:         """

```

```

145:         features: dict with keys like "p2", "p4" containing tensors (B, C, H, W)
146:         segmentation: optional (B, 1, H_img, W_img) or similar (not strictly required)
147:         return_all_steps: if True returns per-step intermediate polygons & displacements
148:         """
149:         # pick features for init and refinement
150:         p4 = features.get("p4", None)
151:         p2 = features.get("p2", None)
152:
153:         if p4 is None and p2 is None:
154:             raise ValueError("At least one of 'p4' or 'p2' must be present in features.")
155:
156:         # prefer p4 for init; fallback to p2 if not present
157:         init_feat = p4 if p4 is not None else p2
158:         refine_feat = p2 if p2 is not None else p4
159:
160:         B = init_feat.shape[0]
161:
162:         # Project features to feature_dim (if needed)
163:         init_feat = self._project_feature("p4_init", init_feat)
164:         refine_feat = self._project_feature("p2_refine", refine_feat)
165:
166:         # -- Initialize polygons --
167:         # Pool then MLP; ensure init_mlp expects feature_dim channels
168:         pooled = self.init_pool(init_feat) # [B, C', 8, 8]
169:         if pooled.shape[1] != self.feature_dim:
170:             # If the projector returned Identity but pooled channels mismatch, try to apply a runtime project
171:             pooled = self._project_feature("p4_init_postpool", pooled)
172:
173:         init_logits = self.init_mlp(pooled) # [B, max_polygons * max_points * 2]
174:         init_polygons = init_logits.view(B, self.max_polygons, self.max_points, 2) # normalized [0,1]
175:
176:         # Iterative refinement
177:         polygons = init_polygons.clone()
178:         per_step_displacements = []
179:         for step in range(self.num_refinement_steps):
180:             # sample features at the polygon control-point locations
181:             displ = self._single_refine_step(polygons, refine_feat)
182:             per_step_displacements.append(displ)
183:             polygons = torch.clamp(polygons + displ, 0.0, 1.0)
184:
185:         # final validity
186:         validity = self._predict_validity(polygons)
187:
188:         out: Dict[str, Any] = {
189:             "polygons": polygons, # [B, P, N, 2]
190:             "validity": validity, # [B, P]
191:             "init_polygons": init_polygons,
192:             "refinement_displacements": per_step_displacements, # list of [B, P, N, 2]
193:         }
194:
195:         if return_all_steps:
196:             out["all_step_polygons"] = [
197:                 torch.clamp(init_polygons + sum(per_step_displacements[:i + 1]), 0.0, 1.0)
198:                 for i in range(len(per_step_displacements))
199:             ]
200:
201:         return out
202:
203:     def _single_refine_step(self, polygons: torch.Tensor, feature_map: torch.Tensor) -> torch.Tensor:
204:         """
205:         One refinement step: sample features at polygon points, predict displacement (scaled), return displ.
206:         polygons: [B, P, N, 2] in [0,1]
207:         feature_map: [B, C, H, W] with C == feature_dim (or projected)
208:         returns displacement: [B, P, N, 2] in [-displacement_scale, displacement_scale]
209:         """
210:         B, P, N, _ = polygons.shape
211:         # flatten pts to sample
212:         coords = polygons.view(B, -1, 2) # [B, P*N, 2], coords in [0,1]
213:         grid = coords * 2.0 - 1.0 # to [-1,1]
214:         # grid_sample expects (B, H_out, W_out, 2); use W_out=1
215:         grid_sample = grid.view(B, -1, 1, 2)
216:         sampled = F.grid_sample(
217:             feature_map,

```

```

218:         grid_sample,
219:         mode="bilinear",
220:         padding_mode=self.padding_mode,
221:         align_corners=self.align_corners,
222:     ) # [B, C, P*N, 1]
223:     sampled = sampled.squeeze(-1).permute(0, 2, 1).contiguous() # [B, P*N, C]
224:
225:     # combine sampled features and coords (coords in [0,1])
226:     input_feats = torch.cat([sampled, coords], dim=-1) # [B, P*N, C+2]
227:     # predict displacements in [-1,1] via tanh on last layer
228:     disp = self.refine_net(input_feats) # [B, P*N, 2], values ~[-1,1]
229:     disp = disp.view(B, P, N, 2)
230:     disp = disp * self.displacement_scale # scale
231:     return disp
232:
233: def _predict_validity(self, polygons: torch.Tensor) -> torch.Tensor:
234:     B, P, N, _ = polygons.shape
235:     if N != self.max_points:
236:         # If someone truncated or padded points, adapt: flatten to last dim whatever it is
237:         poly_flat = polygons.view(B * P, -1)
238:     else:
239:         poly_flat = polygons.view(B * P, -1)
240:     validity = self.validity_net(poly_flat) # [B*P, 1]
241:     validity = validity.view(B, P)
242:     return validity
243:
244:
245: # ----- quick unit test / smoke test -----
246: def _smoke_test():
247:     torch.manual_seed(0)
248:     B = 2
249:     C1 = 384 # different from feature_dim to test projector conv
250:     C2 = 128
251:     H2, W2 = 64, 64
252:     H4, W4 = 16, 16
253:
254:     # create dummy backbone features with different channels
255:     p2 = torch.randn(B, C1, H2, W2)
256:     p4 = torch.randn(B, C2, H4, W4)
257:     seg = torch.rand(B, 1, H2 * 4, W2 * 4) # just a placeholder
258:
259:     dvx = DifferentiableVectorization(
260:         max_polygons=4,
261:         max_points=16,
262:         feature_dim=256,
263:         displacement_scale=0.08,
264:         num_refinement_steps=3,
265:         align_corners=False,
266:         padding_mode="border",
267:         use_proj_conv=True,
268:     )
269:
270:     # ensure module moves projector convs to device when dvx.to(device) called
271:     dvx = dvx.eval() # inference mode ok
272:     # Forward pass
273:     out = dvx({"p2": p2, "p4": p4}, seg, return_all_steps=True)
274:     print("polygons shape:", out["polygons"].shape) # expected [B, P, N, 2]
275:     print("validity shape:", out["validity"].shape) # expected [B, P]
276:     print("init shape:", out["init_polygons"].shape)
277:     print("refinement steps:", len(out["refinement_displacements"]))
278:     # check ranges
279:     assert out["polygons"].min().item() >= 0.0 - 1e-6
280:     assert out["polygons"].max().item() <= 1.0 + 1e-6
281:     print("smoke test passed")
282:
283:
284: if __name__ == "__main__":
285:     _smoke_test()

```

■ File: models\encoder.py

```
=====
1: """
2: Encoder architecture for multi-scale feature extraction
3: """
4:
5: import torch
6: import torch.nn as nn
7: import torch.nn.functional as F
8:
9:
10: class ResidualBlock(nn.Module):
11:     """Basic residual block for the encoder"""
12:
13:     def __init__(self, in_channels, out_channels, stride=1):
14:         super().__init__()
15:
16:         self.conv1 = nn.Conv2d(in_channels, out_channels, 3, stride, 1, bias=False)
17:         self.bn1 = nn.BatchNorm2d(out_channels)
18:         self.conv2 = nn.Conv2d(out_channels, out_channels, 3, 1, 1, bias=False)
19:         self.bn2 = nn.BatchNorm2d(out_channels)
20:
21:         self.shortcut = nn.Sequential()
22:         if stride != 1 or in_channels != out_channels:
23:             self.shortcut = nn.Sequential(
24:                 nn.Conv2d(in_channels, out_channels, 1, stride, bias=False),
25:                 nn.BatchNorm2d(out_channels),
26:             )
27:
28:     def forward(self, x):
29:         out = F.relu(self.bn1(self.conv1(x)))
30:         out = self.bn2(self.conv2(out))
31:         out += self.shortcut(x)
32:         return F.relu(out)
33:
34:
35: class MultiScaleEncoder(nn.Module):
36:     """
37:     Advanced encoder with skip connections and multi-scale feature extraction
38:     Based on ResNet architecture with Feature Pyramid Network (FPN)
39:     """
40:
41:     def __init__(self, input_channels=3, feature_dim=512):
42:         super().__init__()
43:
44:         # Stem
45:         self.stem = nn.Sequential(
46:             nn.Conv2d(input_channels, 64, 7, 2, 3, bias=False),
47:             nn.BatchNorm2d(64),
48:             nn.ReLU(inplace=True),
49:             nn.MaxPool2d(3, 2, 1),
50:         )
51:
52:         # ResNet blocks
53:         self.layer1 = self._make_layer(64, 64, 2, stride=1) # 64x64
54:         self.layer2 = self._make_layer(64, 128, 2, stride=2) # 32x32
55:         self.layer3 = self._make_layer(128, 256, 2, stride=2) # 16x16
56:         self.layer4 = self._make_layer(256, 512, 2, stride=2) # 8x8
57:
58:         # FPN lateral connections
59:         self.lateral4 = nn.Conv2d(512, feature_dim, 1)
60:         self.lateral3 = nn.Conv2d(256, feature_dim, 1)
61:         self.lateral2 = nn.Conv2d(128, feature_dim, 1)
62:         self.lateral1 = nn.Conv2d(64, feature_dim, 1)
63:
64:         # FPN output layers
65:         self.smooth4 = nn.Conv2d(feature_dim, feature_dim, 3, 1, 1)
66:         self.smooth3 = nn.Conv2d(feature_dim, feature_dim, 3, 1, 1)
67:         self.smooth2 = nn.Conv2d(feature_dim, feature_dim, 3, 1, 1)
68:         self.smooth1 = nn.Conv2d(feature_dim, feature_dim, 3, 1, 1)
69:
70:         # Global context
71:         self.global_pool = nn.AdaptiveAvgPool2d(1)
```

```

72:         self.global_fc = nn.Sequential(
73:             nn.Linear(512, feature_dim),
74:             nn.ReLU(),
75:             nn.Linear(feature_dim, feature_dim)
76:         )
77:
78:     def _make_layer(self, in_channels, out_channels, blocks, stride=1):
79:         layers = []
80:         layers.append(ResidualBlock(in_channels, out_channels, stride))
81:         for _ in range(1, blocks):
82:             layers.append(ResidualBlock(out_channels, out_channels))
83:         return nn.Sequential(*layers)
84:
85:     def forward(self, x):
86:         # Bottom-up pathway
87:         x = self.stem(x) # 64x64
88:
89:         c1 = self.layer1(x) # 64x64
90:         c2 = self.layer2(c1) # 32x32
91:         c3 = self.layer3(c2) # 16x16
92:         c4 = self.layer4(c3) # 8x8
93:
94:         # Global context
95:         global_feat = self.global_pool(c4).flatten(1)
96:         global_feat = self.global_fc(global_feat)
97:
98:         # Top-down pathway (FPN)
99:         p4 = self.lateral4(c4)
100:        p3 = self.lateral3(c3) + F.interpolate(p4, scale_factor=2)
101:        p2 = self.lateral2(c2) + F.interpolate(p3, scale_factor=2)
102:        p1 = self.lateral1(c1) + F.interpolate(p2, scale_factor=2)
103:
104:        # Smooth
105:        p4 = self.smooth4(p4)
106:        p3 = self.smooth3(p3)
107:        p2 = self.smooth2(p2)
108:        p1 = self.smooth1(p1)
109:
110:        return {
111:            "p1": p1, # 64x64
112:            "p2": p2, # 32x32
113:            "p3": p3, # 16x16
114:            "p4": p4, # 8x8
115:            "global": global_feat,
116:        }

```

File: models\extrusion.py

```

=====
1: """
2: Vectorized Differentiable 3D extrusion module for converting polygons to 3D occupancy
3: Optimized version with GPU-accelerated vectorized operations
4: """
5:
6: import torch
7: import torch.nn as nn
8: import torch.nn.functional as F
9: import math
10: import logging
11:
12:
13: # -----
14: # Logging and sanitization helper
15: # -----
16: logger = logging.getLogger(__name__)
17: if not logger.handlers:
18:     handler = logging.StreamHandler()
19:     handler.setFormatter(logging.Formatter("%(asctime)s | %(levelname)s | %(message)s"))
20:     logger.addHandler(handler)
21: logger.setLevel(logging.INFO)
22:
23:

```

```

24: def _sanitize_normalized_height(value, sample_id=None, default=0.6):
25:     """
26:     Ensure normalized height value is finite and in [0,1].
27:     Returns a float in [0,1].
28:
29:     Args:
30:         value: torch scalar tensor or float
31:         sample_id: optional identifier for logging (string or int)
32:         default: fallback normalized height
33:     """
34:     try:
35:         if isinstance(value, torch.Tensor):
36:             raw = float(value.item())
37:         else:
38:             raw = float(value)
39:     except Exception:
40:         raw = float("nan")
41:
42:     # Build label for logging
43:     sid = f"[sample={sample_id}]" if sample_id is not None else ""
44:
45:     # Check finite
46:     if not math.isfinite(raw):
47:         logger.warning(f"{sid} Invalid wall height value (not finite): {raw}; using default {default}")
48:         raw = default
49:
50:     # Clamp to [0,1]
51:     if raw < 0.0 or raw > 1.0:
52:         logger.warning(f"{sid} Wall height normalized {raw} out of [0,1]; clamping.")
53:         raw = max(0.0, min(1.0, raw))
54:
55:     return raw
56:
57:
58: def _sanitize_tensor(tensor, default_value=0.0, name="tensor"):
59:     """
60:     Sanitize an entire tensor by replacing NaN/Inf values with default.
61:
62:     Args:
63:         tensor: Input tensor
64:         default_value: Value to replace invalid entries with
65:         name: Name for logging
66:
67:     Returns:
68:         Sanitized tensor
69:     """
70:     if tensor.numel() == 0:
71:         return tensor
72:
73:     # Check for any invalid values
74:     invalid_mask = ~torch.isfinite(tensor)
75:     num_invalid = invalid_mask.sum().item()
76:
77:     if num_invalid > 0:
78:         logger.warning(f"Found {num_invalid} invalid values in {name}, replacing with {default_value}")
79:         tensor = tensor.clone()
80:         tensor[invalid_mask] = default_value
81:
82:     return tensor
83:
84:
85: # -----
86: # Main extrusion module
87: # -----
88: class DifferentiableExtrusion(nn.Module):
89:     """
90:     Vectorized Differentiable 3D extrusion module
91:     Converts polygons + attributes to soft 3D occupancy grids
92:     """
93:
94:     def __init__(self, voxel_size: int = 64):
95:         super().__init__()
96:         self.voxel_size = int(voxel_size)

```

```

97:         self.register_buffer("_coords", None)
98:
99:     def _ensure_coords(self, device):
100:         """Initialize or update coordinate grid if needed"""
101:         if (self._coords is None or
102:             self._coords.device != device or
103:             self._coords.shape[0] != (self.voxel_size * self.voxel_size)):
104:
105:             H = W = self.voxel_size
106:             y, x = torch.meshgrid(
107:                 torch.arange(H, device=device),
108:                 torch.arange(W, device=device),
109:                 indexing="ij"
110:             )
111:             coords = torch.stack([x.flatten().float(), y.flatten().float()], dim=1) # [H*W, 2]
112:             coords = coords / float(self.voxel_size - 1)
113:             self.register_buffer("_coords", coords)
114:
115:     def polygon_sdf(self, polygon_xy):
116:         """
117:         Compute signed distance field for a polygon using vectorized operations.
118:         """
119:         device = polygon_xy.device
120:         self._ensure_coords(device)
121:         pts = self._coords # [M, 2]
122:         P = polygon_xy.shape[0]
123:
124:         if P < 2:
125:             return torch.full((pts.shape[0],), 1.0, device=device)
126:
127:         # Sanitize polygon coordinates
128:         polygon_xy = _sanitize_tensor(polygon_xy, default_value=0.0, name="polygon_xy")
129:
130:         v0 = polygon_xy.unsqueeze(1)
131:         v1 = torch.roll(polygon_xy, shifts=-1, dims=0).unsqueeze(1)
132:         pts_exp = pts.unsqueeze(0)
133:
134:         e = v1 - v0
135:         v = pts_exp - v0
136:         e_norm_sq = (e**2).sum(dim=2, keepdim=True) + 1e-8
137:         t = (v * e).sum(dim=2, keepdim=True) / e_norm_sq
138:         t_clamped = t.clamp(0.0, 1.0)
139:
140:         proj = v0 + t_clamped * e
141:         diff = pts_exp - proj
142:         dists = torch.norm(diff, dim=2)
143:
144:         # Sanitize distances before min operation
145:         dists = _sanitize_tensor(dists, default_value=1.0, name="distances")
146:         min_dist_per_point, _ = dists.min(dim=0)
147:
148:         x_pts = pts[:, 0].unsqueeze(0)
149:         y_pts = pts[:, 1].unsqueeze(0)
150:         x0, y0 = v0[..., 0], v0[..., 1]
151:         x1, y1 = v1[..., 0], v1[..., 1]
152:
153:         y_crosses = ((y0 <= y_pts) & (y1 > y_pts)) | ((y1 <= y_pts) & (y0 > y_pts))
154:         inter_x = x0 + (x1 - x0) * ((y_pts - y0) / (y1 - y0 + 1e-8))
155:         crossings = (inter_x > x_pts) & y_crosses
156:         crossing_count = crossings.sum(dim=0)
157:         inside = (crossing_count % 2 == 1)
158:
159:         sdf = min_dist_per_point.clone()
160:         sdf[inside] = -sdf[inside]
161:
162:         # Final sanitization of SDF output
163:         sdf = _sanitize_tensor(sdf, default_value=1.0, name="sdf")
164:         return sdf
165:
166:     def forward(self, polygons, attributes, validity_scores, sample_ids=None):
167:         """
168:         Convert polygons to 3D voxel occupancy.
169:         sample_ids: optional list/array of identifiers (e.g., filenames or dataset indices)

```

```

170:         """
171:         device = polygons.device
172:         B, N, P, _ = polygons.shape
173:         D = H = W = self.voxel_size
174:
175:         # Sanitize input tensors
176:         polygons = _sanitize_tensor(polygons, default_value=0.0, name="input_polygons")
177:         attributes = _sanitize_tensor(attributes, default_value=0.6, name="input_attributes")
178:         validity_scores = _sanitize_tensor(validity_scores, default_value=0.0, name="input_validity_scores")
179:
180:         voxels = torch.zeros((B, D, H, W), device=device)
181:
182:         for b in range(B):
183:             # pick identifier if available
184:             sid = sample_ids[b] if sample_ids is not None else b
185:
186:             # Sanitize height with logging
187:             wall_height_normalized = attributes[b, 0]
188:             sanitized_norm = _sanitize_normalized_height(
189:                 wall_height_normalized, sample_id=sid, default=0.6
190:             )
191:
192:             wall_height_m = sanitized_norm * 5.0
193:             height_frac = wall_height_m / 5.0
194:             height_voxels = int(round(height_frac * D))
195:             height_voxels = max(1, min(D, height_voxels))
196:
197:             # Process each polygon for this batch
198:             validity_mask = validity_scores[b] >= 0.5
199:             if not validity_mask.any():
200:                 continue
201:
202:             combined_mask = torch.zeros((H, W), device=device)
203:             sharpness = 100.0
204:
205:             for n in range(N):
206:                 if not validity_mask[n]:
207:                     continue
208:
209:                 polygon = polygons[b, n] # [P, 2]
210:
211:                 # Filter out zero-padded vertices
212:                 vertex_mask = (polygon.sum(dim=1) != 0.0)
213:                 if vertex_mask.sum().item() < 3:
214:                     continue
215:
216:                 valid_polygon = polygon[vertex_mask]
217:
218:                 # Compute SDF for this polygon
219:                 sdf = self.polygon_sdf(valid_polygon)
220:                 mask = torch.sigmoid(-sdf * sharpness)
221:                 mask_2d = mask.view(H, W)
222:
223:                 # Sanitize mask before combining
224:                 mask_2d = _sanitize_tensor(mask_2d, default_value=0.0, name=f"mask_2d_b{b}_n{n}")
225:                 combined_mask = torch.maximum(combined_mask, mask_2d)
226:
227:             # Create 3D mask by extruding to the computed height
228:             mask_3d = combined_mask.unsqueeze(0).expand(height_voxels, -1, -1)
229:
230:             # Sanitize final mask before assignment
231:             mask_3d = _sanitize_tensor(mask_3d, default_value=0.0, name=f"final_mask_3d_b{b}")
232:             voxels[b, :height_voxels] = mask_3d
233:
234:         # Final sanitization of output
235:         voxels = _sanitize_tensor(voxels, default_value=0.0, name="output_voxels")
236:         return voxels
237:
238:
239: # -----
240: # Fast extrusion module
241: # -----
242: class DifferentiableExtrusionFast(nn.Module):

```



```

243:     """
244:     Optimized version that batches polygon processing.
245:     """
246:
247:     def __init__(self, voxel_size: int = 64):
248:         super().__init__()
249:         self.voxel_size = int(voxel_size)
250:         self.register_buffer("_coords", None)
251:
252:     def _ensure_coords(self, device):
253:         if (self._coords is None or
254:             self._coords.device != device or
255:             self._coords.shape[0] != (self.voxel_size * self.voxel_size)):
256:
257:             H = W = self.voxel_size
258:             y, x = torch.meshgrid(
259:                 torch.arange(H, device=device),
260:                 torch.arange(W, device=device),
261:                 indexing="ij"
262:             )
263:             coords = torch.stack([x.flatten().float(), y.flatten().float()], dim=1)
264:             coords = coords / float(self.voxel_size - 1)
265:             self.register_buffer("_coords", coords)
266:
267:     def batch_polygon_sdf(self, polygons_batch, validity_mask):
268:         device = polygons_batch.device
269:         self._ensure_coords(device)
270:
271:         N, P, _ = polygons_batch.shape
272:         M = self._coords.shape[0]
273:         sdfs = torch.full((N, M), 1.0, device=device)
274:
275:         valid_indices = torch.where(validity_mask)[0]
276:         if len(valid_indices) == 0:
277:             return sdfs
278:
279:         valid_polygons = polygons_batch[valid_indices]
280:         for i, poly_idx in enumerate(valid_indices):
281:             poly = valid_polygons[i]
282:             vertex_mask = (poly.sum(dim=1) != 0.0)
283:             if vertex_mask.sum().item() >= 3:
284:                 valid_poly = poly[vertex_mask]
285:                 sdf = self.polygon_sdf(valid_poly)
286:                 sdfs[poly_idx] = sdf
287:
288:         return sdfs
289:
290:     def polygon_sdf(self, polygon_xy):
291:         device = polygon_xy.device
292:         self._ensure_coords(device)
293:         pts = self._coords
294:         P = polygon_xy.shape[0]
295:
296:         if P < 2:
297:             return torch.full((pts.shape[0],), 1.0, device=device)
298:
299:         # Sanitize polygon coordinates
300:         polygon_xy = _sanitize_tensor(polygon_xy, default_value=0.0, name="polygon_xy_fast")
301:
302:         v0 = polygon_xy.unsqueeze(1)
303:         v1 = torch.roll(polygon_xy, shifts=-1, dims=0).unsqueeze(1)
304:         pts_exp = pts.unsqueeze(0)
305:
306:         e = v1 - v0
307:         v = pts_exp - v0
308:         e_norm_sq = (e**2).sum(dim=2, keepdim=True) + 1e-8
309:         t = (v * e).sum(dim=2, keepdim=True) / e_norm_sq
310:         t_clamped = t.clamp(0.0, 1.0)
311:
312:         proj = v0 + t_clamped * e
313:         diff = pts_exp - proj
314:         dists = torch.norm(diff, dim=2)
315:

```

```

316:         # Sanitize distances before min operation
317:         dists = _sanitize_tensor(dists, default_value=1.0, name="distances_fast")
318:         min_dist_per_point, _ = dists.min(dim=0)
319:
320:         x_pts = pts[:, 0].unsqueeze(0)
321:         y_pts = pts[:, 1].unsqueeze(0)
322:         x0, y0 = v0[..., 0], v0[..., 1]
323:         x1, y1 = v1[..., 0], v1[..., 1]
324:
325:         y_crosses = ((y0 <= y_pts) & (y1 > y_pts)) | ((y1 <= y_pts) & (y0 > y_pts))
326:         inter_x = x0 + (x1 - x0) * ((y_pts - y0) / (y1 - y0 + 1e-8))
327:         crossings = (inter_x > x_pts) & y_crosses
328:         crossing_count = crossings.sum(dim=0)
329:         inside = (crossing_count % 2 == 1)
330:
331:         sdf = min_dist_per_point.clone()
332:         sdf[inside] = -sdf[inside]
333:
334:         # Final sanitization of SDF output
335:         sdf = _sanitize_tensor(sdf, default_value=1.0, name="sdf_fast")
336:         return sdf
337:
338: def forward(self, polygons: torch.Tensor, attributes: torch.Tensor, validity_scores: torch.Tensor) -> torch.Tensor:
339:     device = polygons.device
340:     B, N, P, _ = polygons.shape
341:     D = H = W = self.voxel_size
342:
343:     # Sanitize input tensors
344:     polygons = _sanitize_tensor(polygons, default_value=0.0, name="input_polygons_fast")
345:     attributes = _sanitize_tensor(attributes, default_value=0.6, name="input_attributes_fast")
346:     validity_scores = _sanitize_tensor(validity_scores, default_value=0.0, name="input_validity_scores_fast")
347:
348:     voxels = torch.zeros((B, D, H, W), device=device)
349:
350:     for b in range(B):
351:         validity_mask = validity_scores[b] >= 0.5
352:         if not validity_mask.any():
353:             continue
354:
355:         sdfs = self.batch_polygon_sdf(polygons[b], validity_mask)
356:
357:         # Sanitize SDFs before sigmoid
358:         sdfs = _sanitize_tensor(sdfs, default_value=1.0, name=f"batch_sdfs_b{b}")
359:
360:         sharpness = 100.0
361:         masks = torch.sigmoid(-sdfs * sharpness)
362:         masks_2d = masks.view(N, H, W)
363:
364:         # Sanitize masks
365:         masks_2d = _sanitize_tensor(masks_2d, default_value=0.0, name=f"masks_2d_b{b}")
366:
367:         # Sanitize height
368:         wall_height_normalized = attributes[b, 0]
369:         sanitized_norm = _sanitize_normalized_height(wall_height_normalized, sample_id=b, default=0.6)
370:         wall_height_m = sanitized_norm * 5.0
371:         height_frac = wall_height_m / 5.0
372:         height_voxels = int(round(height_frac * D))
373:         height_voxels = max(1, min(D, height_voxels))
374:
375:         combined_mask = torch.zeros((H, W), device=device)
376:         for n in range(N):
377:             if validity_mask[n]:
378:                 combined_mask = torch.maximum(combined_mask, masks_2d[n])
379:
380:         mask_3d = combined_mask.unsqueeze(0).expand(height_voxels, -1, -1)
381:
382:         # Sanitize final mask before assignment
383:         mask_3d = _sanitize_tensor(mask_3d, default_value=0.0, name=f"final_mask_3d_fast_b{b}")
384:         voxels[b, :height_voxels] = mask_3d
385:
386:     # Final sanitization of output
387:     voxels = _sanitize_tensor(voxels, default_value=0.0, name="output_voxels_fast")
388:     return voxels

```

■ File: models\heads.py

```
=====
1: """
2: Multi-task prediction heads for the Neural-Geometric 3D Model Generator
3: """
4:
5: import torch
6: import torch.nn as nn
7: import torch.nn.functional as F
8:
9:
10: class SegmentationHead(nn.Module):
11:     """Semantic segmentation head with multi-scale fusion"""
12:
13:     def __init__(self, feature_dim=512, num_classes=5, dropout=0.1):
14:         super().__init__()
15:
16:         # Multi-scale fusion
17:         self.fusion = nn.Sequential(
18:             nn.Conv2d(feature_dim * 4, feature_dim, 3, 1, 1),
19:             nn.BatchNorm2d(feature_dim),
20:             nn.ReLU(),
21:             nn.Dropout2d(dropout),
22:         )
23:
24:         # Segmentation decoder
25:         self.decoder = nn.Sequential(
26:             nn.Conv2d(feature_dim, feature_dim // 2, 3, 1, 1),
27:             nn.BatchNorm2d(feature_dim // 2),
28:             nn.ReLU(),
29:             nn.Conv2d(feature_dim // 2, feature_dim // 4, 3, 1, 1),
30:             nn.BatchNorm2d(feature_dim // 4),
31:             nn.ReLU(),
32:             nn.Conv2d(feature_dim // 4, num_classes, 1),
33:         )
34:
35:     def forward(self, features):
36:         # Fuse multi-scale features
37:         p1, p2, p3, p4 = features["p1"], features["p2"], features["p3"], features["p4"]
38:
39:         # Upsample all to p1 resolution
40:         p2_up = F.interpolate(
41:             p2, size=p1.shape[-2:], mode="bilinear", align_corners=False
42:         )
43:         p3_up = F.interpolate(
44:             p3, size=p1.shape[-2:], mode="bilinear", align_corners=False
45:         )
46:         p4_up = F.interpolate(
47:             p4, size=p1.shape[-2:], mode="bilinear", align_corners=False
48:         )
49:
50:         fused = torch.cat([p1, p2_up, p3_up, p4_up], dim=1)
51:         fused = self.fusion(fused)
52:
53:         # Final segmentation
54:         seg = self.decoder(fused)
55:         return F.interpolate(seg, scale_factor=4, mode="bilinear", align_corners=False)
56:
57:
58: class AttributeHead(nn.Module):
59:     """Attribute regression head for geometric parameters"""
60:
61:     def __init__(self, feature_dim=512, num_attributes=6, dropout=0.2):
62:         super().__init__()
63:
64:         self.regressor = nn.Sequential(
65:             nn.Linear(feature_dim, feature_dim),
66:             nn.ReLU(),
67:             nn.Dropout(dropout),
68:             nn.Linear(feature_dim, feature_dim // 2),
69:             nn.ReLU(),
70:             nn.Dropout(dropout),
71:             nn.Linear(feature_dim // 2, num_attributes),
```

```

72:         nn.Sigmoid(), # Output in [0,1] range
73:     )
74:
75:     def forward(self, global_features):
76:         return self.regressor(global_features)
77:
78:
79: class SDFHead(nn.Module):
80:     """Signed Distance Field prediction for sharp boundaries"""
81:
82:     def __init__(self, feature_dim=512, dropout=0.1):
83:         super().__init__()
84:
85:         self.sdf_decoder = nn.Sequential(
86:             nn.Conv2d(feature_dim, feature_dim // 2, 3, 1, 1),
87:             nn.BatchNorm2d(feature_dim // 2),
88:             nn.ReLU(),
89:             nn.Dropout2d(dropout),
90:             nn.Conv2d(feature_dim // 2, feature_dim // 4, 3, 1, 1),
91:             nn.BatchNorm2d(feature_dim // 4),
92:             nn.ReLU(),
93:             nn.Conv2d(feature_dim // 4, 1, 1),
94:             nn.Tanh(), # SDF in [-1, 1]
95:         )
96:
97:     def forward(self, features):
98:         # Use highest resolution features
99:         p1 = features["p1"]
100:         sdf = self.sdf_decoder(p1)
101:         return F.interpolate(sdf, scale_factor=4, mode="bilinear", align_corners=False)

```

File: models\model.py

```

=====
1: """
2: Advanced loss functions for multi-task training with dynamic weighting
3: Enhanced with cross-modal consistency, graph constraints, and GradNorm
4: Modified to support conditional geometric losses via run_full_geometric flag
5: """
6: import torch
7: import torch.nn as nn
8: import torch.nn.functional as F
9: from .encoder import MultiScaleEncoder
10: from .heads import SegmentationHead, AttributeHead, SDFHead
11: from .dvx import DifferentiableVectorization
12: from .extrusion import DifferentiableExtrusion
13:
14:
15: class L2Normalize(nn.Module):
16:     """L2 normalization layer"""
17:
18:     def __init__(self, dim=1):
19:         super().__init__()
20:         self.dim = dim
21:
22:     def forward(self, x):
23:         return F.normalize(x, p=2, dim=self.dim)
24:
25:
26: class LatentEmbeddingHead(nn.Module):
27:     """Auxiliary head for cross-modal latent consistency"""
28:
29:     def __init__(self, feature_dim: int, embedding_dim: int = 256):
30:         super().__init__()
31:         self.embedding_dim = embedding_dim
32:
33:         # 2D embedding path
34:         self.embedding_2d = nn.Sequential(
35:             nn.AdaptiveAvgPool2d((1, 1)),
36:             nn.Flatten(),
37:             nn.Linear(feature_dim, embedding_dim * 2),
38:             nn.ReLU(),

```

```

39:         nn.Dropout(0.1),
40:         nn.Linear(embedding_dim * 2, embedding_dim),
41:         L2Normalize(dim=1), # L2 normalize for cosine similarity
42:     )
43:
44:     # 3D embedding path (from voxel features)
45:     self.embedding_3d = nn.Sequential(
46:         nn.AdaptiveAvgPool3d((1, 1, 1)),
47:         nn.Flatten(),
48:         nn.Linear(feature_dim, embedding_dim * 2),
49:         nn.ReLU(),
50:         nn.Dropout(0.1),
51:         nn.Linear(embedding_dim * 2, embedding_dim),
52:         L2Normalize(dim=1),
53:     )
54:
55:     def forward(
56:         self, features_2d: torch.Tensor, features_3d: torch.Tensor = None
57:     ) -> tuple:
58:         """
59:         Generate 2D and 3D embeddings for consistency loss
60:
61:         Args:
62:             features_2d: [B, C, H, W] - 2D feature maps
63:             features_3d: [B, C, D, H, W] - 3D feature maps (optional)
64:
65:         Returns:
66:             tuple: (embedding_2d, embedding_3d)
67:         """
68:         # 2D embedding
69:         emb_2d = self.embedding_2d(features_2d)
70:
71:         # 3D embedding (if available, otherwise use 2D features reshaped)
72:         if features_3d is not None:
73:             emb_3d = self.embedding_3d(features_3d)
74:         else:
75:             # Create pseudo-3D from 2D features
76:             B, C, H, W = features_2d.shape
77:             pseudo_3d = features_2d.unsqueeze(2).expand(
78:                 B, C, 4, H, W
79:             ) # Duplicate across depth
80:             emb_3d = self.embedding_3d(pseudo_3d)
81:
82:         return emb_2d, emb_3d
83:
84:
85: class GraphStructureHead(nn.Module):
86:     """Head for predicting graph structure (room connectivity)"""
87:
88:     def __init__(self, feature_dim: int, max_rooms: int = 10):
89:         super().__init__()
90:         self.max_rooms = max_rooms
91:
92:         # Room detection branch
93:         self.room_detector = nn.Sequential(
94:             nn.Conv2d(feature_dim, feature_dim // 2, 3, padding=1),
95:             nn.ReLU(),
96:             nn.Conv2d(feature_dim // 2, max_rooms, 3, padding=1),
97:             nn.Sigmoid(), # Room probability maps
98:         )
99:
100:        # Room feature extractor
101:        self.room_features = nn.Sequential(
102:            nn.AdaptiveAvgPool2d((8, 8)), # Pool to fixed size
103:            nn.Flatten(),
104:            nn.Linear(feature_dim * 64, 256),
105:            nn.ReLU(),
106:            nn.Linear(256, 128), # Room feature vectors
107:        )
108:
109:        # Adjacency predictor
110:        self.adjacency_net = nn.Sequential(
111:            nn.Linear(128 * 2, 64), # Pairwise room features

```

```

112:         nn.ReLU(),
113:         nn.Linear(64, 32),
114:         nn.ReLU(),
115:         nn.Linear(32, 1),
116:         nn.Sigmoid(), # Adjacency probability
117:     )
118:
119: def forward(self, features: torch.Tensor) -> dict:
120:     """
121:     Predict room graph structure
122:
123:     Args:
124:         features: [B, C, H, W] - Feature maps
125:
126:     Returns:
127:         dict with 'room_maps', 'room_features', 'adjacency_matrix'
128:     """
129:     B = features.shape[0]
130:
131:     # Detect room probability maps
132:     room_maps = self.room_detector(features) # [B, max_rooms, H, W]
133:
134:     # Extract room features
135:     room_feats = self.room_features(features) # [B, 128]
136:
137:     # Create adjacency matrix for all room pairs
138:     adjacency_matrices = []
139:
140:     for b in range(B):
141:         # Get room features for this batch item
142:         feat_b = room_feats[b : b + 1] # [1, 128]
143:
144:         # Create pairwise combinations
145:         adj_matrix = torch.zeros(
146:             (self.max_rooms, self.max_rooms), device=features.device
147:         )
148:
149:         for i in range(self.max_rooms):
150:             for j in range(i + 1, self.max_rooms):
151:                 # Concatenate features for room pair
152:                 pair_feat = torch.cat([feat_b, feat_b], dim=1) # [1, 256]
153:
154:                 # Predict adjacency
155:                 adj_prob = self.adjacency_net(pair_feat) # [1, 1]
156:
157:                 # Fill symmetric matrix
158:                 adj_matrix[i, j] = adj_prob.squeeze()
159:                 adj_matrix[j, i] = adj_prob.squeeze()
160:
161:             adjacency_matrices.append(adj_matrix)
162:
163:     return {
164:         "room_maps": room_maps,
165:         "room_features": room_feats,
166:         "adjacency_matrices": torch.stack(adjacency_matrices),
167:     }
168:
169:
170: class NeuralGeometric3DGenerator(nn.Module):
171:     """
172:     Enhanced neural-geometric system with auxiliary heads for novel training strategies:
173:     - Cross-modal latent consistency
174:     - Graph structure prediction
175:     - Multi-view embeddings for dynamic curriculum
176:     - Conditional geometric computation via run_full_geometric flag
177:     """
178:
179:     def __init__(
180:         self,
181:         input_channels=3,
182:         num_classes=5,
183:         feature_dim=512,
184:         num_attributes=6,

```

```

185:         voxel_size=64,
186:         max_polygons=20,
187:         max_points=50,
188:         use_latent_consistency=True,
189:         use_graph_constraints=True,
190:         latent_embedding_dim=256,
191:     ):
192:         super().__init__()
193:
194:         # Store configuration
195:         self.use_latent_consistency = use_latent_consistency
196:         self.use_graph_constraints = use_graph_constraints
197:         self.feature_dim = feature_dim
198:
199:         # Core components
200:         self.encoder = MultiScaleEncoder(input_channels, feature_dim)
201:         self.seg_head = SegmentationHead(feature_dim, num_classes)
202:         self.attr_head = AttributeHead(feature_dim, num_attributes)
203:         self.sdf_head = SDFHead(feature_dim)
204:         self.dvx = DifferentiableVectorization(max_polygons, max_points, feature_dim)
205:         self.extrusion = DifferentiableExtrusion(voxel_size)
206:
207:         # NEW: Auxiliary heads for novel training strategies
208:         if use_latent_consistency:
209:             self.latent_head = LatentEmbeddingHead(feature_dim, latent_embedding_dim)
210:
211:         if use_graph_constraints:
212:             self.graph_head = GraphStructureHead(feature_dim)
213:
214:         # Enhanced feature processing for multi-stage training
215:         self.feature_enhancer = nn.Sequential(
216:             nn.Conv2d(feature_dim, feature_dim, 3, padding=1),
217:             nn.GroupNorm(32, feature_dim),
218:             nn.ReLU(),
219:             nn.Conv2d(feature_dim, feature_dim, 3, padding=1),
220:             nn.GroupNorm(32, feature_dim),
221:         )
222:
223:         # lazy-created 3d voxel processor will be attached on first use
224:         self._voxel_processor = None
225:
226:     def _select_spatial_feature(self, features):
227:         """
228:         Given encoder output (dict or tensor), select a spatial 4-D feature map
229:         Prefer high-resolution feature maps (p1) and avoid selecting 'global' vector.
230:         """
231:         # If encoder returned a tensor already, make sure it's 4D
232:         if not isinstance(features, dict):
233:             if features.dim() == 4:
234:                 return features
235:             else:
236:                 raise ValueError(
237:                     f"Encoder returned a tensor with shape {tuple(features.shape)}; "
238:                     "expected a 4D feature map [B, C, H, W].")
239:
240:         # Encoder returned dict: prefer p1,p2,p3,p4,high_res,out,main but NEVER 'global'
241:         preferred_keys = ["p1", "p2", "p3", "p4", "high_res", "out", "main"]
242:         for k in preferred_keys:
243:             if k in features:
244:                 candidate = features[k]
245:                 if isinstance(candidate, torch.Tensor) and candidate.dim() == 4:
246:                     return candidate
247:
248:         # As a last resort, scan dict values for the first 4D tensor that isn't 'global'
249:         for k, v in features.items():
250:             if k == "global":
251:                 continue
252:             if isinstance(v, torch.Tensor) and v.dim() == 4:
253:                 return v
254:
255:         # If nothing found, raise informative error rather than silently picking wrong shape
256:         raise RuntimeError(

```

```

258:         "No spatial 4D feature map found in encoder output. Encoder returned keys: "
259:         f"{list(features.keys())}. Ensure encoder provides at least one [B,C,H,W] tensor "
260:         "under keys like 'p1','p2','p3','p4','out', or 'high_res'."
261:     )
262:
263: def forward(self, image, run_full_geometric=True, return_aux=True):
264:     """
265:     Enhanced forward pass with auxiliary outputs and conditional geometric computation
266:
267:     Args:
268:         image: [B, C, H, W] input images
269:         run_full_geometric: Whether to run heavy DVX and extrusion computations
270:         return_aux: Whether to compute auxiliary outputs
271:
272:     Returns:
273:         dict with predictions, conditionally including geometric outputs
274:     """
275:     # Multi-scale feature extraction
276:     features = self.encoder(image)
277:
278:     # Enhance features
279:     spatial_feat = self._select_spatial_feature(features)
280:     enhanced_features = self.feature_enhancer(spatial_feat)
281:
282:     # keep structured features dict for heads that expect multi-scale inputs
283:     if isinstance(features, dict):
284:         features["enhanced"] = enhanced_features
285:         main_features = enhanced_features
286:     else:
287:         features = {"main": enhanced_features, "enhanced": enhanced_features}
288:         main_features = enhanced_features
289:
290:     # Core predictions (always computed - these are fast)
291:     segmentation = self.seg_head(features)
292:     attributes = self.attr_head(
293:         features.get("global")
294:         if isinstance(features, dict) and "global" in features
295:         else main_features.mean(dim=[2, 3])
296:     )
297:     sdf = self.sdf_head(features)
298:
299:     # Base outputs
300:     outputs = {
301:         "segmentation": segmentation,
302:         "attributes": attributes,
303:         "sdf": sdf,
304:         "features": features,
305:     }
306:
307:     # Conditional geometric computation (heavy operations)
308:     if run_full_geometric:
309:         # DVX polygon fitting
310:         dvx_output = self.dvx(features, segmentation)
311:         polygons = dvx_output.get("polygons", None)
312:         validity = dvx_output.get("validity", None)
313:
314:         # 3D extrusion (defensive: ensure inputs exist)
315:         try:
316:             voxels_pred = self.extrusion(polygons, attributes, validity)
317:         except Exception as e:
318:             # Log or print a helpful message for debugging; avoid crashing training
319:             # (Replace print with logger if you have one)
320:             print(f"[Warning] extrusion failed: {e}")
321:             voxels_pred = None
322:
323:         # Add geometric outputs
324:         outputs.update({
325:             "polygons": polygons,
326:             "polygon_validity": validity,
327:             "voxels_pred": voxels_pred,
328:         })
329:
330:     # NEW: Auxiliary outputs for novel training strategies (only when geometric is enabled)

```



```

331:         if return_aux:
332:             # Cross-modal consistency embeddings
333:             if self.use_latent_consistency:
334:                 if voxels_pred is not None:
335:                     voxel_features = self._create_3d_features_from_voxels(voxels_pred)
336:                     latent_2d, latent_3d = self.latent_head(main_features, voxel_features)
337:                 else:
338:                     # Fall back to pseudo-3D built from 2D features if voxels not available
339:                     latent_2d, latent_3d = self.latent_head(main_features, None)
340:                     outputs["latent_2d_embedding"] = latent_2d
341:                     outputs["latent_3d_embedding"] = latent_3d
342:             else:
343:                 # Geometric path explicitly skipped for this stage
344:                 outputs.update({
345:                     "polygons": None,
346:                     "polygon_validity": None,
347:                     "voxels_pred": None,
348:                 })
349:
350:             # Still compute some auxiliary outputs that don't depend on geometry
351:             if return_aux and self.use_latent_consistency:
352:                 # Use pseudo-3D features for 2D-only consistency inside latent head
353:                 latent_2d, latent_3d = self.latent_head(main_features, None)
354:                 outputs["latent_2d_embedding"] = latent_2d
355:                 outputs["latent_3d_embedding"] = latent_3d
356:
357:             # Graph structure predictions (independent of geometric computation)
358:             if return_aux and self.use_graph_constraints:
359:                 graph_output = self.graph_head(main_features)
360:                 outputs.update(graph_output)
361:
362:         return outputs
363:
364:     def get_latent_embeddings(self, image):
365:         """
366:         Convenience method to get just the latent embeddings
367:         Used by trainer for consistency loss
368:         """
369:         if not self.use_latent_consistency:
370:             return None, None
371:
372:         with torch.no_grad():
373:             features = self.encoder(image)
374:             spatial_feat = self._select_spatial_feature(features)
375:             main_features = self.feature_enhancer(spatial_feat)
376:
377:             # Quick forward to get segmentation/attributes
378:             segmentation = self.seg_head(features)
379:             attributes = self.attr_head(
380:                 features.get("global")
381:                 if isinstance(features, dict) and "global" in features
382:                 else main_features.mean(dim=[2, 3])
383:             )
384:
385:             # Attempt DVX + extrusion, but be defensive (may be expensive)
386:             dvx_output = self.dvx(features, segmentation)
387:             polygons = dvx_output.get("polygons", None)
388:             validity = dvx_output.get("validity", None)
389:
390:             try:
391:                 voxels_pred = self.extrusion(polygons, attributes, validity)
392:             except Exception as e:
393:                 print(f"[Warning] get_latent_embeddings: extrusion failed: {e}")
394:                 voxels_pred = None
395:
396:             # If voxels not available, latent_head will fall back to pseudo-3D
397:             if voxels_pred is not None:
398:                 voxel_features = self._create_3d_features_from_voxels(voxels_pred)
399:             else:
400:                 voxel_features = None
401:
402:         return self.latent_head(main_features, voxel_features)
403:

```

```

404:     def _create_3d_features_from_voxels(self, voxels):
405:         """
406:         Create 3D feature representation from voxel predictions
407:
408:         Args:
409:             voxels: [B, D, H, W] voxel predictions
410:
411:         Returns:
412:             [B, C, D, H, W] 3D features
413:         """
414:         # Defensive check
415:         if voxels is None:
416:             raise ValueError(
417:                 "Received voxels=None in _create_3d_features_from_voxels(). "
418:                 "This indicates that the geometric pipeline was skipped or extrusion failed. "
419:                 "Call this method only when voxels are available, or use latent_head(..., None) to "
420:                 "compute pseudo-3D features from 2D."
421:             )
422:
423:         # Ensure expected shape
424:         if voxels.dim() != 4:
425:             raise ValueError(f"voxels must have shape [B,D,H,W], got {tuple(voxels.shape)}")
426:
427:         B, D, H, W = voxels.shape
428:
429:         # Expand voxels to have feature channels
430:         # Simple approach: repeat voxel values across feature dimension
431:         rep_ch = max(1, self.feature_dim // 4)
432:         voxel_features = voxels.unsqueeze(1).expand(B, rep_ch, D, H, W).contiguous()
433:
434:         # Add some learned 3D processing
435:         if self._voxel_processor is None:
436:             # Build with correct device
437:             device = voxels.device
438:             self._voxel_processor = nn.Sequential(
439:                 nn.Conv3d(rep_ch, max(rep_ch, self.feature_dim // 2), 3, padding=1),
440:                 nn.ReLU(),
441:                 nn.Conv3d(max(rep_ch, self.feature_dim // 2), self.feature_dim, 3, padding=1),
442:             ).to(device)
443:
444:         return self._voxel_processor(voxel_features)
445:
446:     def get_stage_parameters(self, stage: int):
447:         """
448:         Get parameters for specific training stage
449:         Useful for stage-specific optimization
450:         """
451:         if stage == 1:
452:             # Stage 1: 2D components only
453:             params = []
454:             params.extend(list(self.encoder.parameters()))
455:             params.extend(list(self.seg_head.parameters()))
456:             params.extend(list(self.attr_head.parameters()))
457:             params.extend(list(self.sdf_head.parameters()))
458:             params.extend(list(self.feature_enhancer.parameters()))
459:
460:             if self.use_latent_consistency:
461:                 params.extend(list(self.latent_head.parameters()))
462:
463:         elif stage == 2:
464:             # Stage 2: DVX components
465:             params = list(self.dvx.parameters())
466:
467:         else: # stage == 3
468:             # Stage 3: All parameters
469:             params = list(self.parameters())
470:
471:         return params
472:
473:     def freeze_stage_parameters(self, stages_to_freeze: list):
474:         """
475:         Freeze parameters for specific stages
476:

```

```

477:         Args:
478:             stages_to_freeze: List of stage numbers to freeze
479:             """
480:         for stage in stages_to_freeze:
481:             stage_params = self.get_stage_parameters(stage)
482:             for param in stage_params:
483:                 param.requires_grad = False
484:
485:     def unfreeze_stage_parameters(self, stages_to_unfreeze: list):
486:         """
487:         Unfreeze parameters for specific stages
488:
489:         Args:
490:             stages_to_unfreeze: List of stage numbers to unfreeze
491:             """
492:         for stage in stages_to_unfreeze:
493:             stage_params = self.get_stage_parameters(stage)
494:             for param in stage_params:
495:                 param.requires_grad = True
496:
497:     def get_curriculum_metrics(self):
498:         """
499:         Get metrics useful for curriculum learning decisions
500:         """
501:         metrics = {}
502:
503:         # Parameter counts per stage
504:         for stage in [1, 2, 3]:
505:             stage_params = self.get_stage_parameters(stage)
506:             metrics[f"stage_{stage}_params"] = sum(p.numel() for p in stage_params)
507:
508:         # Feature dimensions
509:         metrics["feature_dim"] = self.feature_dim
510:         metrics["has_latent_consistency"] = self.use_latent_consistency
511:         metrics["has_graph_constraints"] = self.use_graph_constraints
512:
513:         return metrics

```

File: param.py

```

=====
1: import json
2: import numpy as np
3: from pathlib import Path
4:
5: # Adjust this to your dataset path
6: data_root = Path("./data/floorplans")
7:
8: # Expected attributes with their default values
9: expected_keys = {
10:     "wall_height": 2.6,
11:     "wall_thickness": 0.15,
12:     "window_base_height": 0.7,
13:     "window_height": 0.95,
14:     "door_height": 2.6,
15:     "pixel_scale": 0.01
16: }
17:
18: def is_valid_number(value):
19:     if value is None:
20:         return False
21:     if isinstance(value, (int, float)):
22:         return np.isfinite(value)
23:     return False
24:
25: def check_params_file(params_file):
26:     invalid_entries = []
27:     try:
28:         with open(params_file, "r") as f:
29:             params = json.load(f)
30:     except Exception as e:
31:         invalid_entries.append(f"Could not load JSON: {e}")

```

```

32:         return invalid_entries
33:
34:     for key in expected_keys.keys():
35:         val = params.get(key)
36:         if val is None:
37:             invalid_entries.append(f"missing '{key}'")
38:         elif not is_valid_number(val):
39:             invalid_entries.append(f"{key}={val} (invalid)")
40:     return invalid_entries
41:
42: def check_split(split="train"):
43:     split_dir = data_root / split
44:     total_files = 0
45:     good_files = 0
46:     bad_files = 0
47:
48:     print(f"\nChecking split: {split}")
49:     if not split_dir.exists():
50:         print(f"Warning: {split_dir} does not exist")
51:         return
52:
53:     # Recursively find all params.json files
54:     for params_file in split_dir.rglob("params.json"):
55:         total_files += 1
56:         invalid_entries = check_params_file(params_file)
57:
58:         if invalid_entries:
59:             print(f"[BAD] {params_file}")
60:             for entry in invalid_entries:
61:                 print(f"    - {entry}")
62:             bad_files += 1
63:         else:
64:             good_files += 1
65:
66:     print(f"\nSummary for split: {split}")
67:     print(f"Total files checked: {total_files}")
68:     print(f"Good files: {good_files}")
69:     print(f"Bad files: {bad_files}")
70:
71: def main():
72:     for split in ["train", "val", "test"]:
73:         check_split(split)
74:
75: if __name__ == "__main__":
76:     main()

```

■ File: setup.py

```

1: """
2: Setup script for the Neural-Geometric 3D Model Generator project
3: """
4:
5: from pathlib import Path
6: import os
7:
8: def create_project_structure():
9:     """Create the complete project directory structure"""
10:
11:     # Define directory structure
12:     directories = [
13:         "models",
14:         "training",
15:         "inference",
16:         "evaluation",
17:         "utils",
18:         "data/floorplans/train",
19:         "data/floorplans/val",
20:         "data/floorplans/test",
21:         "data/test_images",
22:         "checkpoints",
23:         "outputs",

```

```

24:         "demo_outputs",
25:         "evaluation_results",
26:         "logs"
27:     ]
28:
29:     # Create directories
30:     for directory in directories:
31:         Path(directory).mkdir(parents=True, exist_ok=True)
32:         print(f"Created directory: {directory}")
33:
34:     # Create __init__.py files
35:     init_files = [
36:         "models/__init__.py",
37:         "training/__init__.py",
38:         "inference/__init__.py",
39:         "evaluation/__init__.py",
40:         "utils/__init__.py"
41:     ]
42:
43:     init_content = {
44:         "models/__init__.py": '''"""
45: Model components for Neural-Geometric 3D Model Generator
46: """
47:
48: from .encoder import MultiScaleEncoder, ResidualBlock
49: from .heads import SegmentationHead, AttributeHead, SDFHead
50: from .dvx import DifferentiableVectorization
51: from .extrusion import DifferentiableExtrusion
52: from .model import NeuralGeometric3DGenerator
53:
54: __all__ = [
55:     'MultiScaleEncoder',
56:     'ResidualBlock',
57:     'SegmentationHead',
58:     'AttributeHead',
59:     'SDFHead',
60:     'DifferentiableVectorization',
61:     'DifferentiableExtrusion',
62:     'NeuralGeometric3DGenerator'
63: ]''',
64:
65:         "training/__init__.py": '''"""
66: Training components for Neural-Geometric 3D Model Generator
67: """
68:
69: from .losses import ResearchGradeLoss
70: from .trainer import MultiStageTrainer
71:
72: __all__ = [
73:     'ResearchGradeLoss',
74:     'MultiStageTrainer'
75: ]''',
76:
77:         "inference/__init__.py": '''"""
78: Inference components for Neural-Geometric 3D Model Generator
79: """
80:
81: from .engine import ResearchInferenceEngine
82:
83: __all__ = [
84:     'ResearchInferenceEngine'
85: ]''',
86:
87:         "evaluation/__init__.py": '''"""
88: Evaluation components for Neural-Geometric 3D Model Generator
89: """
90:
91: from .metrics import ModelEvaluator, evaluate_model, compute_iou, compute_3d_iou
92:
93: __all__ = [
94:     'ModelEvaluator',
95:     'evaluate_model',
96:     'compute_iou',

```

```

97:     'compute_3d_iou'
98: ]'''
99:
100:     "utils/__init__.py": '''"""
101: Utility functions for Neural-Geometric 3D Model Generator
102: """
103:
104: from .visualization import (
105:     plot_training_history,
106:     visualize_predictions,
107:     visualize_polygons,
108:     save_model_outputs,
109:     create_comparison_grid,
110:     create_3d_visualization
111: )
112:
113: __all__ = [
114:     'plot_training_history',
115:     'visualize_predictions',
116:     'visualize_polygons',
117:     'save_model_outputs',
118:     'create_comparison_grid',
119:     'create_3d_visualization'
120: ]'''
121: }
122:
123: for file_path, content in init_content.items():
124:     with open(file_path, 'w') as f:
125:         f.write(content)
126:         print(f"Created: {file_path}")
127:
128:
129: def create_sample_config():
130:     """Create a sample configuration file for easy customization"""
131:
132:     sample_config = '''"""
133: Sample configuration for quick customization
134: Copy this to config_custom.py and modify as needed
135: """
136:
137: from config import *
138:
139: # Custom configuration example
140: CUSTOM_DATA_CONFIG = DataConfig(
141:     data_dir="./my_data/floorplans",
142:     batch_size=16, # Larger batch if you have more GPU memory
143:     num_workers=8, # More workers if you have more CPU cores
144: )
145:
146: CUSTOM_TRAINING_CONFIG = TrainingConfig(
147:     stage1_epochs=30, # More epochs for better 2D learning
148:     stage2_epochs=20, # More DVX training
149:     stage3_epochs=50, # Longer end-to-end training
150:     stage1_lr=2e-4, # Higher learning rate
151: )
152:
153: CUSTOM_MODEL_CONFIG = ModelConfig(
154:     feature_dim=768, # Larger model
155:     voxel_size=128, # Higher resolution 3D
156:     max_polygons=30, # More polygons
157: )
158: '''
159:
160: with open("config_custom_example.py", "w") as f:
161:     f.write(sample_config)
162:     print("Created: config_custom_example.py")
163:
164:
165: def create_gitignore():
166:     """Create .gitignore file"""
167:
168:     gitignore_content = '''# Python
169: __pycache__/

```

```
170: *.py[cod]
171: *$py.class
172: *.so
173: .Python
174: build/
175: develop-eggs/
176: dist/
177: downloads/
178: eggs/
179: .eggs/
180: lib/
181: lib64/
182: parts/
183: sdist/
184: var/
185: wheels/
186: *.egg-info/
187: .installed.cfg
188: *.egg
189:
190: # PyTorch
191: *.pth
192: *.pt
193:
194: # Data
195: data/
196: datasets/
197: *.npz
198: *.obj
199: *.off
200: *.ply
201:
202: # Outputs
203: outputs/
204: results/
205: checkpoints/
206: logs/
207: demo_outputs/
208: evaluation_results/
209: training_progress/
210:
211: # IDE
212: .vscode/
213: .idea/
214: *.swp
215: *.swo
216:
217: # OS
218: .DS_Store
219: Thumbs.db
220:
221: # Jupyter
222: .ipynb_checkpoints
223:
224: # Environment
225: .env
226: .venv
227: env/
228: venv/
229:
230: # Images and videos
231: *.png
232: *.jpg
233: *.jpeg
234: *.gif
235: *.mp4
236: *.avi
237:
238: # Except sample images
239: !sample_images/
240: !docs/images/
241: ''
242:
```

```

243:     with open(".gitignore", "w") as f:
244:         f.write(gitignore_content)
245:     print("Created: .gitignore")
246:
247:
248: def main():
249:     """Main setup function"""
250:     print("Setting up Neural-Geometric 3D Model Generator project...")
251:     print("=" * 60)
252:
253:     # Create directory structure
254:     create_project_structure()
255:     print()
256:
257:     # Create sample config
258:     create_sample_config()
259:     print()
260:
261:     # Create gitignore
262:     create_gitignore()
263:     print()
264:
265:     print("Project setup completed!")
266:     print("=" * 60)
267:     print("IMPORTANT ACCURACY EXPECTATIONS:")
268:     print("- 90%+ accuracy in 2D-to-3D generation is extremely challenging")
269:     print("- Actual accuracy depends heavily on:")
270:     print("  * Dataset quality and size (need 10K+ samples)")
271:     print("  * Ground truth accuracy")
272:     print("  * Problem complexity (simple vs complex floorplans)")
273:     print("  * Evaluation metrics used")
274:     print("- Realistic expectations:")
275:     print("  * Segmentation: 75-85% mIoU with good data")
276:     print("  * 3D reconstruction: 60-75% IoU for architectural scenes")
277:     print("  * Polygon fitting: 70-80% accuracy")
278:     print("=" * 60)
279:     print("\nNext steps:")
280:     print("1. Install dependencies: pip install -r requirements.txt")
281:     print("2. Prepare high-quality dataset (critical for accuracy)")
282:     print("3. Run demo: python demo.py")
283:     print("4. Start training: python train.py")
284:
285:
286: if __name__ == "__main__":
287:     main()

```

File: train.py

```

=====
1: """
2: Enhanced training script for the Neural-Geometric 3D Model Generator
3: Implements novel training strategies: dynamic curriculum, adaptive weighting, cross-modal consistency
4: """
5:
6: import argparse
7: import torch
8: from torch.utils.data import DataLoader
9: from pathlib import Path
10: import json
11: import matplotlib.pyplot as plt
12: import numpy as np
13:
14: from dataset import AdvancedFloorPlanDataset
15: from models.model import NeuralGeometric3DGenerator
16: from training.trainer import AdaptiveMultiStageTrainer, MultiStageTrainer
17: from utils.visualization import plot_training_history, plot_curriculum_analysis
18: from config import (
19:     DEFAULT_DATA_CONFIG,
20:     DEFAULT_MODEL_CONFIG,
21:     DEFAULT_TRAINING_CONFIG,
22:     DEFAULT_LOSS_CONFIG,
23:     TrainingConfig,

```



```

24: CurriculumConfig
25: )
26:
27:
28: def create_enhanced_config(args):
29:     """Create enhanced training configuration with novel strategies"""
30:     config = TrainingConfig()
31:
32:     # Basic settings
33:     config.device = args.device or ("cuda" if torch.cuda.is_available() else "cpu")
34:
35:     # Dynamic curriculum settings
36:     if args.dynamic_curriculum:
37:         config.curriculum = CurriculumConfig()
38:         config.curriculum.use_dynamic_curriculum = True
39:         config.curriculum.stage_switch_patience = args.patience
40:         config.curriculum.min_improvement_threshold = args.min_improvement
41:
42:         # Adjust epoch limits for dynamic training
43:         config.max_stage1_epochs = args.max_stage1_epochs
44:         config.max_stage2_epochs = args.max_stage2_epochs
45:         config.max_stage3_epochs = args.max_stage3_epochs
46:
47:         print("Dynamic curriculum learning enabled")
48:         print(f" Stage switch patience: {config.curriculum.stage_switch_patience}")
49:         print(f" Min improvement threshold: {config.curriculum.min_improvement_threshold}")
50:     else:
51:         # Disable dynamic curriculum for traditional training
52:         config.curriculum.use_dynamic_curriculum = False
53:         print("Using traditional fixed-epoch training")
54:
55:     # GradNorm dynamic weighting
56:     if args.gradnorm:
57:         config.curriculum.use_gradnorm = True
58:         config.curriculum.gradnorm_alpha = args.gradnorm_alpha
59:         config.curriculum.gradnorm_update_freq = args.gradnorm_freq
60:         print(f"GradNorm dynamic weighting enabled (alpha={args.gradnorm_alpha})")
61:
62:     # Topology-aware scheduling
63:     if args.topology_schedule != "static":
64:         config.curriculum.topology_schedule = args.topology_schedule
65:         config.curriculum.topology_start_weight = args.topology_start
66:         config.curriculum.topology_end_weight = args.topology_end
67:         print(f"Topology-aware scheduling: {args.topology_schedule}")
68:         print(f" Weights: {args.topology_start} -> {args.topology_end}")
69:
70:     return config
71:
72:
73: def create_enhanced_model(args):
74:     """Create enhanced model with auxiliary heads"""
75:     model = NeuralGeometric3DGenerator(
76:         input_channels=args.input_channels,
77:         num_classes=args.num_classes,
78:         feature_dim=args.feature_dim,
79:         num_attributes=args.num_attributes,
80:         voxel_size=args.voxel_size,
81:         max_polygons=args.max_polygons,
82:         max_points=args.max_points,
83:         use_latent_consistency=args.latent_consistency,
84:         use_graph_constraints=args.graph_constraints,
85:         latent_embedding_dim=args.embedding_dim
86:     )
87:
88:     print(f"Enhanced model created:")
89:     print(f" Feature dim: {args.feature_dim}")
90:     print(f" Latent consistency: {args.latent_consistency}")
91:     print(f" Graph constraints: {args.graph_constraints}")
92:
93:     # Print parameter counts
94:     total_params = sum(p.numel() for p in model.parameters())
95:     trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
96:     print(f" Total parameters: {total_params:,}")

```

```

97:     print(f" Trainable parameters: {trainable_params:,}")
98:
99:     return model
100:
101:
102: def visualize_training_results(history, output_dir):
103:     """Create comprehensive training visualizations"""
104:     output_dir = Path(output_dir)
105:     output_dir.mkdir(exist_ok=True)
106:
107:     # Traditional loss curves
108:     plot_training_history(history, save_path=str(output_dir / "training_history.png"))
109:
110:     # Novel curriculum analysis plots
111:     if "stage_transitions" in history and history["stage_transitions"]:
112:         plot_curriculum_analysis(history, save_path=str(output_dir / "curriculum_analysis.png"))
113:
114:     # Dynamic weight evolution
115:     if "dynamic_weights" in history and history["dynamic_weights"]:
116:         plt.figure(figsize=(12, 8))
117:
118:         # Extract weight evolution data
119:         epochs = [entry["epoch"] for entry in history["dynamic_weights"]]
120:         weight_names = list(history["dynamic_weights"][0]["weights"].keys())
121:
122:         for weight_name in weight_names:
123:             weights = [entry["weights"].get(weight_name, 0) for entry in history["dynamic_weights"]]
124:             if any(w > 0.001 for w in weights): # Only plot significant weights
125:                 plt.plot(epochs, weights, label=weight_name, linewidth=2)
126:
127:             plt.xlabel("Global Epoch")
128:             plt.ylabel("Loss Weight")
129:             plt.title("Dynamic Loss Weight Evolution")
130:             plt.legend()
131:             plt.grid(True, alpha=0.3)
132:             plt.tight_layout()
133:             plt.savefig(output_dir / "weight_evolution.png", dpi=300)
134:             plt.close()
135:
136:     # Component loss breakdown
137:     fig, axes = plt.subplots(1, 3, figsize=(18, 5))
138:     stage_names = ["stage1", "stage2", "stage3"]
139:
140:     for idx, stage_name in enumerate(stage_names):
141:         if stage_name in history and "component_losses" in history[stage_name]:
142:             component_data = history[stage_name]["component_losses"]
143:             if component_data:
144:                 # Get component names from first entry
145:                 component_names = list(component_data[0].keys())
146:
147:                 for comp_name in component_names:
148:                     if comp_name in ['seg', 'dice', 'polygon', 'voxel', 'topology',
149:                                     'latent_consistency', 'graph']:
150:                         values = [entry.get(comp_name, 0) for entry in component_data]
151:                         if any(v > 0.001 for v in values): # Only plot significant losses
152:                             axes[idx].plot(values, label=comp_name, linewidth=2)
153:
154:                     axes[idx].set_title(f"{stage_name.upper()} Component Losses")
155:                     axes[idx].set_xlabel("Epoch")
156:                     axes[idx].set_ylabel("Loss Value")
157:                     axes[idx].legend()
158:                     axes[idx].grid(True, alpha=0.3)
159:
160:     plt.tight_layout()
161:     plt.savefig(output_dir / "component_losses.png", dpi=300)
162:     plt.close()
163:
164:     print(f"Training visualizations saved to {output_dir}")
165:
166:
167: def save_training_summary(history, config, output_dir):
168:     """Save comprehensive training summary"""
169:     output_dir = Path(output_dir)

```

```

170:
171:     summary = {
172:         "training_config": {
173:             "dynamic_curriculum": config.curriculum.use_dynamic_curriculum,
174:             "gradnorm_enabled": config.curriculum.use_gradnorm,
175:             "topology_schedule": config.curriculum.topology_schedule,
176:             "max_epochs": [config.max_stage1_epochs, config.max_stage2_epochs, config.max_stage3_epochs]
177:         },
178:         "training_results": {},
179:         "novel_strategies_summary": {}
180:     }
181:
182:     # Training results
183:     for stage_name, data in history.items():
184:         if isinstance(data, dict) and "val_loss" in data and data["val_loss"]:
185:             summary["training_results"][stage_name] = {
186:                 "final_val_loss": data["val_loss"][-1],
187:                 "best_val_loss": min(data["val_loss"]),
188:                 "epochs_trained": len(data["val_loss"])
189:             }
190:
191:     # Novel strategies summary
192:     if "stage_transitions" in history:
193:         summary["novel_strategies_summary"]["adaptive_transitions"] = len(history["stage_transitions"])
194:
195:     if "dynamic_weights" in history:
196:         summary["novel_strategies_summary"]["weight_updates"] = len(history["dynamic_weights"])
197:
198:     if "curriculum_events" in history:
199:         summary["novel_strategies_summary"]["curriculum_events"] = len(history["curriculum_events"])
200:
201:     # Save as JSON
202:     with open(output_dir / "training_summary.json", 'w') as f:
203:         json.dump(summary, f, indent=2)
204:
205:     print(f"Training summary saved to {output_dir / 'training_summary.json'}")
206:
207:
208: def main():
209:     parser = argparse.ArgumentParser(description="Enhanced Neural-Geometric 3D Model Generator Training")
210:
211:     # Basic arguments
212:     parser.add_argument("--data_dir", type=str, default="./data/floorplans",
213:                         help="Path to dataset directory")
214:     parser.add_argument("--batch_size", type=int, default=2, help="Batch size")
215:     parser.add_argument("--num_workers", type=int, default=4, help="Number of data workers")
216:     parser.add_argument("--device", type=str, default=None, help="Training device")
217:     parser.add_argument("--resume", type=str, default=None, help="Resume from checkpoint")
218:     parser.add_argument("--output_dir", type=str, default="./checkpoints",
219:                         help="Output directory for checkpoints")
220:
221:     # Training mode selection
222:     parser.add_argument("--training_mode", type=str, choices=["traditional", "adaptive"],
223:                         default="adaptive", help="Training mode (traditional fixed epochs vs adaptive)")
224:     parser.add_argument("--stage", type=str, choices=["1", "2", "3", "all"], default="all",
225:                         help="Training stage to run (only for traditional mode)")
226:
227:     # Novel training strategies
228:     parser.add_argument("--dynamic-curriculum", action="store_true", default=True,
229:                         help="Enable adaptive stage transitioning")
230:     parser.add_argument("--patience", type=int, default=5,
231:                         help="Epochs without improvement before stage transition")
232:     parser.add_argument("--min-improvement", type=float, default=0.001,
233:                         help="Minimum relative improvement threshold")
234:
235:     parser.add_argument("--gradnorm", action="store_true", default=True,
236:                         help="Enable GradNorm dynamic loss weighting")
237:     parser.add_argument("--gradnorm-alpha", type=float, default=0.12,
238:                         help="GradNorm restoring force parameter")
239:     parser.add_argument("--gradnorm-freq", type=int, default=5,
240:                         help="GradNorm update frequency (batches)")
241:
242:     parser.add_argument("--topology-schedule", type=str,

```

```

243:             choices=["static", "progressive", "linear_ramp"],
244:             default="progressive", help="Topology loss scheduling strategy")
245: parser.add_argument("--topology-start", type=float, default=0.1,
246:                     help="Starting weight for topology loss")
247: parser.add_argument("--topology-end", type=float, default=1.0,
248:                     help="Ending weight for topology loss")
249:
250: # Model enhancements
251: parser.add_argument("--latent-consistency", action="store_true", default=True,
252:                     help="Enable cross-modal latent consistency loss")
253: parser.add_argument("--graph-constraints", action="store_true", default=True,
254:                     help="Enable graph-based topology constraints")
255: parser.add_argument("--embedding-dim", type=int, default=256,
256:                     help="Latent embedding dimension")
257:
258: # Model architecture
259: parser.add_argument("--input_channels", type=int, default=3, help="Input image channels")
260: parser.add_argument("--num_classes", type=int, default=5, help="Number of segmentation classes")
261: parser.add_argument("--feature_dim", type=int, default=768, help="Feature dimension")
262: parser.add_argument("--num_attributes", type=int, default=6, help="Number of attribute predictions")
263: parser.add_argument("--voxel_size", type=int, default=64, help="3D voxel grid size")
264: parser.add_argument("--max_polygons", type=int, default=30, help="Maximum number of polygons")
265: parser.add_argument("--max_points", type=int, default=100, help="Maximum points per polygon")
266:
267: # Dynamic epoch limits
268: parser.add_argument("--max-stage1-epochs", type=int, default=50, help="Max epochs for Stage 1")
269: parser.add_argument("--max-stage2-epochs", type=int, default=35, help="Max epochs for Stage 2")
270: parser.add_argument("--max-stage3-epochs", type=int, default=100, help="Max epochs for Stage 3")
271:
272: parser.add_argument("--persistent_workers", action="store_true", default=False, help="Keep DataLoader workers")
273:
274: parser.add_argument("--prefetch_factor", type=int, default=2, help="Number of batches preloaded by each worker")
275:
276:
277: args = parser.parse_args()
278:
279: # Setup device
280: device = args.device or ("cuda" if torch.cuda.is_available() else "cpu")
281: print(f"Using device: {device}")
282:
283: import torch.backends.cudnn as cudnn
284: if device == "cuda":
285:     cudnn.benchmark = True
286:
287: # Create output directory
288: output_dir = Path(args.output_dir)
289: output_dir.mkdir(exist_ok=True)
290:
291: # Create enhanced configuration
292: config = create_enhanced_config(args)
293:
294: print("\n" + "="*80)
295: print("NEURAL-GEOMETRIC 3D MODEL GENERATOR - ENHANCED TRAINING")
296: print("="*80)
297: print("Novel Training Strategies Enabled:")
298: if config.curriculum.use_dynamic_curriculum:
299:     print("? Adaptive Stage Transitioning (Dynamic Curriculum)")
300: if config.curriculum.use_gradnorm:
301:     print("? Multi-objective Optimization with GradNorm")
302: if config.curriculum.topology_schedule != "static":
303:     print("? Topology-aware Loss Scheduling")
304: if args.latent_consistency:
305:     print("? Cross-modal Latent Consistency Learning")
306: if args.graph_constraints:
307:     print("? Graph-based Topology Constraints")
308: print("="*80)
309:
310: # Create datasets
311: print("\nLoading datasets...")
312: train_dataset = AdvancedFloorPlanDataset(
313:     args.data_dir, split="train", augment=True
314: )
315: val_dataset = AdvancedFloorPlanDataset(

```

```

316:         args.data_dir, split="val", augment=False
317:     )
318:
319:     print(f"Train samples: {len(train_dataset)}")
320:     print(f"Validation samples: {len(val_dataset)}")
321:
322:     if len(train_dataset) == 0:
323:         print("Error: No training samples found!")
324:         return
325:
326:     # Create data loaders
327:     train_loader = DataLoader(
328:         train_dataset,
329:         batch_size=args.batch_size,
330:         shuffle=True,
331:         num_workers=args.num_workers,
332:         pin_memory=True,
333:         drop_last=True,
334:         persistent_workers=args.persistent_workers if args.num_workers > 0 else False,
335:         prefetch_factor=args.prefetch_factor if args.num_workers > 0 else None
336:     )
337:
338:     val_loader = DataLoader(
339:         val_dataset,
340:         batch_size=max(1, args.batch_size),
341:         shuffle=False,
342:         num_workers=max(1, args.num_workers // 2),
343:         pin_memory=True,
344:         drop_last=False,
345:         persistent_workers=args.persistent_workers if args.num_workers > 0 else False,
346:         prefetch_factor=args.prefetch_factor if args.num_workers > 0 else None
347:     )
348:
349:     # Create enhanced model
350:     print("\nInitializing enhanced model...")
351:     model = create_enhanced_model(args)
352:
353:     # Create appropriate trainer
354:     if args.training_mode == "adaptive":
355:         print("\nUsing Adaptive Multi-Stage Trainer with Novel Strategies")
356:         trainer = AdaptiveMultiStageTrainer(
357:             model=model,
358:             train_loader=train_loader,
359:             val_loader=val_loader,
360:             device=device,
361:             config=config
362:         )
363:     else:
364:         print("\nUsing Traditional Multi-Stage Trainer")
365:         trainer = MultiStageTrainer(
366:             model=model,
367:             train_loader=train_loader,
368:             val_loader=val_loader,
369:             device=device,
370:             config=config
371:         )
372:
373:     # Resume from checkpoint if specified
374:     if args.resume:
375:         print(f"Resuming from checkpoint: {args.resume}")
376:         trainer.load_checkpoint(args.resume)
377:
378:     # Run training
379:     if args.training_mode == "adaptive" or args.stage == "all":
380:         print("\nStarting adaptive multi-stage training with novel strategies...")
381:         history = trainer.train_all_stages()
382:     else:
383:         # Traditional single-stage training
384:         stage_num = int(args.stage)
385:         print(f"Training Stage {stage_num} only...")
386:         if stage_num == 1:
387:             trainer.train_stage1()
388:         elif stage_num == 2:

```

```

389:         trainer.train_stage2()
390:     elif stage_num == 3:
391:         trainer.train_stage3()
392:         history = trainer.history
393:
394:     # Save final model
395:     final_model_path = output_dir / "final_enhanced_model.pth"
396:     if hasattr(trainer, '_save_checkpoint'):
397:         trainer._save_checkpoint(str(final_model_path))
398:     print(f"Final model saved to: {final_model_path}")
399:
400:     # Create comprehensive visualizations
401:     print("\nGenerating training analysis...")
402:     visualize_training_results(history, output_dir)
403:
404:     # Save training summary
405:     save_training_summary(history, config, output_dir)
406:
407:     print(f"\n? Enhanced training completed successfully!")
408:     print(f"? Results saved to: {output_dir}")
409:     print("\nNovel contributions implemented:")
410:     print("- Dynamic curriculum learning with adaptive stage transitions")
411:     print("- Multi-objective optimization with gradient-based reweighting")
412:     print("- Topology-aware progressive constraint injection")
413:     print("- Cross-modal latent consistency learning")
414:     print("- Graph-based architectural constraint learning")
415:
416:
417: if __name__ == "__main__":
418:     main()

```

■ File: traininglosses.py

```

=====
1: """
2: Advanced loss functions for multi-task training with dynamic weighting
3: Enhanced with cross-modal consistency, graph constraints, and GradNorm
4: Modified to support conditional geometric losses via run_full_geometric flag
5: """
6:
7: import torch
8: import torch.nn as nn
9: import torch.nn.functional as F
10: import cv2
11: import numpy as np
12: from typing import Dict, Optional, Tuple, List
13: import networkx as nx
14:
15:
16: class DynamicLossWeighter:
17:     """Implements GradNorm and other dynamic weighting strategies"""
18:
19:     def __init__(self, loss_names: List[str], alpha: float = 0.12, device: str = 'cuda'):
20:         self.loss_names = loss_names
21:         self.alpha = alpha
22:         self.device = device
23:
24:         # Initialize weights
25:         self.weights = {name: 1.0 for name in loss_names} # FIX: keep floats, easier logging
26:         self.initial_task_losses = None
27:         self.running_mean_losses = {name: 0.0 for name in loss_names}
28:         self.update_count = 0
29:
30:     def update_weights(self, task_losses: Dict[str, torch.Tensor],
31:                       shared_parameters, update_freq: int = 10):
32:         """Update loss weights using GradNorm algorithm with enhanced stability"""
33:         if self.update_count % update_freq != 0:
34:             self.update_count += 1
35:             return self.weights
36:
37:         # Store initial losses on first update
38:         if self.initial_task_losses is None:

```

```

39:         self.initial_task_losses = {name: loss.item() for name, loss in task_losses.items()}
40:
41:     # Update running mean
42:     for name, loss in task_losses.items():
43:         self.running_mean_losses[name] = (0.9 * self.running_mean_losses[name] +
44:                                           0.1 * loss.item())
45:
46:     # Calculate relative decrease rates
47:     loss_ratios = {}
48:     for name in self.loss_names:
49:         if name in task_losses:
50:             current_loss = self.running_mean_losses[name]
51:             initial_loss = self.initial_task_losses[name]
52:             loss_ratios[name] = current_loss / (initial_loss + 1e-8)
53:
54:     # Calculate average relative decrease
55:     if not loss_ratios:
56:         self.update_count += 1
57:         return self.weights
58:     avg_loss_ratio = sum(loss_ratios.values()) / len(loss_ratios)
59:
60:     # Calculate gradient norms with enhanced safety checks
61:     grad_norms = {}
62:     for name in self.loss_names:
63:         if name in task_losses:
64:             # Check if loss is finite before gradient computation
65:             if not torch.isfinite(task_losses[name]):
66:                 print(f"Warning: Non-finite loss for {name}, skipping gradient computation")
67:                 continue
68:
69:             try:
70:                 grads = torch.autograd.grad(
71:                     task_losses[name], shared_parameters,
72:                     retain_graph=True, create_graph=False, allow_unused=True
73:                 )
74:                 grad_norm = 0.0
75:                 for grad in grads:
76:                     if grad is not None and torch.isfinite(grad).all():
77:                         grad_norm += grad.norm().item() ** 2
78:                     elif grad is not None:
79:                         print(f"Warning: Non-finite gradient detected for {name}")
80:                         grad_norm = 0.0
81:                         break
82:
83:                 if grad_norm > 0:
84:                     grad_norms[name] = grad_norm ** 0.5
85:             except Exception as e:
86:                 print(f"Warning: Gradient computation failed for {name}: {e}")
87:                 continue
88:
89:     if not grad_norms:
90:         self.update_count += 1
91:         return self.weights
92:
93:     avg_grad_norm = sum(grad_norms.values()) / len(grad_norms)
94:
95:     # Enhanced safety parameters
96:     epsilon = 1e-6 # Increased minimum threshold
97:     max_target_grad = 10.0 # Reduced maximum allowed gradient
98:     max_weight_update = 2.0 # Reduced maximum weight change
99:
100:    for name in self.loss_names:
101:        if name in grad_norms and name in loss_ratios:
102:            target_grad = avg_grad_norm * (loss_ratios[name] ** self.alpha)
103:            target_grad = min(target_grad, max_target_grad)
104:
105:            safe_grad_norm = max(grad_norms[name], epsilon)
106:            if safe_grad_norm == epsilon:
107:                print(f"? Small gradient norm for {name}, clamped to {epsilon}")
108:
109:            weight_update = target_grad / safe_grad_norm
110:            weight_update = min(weight_update, max_weight_update)
111:            if weight_update == max_weight_update:

```

```

112:         print(f"? Weight update for {name} clamped to {max_weight_update}")
113:
114:         # More conservative momentum update
115:         new_w = 0.95 * self.weights[name] + 0.05 * float(weight_update) # Slower adaptation
116:         self.weights[name] = float(np.clip(new_w, 0.01, 3.0)) # Tighter bounds
117:
118:     self.update_count += 1
119:     return self.weights
120:
121: class GraphTopologyExtractor:
122:     """Extracts graph structure from segmentation for topology constraints"""
123:
124:     @staticmethod
125:     def extract_room_graph(segmentation: torch.Tensor) -> Dict[str, torch.Tensor]:
126:         """Extract room connectivity graph from segmentation mask"""
127:         B, C, H, W = segmentation.shape
128:         device = segmentation.device
129:
130:         # Get room predictions (assume classes: 0=bg, 1=wall, 2=door, 3=window, 4=room)
131:         room_probs = F.softmax(segmentation, dim=1)
132:         room_mask = room_probs[:, 4] if C > 4 else torch.zeros((B, H, W), device=device)
133:         wall_mask = room_probs[:, 1] if C > 1 else torch.zeros((B, H, W), device=device)
134:
135:         # Simple connectivity: rooms connected if they share wall boundary
136:         adjacency_matrices = []
137:         room_features = []
138:
139:         for b in range(B):
140:             room_b = room_mask[b].detach().cpu().numpy()
141:             wall_b = wall_mask[b].detach().cpu().numpy()
142:
143:             # Find connected components (rooms)
144:             try:
145:                 from scipy import ndimage
146:                 labeled_rooms, num_rooms = ndimage.label(room_b > 0.5)
147:
148:                 # Create adjacency matrix
149:                 adj_matrix = np.zeros((max(num_rooms, 1), max(num_rooms, 1)))
150:                 room_centroids = []
151:
152:                 for i in range(1, num_rooms + 1):
153:                     room_i_mask = (labeled_rooms == i)
154:                     if np.sum(room_i_mask) > 0:
155:                         centroid = ndimage.center_of_mass(room_i_mask)
156:                         room_centroids.append(centroid)
157:
158:                 # Check connectivity to other rooms through walls
159:                 for j in range(i + 1, num_rooms + 1):
160:                     room_j_mask = (labeled_rooms == j)
161:                     if np.sum(room_j_mask) > 0:
162:                         # Check if rooms are connected via wall adjacency
163:                         connectivity = GraphTopologyExtractor._check_room_connectivity(
164:                             room_i_mask, room_j_mask, wall_b
165:                         )
166:                         adj_matrix[i-1, j-1] = connectivity
167:                         adj_matrix[j-1, i-1] = connectivity
168:
169:                 # Convert to tensor
170:                 adj_tensor = torch.from_numpy(adj_matrix).float().to(device)
171:                 centroids_tensor = torch.from_numpy(np.array(room_centroids) if room_centroids else np.zeros((0, 2)))
172:
173:             except ImportError:
174:                 # Fallback if scipy not available
175:                 adj_tensor = torch.zeros((1, 1), device=device)
176:                 centroids_tensor = torch.zeros((0, 2), device=device)
177:
178:             adjacency_matrices.append(adj_tensor)
179:             room_features.append(centroids_tensor)
180:
181:         return {
182:             "adjacency_matrices": adjacency_matrices,
183:             "room_features": room_features
184:         }

```



```

185:
186:     @staticmethod
187:     def _check_room_connectivity(room1_mask, room2_mask, wall_mask):
188:         """Check if two rooms are connected through walls"""
189:         try:
190:             from scipy.ndimage import binary_dilation
191:
192:             # Dilate room masks to check wall adjacency
193:             dilated1 = binary_dilation(room1_mask, iterations=2)
194:             dilated2 = binary_dilation(room2_mask, iterations=2)
195:
196:             # Check overlap through wall areas
197:             wall_overlap = (dilated1 & dilated2) & (wall_mask > 0.3)
198:             return float(np.sum(wall_overlap) > 0)
199:         except ImportError:
200:             # Simple distance-based fallback
201:             return 0.0
202:
203:
204: class ResearchGradeLoss(nn.Module):
205:     """
206:     Multi-task loss combining:
207:     - Traditional losses (segmentation, SDF, attributes, polygons, voxels, topology)
208:     - NEW: Cross-modal latent consistency
209:     - NEW: Graph-based topology constraints
210:     - NEW: Dynamic loss weighting via GradNorm
211:     - NEW: Conditional geometric losses based on run_full_geometric flag
212:     """
213:
214:     def __init__(
215:         self,
216:         seg_weight: float = 1.0,
217:         dice_weight: float = 1.0,
218:         sdf_weight: float = 0.5,
219:         attr_weight: float = 1.0,
220:         polygon_weight: float = 1.0,
221:         voxel_weight: float = 1.0,
222:         topology_weight: float = 0.5,
223:         latent_consistency_weight: float = 0.5,
224:         graph_constraint_weight: float = 0.3,
225:         enable_dynamic_weighting: bool = True,
226:         gradnorm_alpha: float = 0.12,
227:         device: str = 'cuda',
228:         weight_update_freq: int = 10,
229:         weight_momentum: float = 0.9,
230:     ):
231:         super().__init__()
232:
233:         # Store initial weights
234:         self.initial_weights = {
235:             'seg': float(seg_weight),
236:             'dice': float(dice_weight),
237:             'sdf': float(sdf_weight),
238:             'attr': float(attr_weight),
239:             'polygon': float(polygon_weight),
240:             'voxel': float(voxel_weight),
241:             'topology': float(topology_weight),
242:             'latent_consistency': float(latent_consistency_weight),
243:             'graph': float(graph_constraint_weight)
244:         }
245:
246:         # Current weights (will be dynamically updated)
247:         self.weights = self.initial_weights.copy()
248:
249:         # Core losses
250:         self.ce_loss = nn.CrossEntropyLoss()
251:         self.mse_loss = nn.MSELoss()
252:         self.l1_loss = nn.L1Loss()
253:         self.cosine_loss = nn.CosineEmbeddingLoss()
254:
255:         # Dynamic weighting
256:         self.enable_dynamic_weighting = enable_dynamic_weighting
257:         if enable_dynamic_weighting:

```

```

258:         self.loss_weighter = DynamicLossWeighter(
259:             list(self.initial_weights.keys()), alpha=gradnorm_alpha, device=device,
260:         )
261:         self.loss_weighter.update_freq = weight_update_freq
262:         self.loss_weighter.momentum = weight_momentum
263:
264:     self.device = device
265:
266:     def update_loss_weights(self, new_weights: Dict[str, float]):
267:         """Update loss weights (called by trainer for curriculum scheduling)"""
268:         for key, value in new_weights.items():
269:             if key in self.weights:
270:                 self.weights[key] = float(value)
271:
272:     def forward(self, predictions: dict, targets: dict, shared_parameters=None, run_full_geometric=True):
273:         """
274:         Compute multi-task loss with conditional geometric computation and dynamic weighting.
275:
276:         Args:
277:             predictions: Model predictions dict
278:             targets: Ground truth targets dict
279:             shared_parameters: Model parameters for GradNorm (optional)
280:             run_full_geometric: Whether geometric losses should be computed
281:
282:         Returns:
283:             tuple: (total_loss, individual_losses_dict)
284:         """
285:         # Input validation - check for NaN/Inf values
286:         for name, tensor in predictions.items():
287:             if torch.is_tensor(tensor) and (torch.isnan(tensor).any() or torch.isinf(tensor).any()):
288:                 print(f"WARNING: NaN/Inf in predictions[{name}] - zeroing out")
289:                 predictions[name] = torch.zeros_like(tensor)
290:
291:         for name, tensor in targets.items():
292:             if torch.is_tensor(tensor) and (torch.isnan(tensor).any() or torch.isinf(tensor).any()):
293:                 print(f"WARNING: NaN/Inf in targets[{name}] - zeroing out")
294:                 targets[name] = torch.zeros_like(tensor)
295:
296:         device = self._get_device_from_inputs(predictions, targets)
297:         losses = {}
298:         total_loss = torch.tensor(0.0, device=device)
299:
300:         # ---- 1) Core losses (always computed - lightweight) ----
301:         if "segmentation" in predictions and "mask" in targets:
302:             seg_pred = predictions["segmentation"]
303:             seg_target = targets["mask"].long()
304:
305:             ce_loss = self.ce_loss(seg_pred, seg_target)
306:             losses["ce"] = ce_loss
307:             losses["seg"] = ce_loss # alias for dynamic weighting
308:
309:             dice_loss = self._dice_loss(seg_pred, seg_target)
310:             losses["dice"] = dice_loss
311:
312:         if "sdf" in predictions and "mask" in targets:
313:             sdf_pred = predictions["sdf"]
314:             sdf_pred = torch.clamp(sdf_pred, -1.0, 1.0) # FIX: prevent huge values
315:             sdf_target = self._mask_to_sdf(targets["mask"])
316:             sdf_target = sdf_target.to(sdf_pred.device).type_as(sdf_pred)
317:             losses["sdf"] = self.mse_loss(sdf_pred, sdf_target)
318:
319:         if "attributes" in predictions and "attributes" in targets:
320:             pred_attr = predictions["attributes"].float()
321:             tgt_attr = targets["attributes"].float().to(pred_attr.device)
322:             losses["attr"] = self.l1_loss(pred_attr, tgt_attr)
323:
324:         # ---- 2) Conditional geometric losses (heavy operations) ----
325:         if run_full_geometric:
326:             # Polygon loss (only if model produced polygons)
327:             if ("polygons" in predictions and predictions["polygons"] is not None and
328:                 "polygons_gt" in targets):
329:                 losses["polygon"] = self._polygon_loss(predictions, targets["polygons_gt"])
330:             else:

```

```

331:         # Zero loss if polygons not available
332:         losses["polygon"] = torch.tensor(0.0, device=device)
333:
334:     # Voxel loss (only if model produced voxels)
335:     if ("voxels_pred" in predictions and predictions["voxels_pred"] is not None and
336:         "voxels_gt" in targets):
337:         pred_vox = predictions["voxels_pred"].float()
338:         tgt_vox = targets["voxels_gt"].float().to(pred_vox.device)
339:         losses["voxel"] = self._voxel_iou_loss(pred_vox, tgt_vox)
340:     else:
341:         # Zero loss if voxels not available
342:         losses["voxel"] = torch.tensor(0.0, device=device)
343:
344:     # Cross-modal latent consistency (only if embeddings available)
345:     if ("latent_2d_embedding" in predictions and "latent_3d_embedding" in predictions and
346:         predictions["latent_2d_embedding"] is not None and predictions["latent_3d_embedding"] is not None):
347:         consistency_loss = self._latent_consistency_loss(
348:             predictions["latent_2d_embedding"],
349:             predictions["latent_3d_embedding"]
350:         )
351:         losses["latent_consistency"] = consistency_loss
352:     else:
353:         losses["latent_consistency"] = torch.tensor(0.0, device=device)
354:
355:     # When geometric computation is skipped, use zero losses
356:     losses["polygon"] = torch.tensor(0.0, device=device)
357:     losses["voxel"] = torch.tensor(0.0, device=device)
358:     losses["latent_consistency"] = torch.tensor(0.0, device=device)
359:
360:     # ---- 3) Independent auxiliary losses (always computed if enabled) ----
361:     # Traditional topology loss
362:     if "segmentation" in predictions:
363:         losses["topology"] = self._topology_loss(predictions["segmentation"])
364:
365:     # Graph-based topology constraints
366:     if "segmentation" in predictions:
367:         graph_loss = self._graph_topology_loss(predictions["segmentation"])
368:         losses["graph"] = graph_loss
369:
370:     # ---- 4) Apply weighting ----
371:     if self.enable_dynamic_weighting and shared_parameters is not None:
372:         # Only include differentiable losses for GradNorm
373:         task_losses = {
374:             name: loss for name, loss in losses.items()
375:             if name in self.weights and isinstance(loss, torch.Tensor) and loss.requires_grad_()
376:         }
377:
378:         dynamic_weights = self.loss_weighter.update_weights(task_losses, shared_parameters)
379:
380:         # Apply weights (dynamic for diff losses, static for non-diff losses)
381:         for name, loss in losses.items():
382:             if name in self.weights:
383:                 if name in dynamic_weights:
384:                     weight = dynamic_weights[name]
385:                 else:
386:                     weight = self.weights[name]
387:                 total_loss = total_loss + weight * loss
388:     else:
389:         # Static weights
390:         for name, loss in losses.items():
391:             if name in self.weights:
392:                 total_loss = total_loss + self.weights[name] * loss
393:
394:     # Final NaN/Inf guard
395:     for k, v in list(losses.items()):
396:         if torch.isnan(v).any() or torch.isinf(v).any():
397:             print(f"[Warning] {k} loss is NaN/Inf ? zeroed out")
398:             losses[k] = torch.tensor(0.0, device=device)
399:
400:     losses["total"] = total_loss
401:     return total_loss, losses
402:
403: def __call__(self, predictions: dict, targets: dict, shared_parameters=None, run_full_geometric=True):

```

```

404:         """
405:         Convenience method for trainer compatibility
406:
407:         Args:
408:             predictions: Model predictions dict
409:             targets: Ground truth targets dict
410:             shared_parameters: Model parameters for GradNorm (optional)
411:             run_full_geometric: Whether to compute geometric losses
412:
413:         Returns:
414:             tuple: (total_loss, individual_losses_dict)
415:         """
416:         return self.forward(predictions, targets, shared_parameters, run_full_geometric)
417:
418:     def _get_device_from_inputs(self, predictions, targets):
419:         """Helper to determine device from inputs"""
420:         for pred_dict in [predictions, targets]:
421:             for value in pred_dict.values():
422:                 if torch.is_tensor(value):
423:                     return value.device
424:         return self.device
425:
426:     # ---- NEW: Cross-modal latent consistency loss ----
427:     def _latent_consistency_loss(self, embedding_2d: torch.Tensor, embedding_3d: torch.Tensor) -> torch.Tensor:
428:         """
429:         Ensure 2D floorplan embeddings match 3D voxelized structure embeddings
430:         embedding_2d: [B, D] - 2D floorplan embeddings
431:         embedding_3d: [B, D] - 3D structure embeddings
432:         """
433:         if embedding_2d.shape != embedding_3d.shape:
434:             # Project to same dimension if needed
435:             min_dim = min(embedding_2d.shape[-1], embedding_3d.shape[-1])
436:             embedding_2d = embedding_2d[..., :min_dim]
437:             embedding_3d = embedding_3d[..., :min_dim]
438:
439:             # Cosine similarity loss (maximize similarity)
440:             target = torch.ones(embedding_2d.shape[0], device=embedding_2d.device)
441:             cosine_loss = self.cosine_loss(embedding_2d, embedding_3d, target)
442:
443:             # L2 consistency loss
444:             l2_loss = F.mse_loss(embedding_2d, embedding_3d)
445:
446:             return 0.7 * cosine_loss + 0.3 * l2_loss
447:
448:     # ---- NEW: Graph-based topology constraints ----
449:     def _graph_topology_loss(self, segmentation_logits: torch.Tensor) -> torch.Tensor:
450:         """
451:         Graph-based topology constraints on room connectivity
452:         segmentation_logits: [B, C, H, W]
453:         """
454:         try:
455:             # Extract graph structure
456:             graph_data = GraphTopologyExtractor.extract_room_graph(segmentation_logits)
457:             device = segmentation_logits.device
458:
459:             total_graph_loss = torch.tensor(0.0, device=device)
460:             batch_size = segmentation_logits.shape[0]
461:
462:             for b in range(batch_size):
463:                 if b < len(graph_data["adjacency_matrices"]):
464:                     adj_matrix = graph_data["adjacency_matrices"][b]
465:                     if adj_matrix.numel() == 0:
466:                         continue
467:
468:                     # Connectivity constraint: encourage reasonable connectivity
469:                     # Penalize isolated rooms (degree 0) and over-connected rooms
470:                     degrees = adj_matrix.sum(dim=1)
471:
472:                     # Isolation penalty (rooms should have at least 1 connection)
473:                     isolation_penalty = torch.exp(-degrees).mean()
474:
475:                     # Over-connection penalty (rooms shouldn't connect to everything)
476:                     max_reasonable_connections = min(4, adj_matrix.shape[0] - 1)

```

```

477:         over_connection_penalty = F.relu(degrees - max_reasonable_connections).mean()
478:
479:         # Graph smoothness (connected rooms should have similar features)
480:         if b < len(graph_data["room_features"]) and graph_data["room_features"][b].numel() > 0:
481:             room_features = graph_data["room_features"][b]
482:             if room_features.shape[0] > 1:
483:                 feature_distances = torch.cdist(room_features, room_features)
484:                 # Weight by adjacency - connected rooms should be similar
485:                 smoothness_loss = (adj_matrix * feature_distances).sum() / (adj_matrix.sum() +
486:                                     else:
487:                     smoothness_loss = torch.tensor(0.0, device=device)
488:             else:
489:                 smoothness_loss = torch.tensor(0.0, device=device)
490:
491:             batch_graph_loss = (0.4 * isolation_penalty +
492:                                 0.3 * over_connection_penalty +
493:                                 0.3 * smoothness_loss)
494:             total_graph_loss = total_graph_loss + batch_graph_loss
495:
496:         return total_graph_loss / batch_size
497:
498:     except Exception as e:
499:         # Fallback to zero loss if graph extraction fails
500:         return torch.tensor(0.0, device=segmentation_logits.device)
501:
502: # ---- Existing helper methods (preserved) ----
503: def _dice_loss(self, pred: torch.Tensor, target: torch.Tensor, smooth: float = 1e-6) -> torch.Tensor:
504:     """Dice loss implementation"""
505:     pred_soft = F.softmax(pred, dim=1)
506:     B = pred_soft.shape[0]
507:     C = pred_soft.shape[1]
508:
509:     dice_losses = []
510:     for c in range(C):
511:         pred_c = pred_soft[:, c, :, :]
512:         target_c = (target == c).float().to(pred_c.device)
513:         intersection = (pred_c * target_c).view(B, -1).sum(dim=1)
514:         union = pred_c.view(B, -1).sum(dim=1) + target_c.view(B, -1).sum(dim=1)
515:         dice = (2.0 * intersection + smooth) / (union + smooth)
516:         dice_losses.append((1.0 - dice).mean())
517:
518:     return torch.stack(dice_losses).mean()
519:
520: def _mask_to_sdf(self, mask: torch.Tensor) -> torch.Tensor:
521:     """Convert mask to SDF with performance warning"""
522:     device = mask.device if torch.is_tensor(mask) else None
523:     if not torch.is_tensor(mask):
524:         mask = torch.tensor(mask, device=device)
525:
526:     B, H, W = mask.shape
527:     sdf = torch.zeros((B, 1, H, W), dtype=torch.float32, device=device)
528:
529:     # FIX: Add performance warning for CV2 bottleneck
530:     if B > 8: # Warn for large batches
531:         print(f"[Performance Warning] SDF conversion with batch_size={B} uses CPU cv2 - consider GPU im")
532:
533:     for b in range(B):
534:         mask_np = mask[b].detach().cpu().numpy().astype(np.uint8) # FIX: explicit detach
535:         try:
536:             dist_inside = cv2.distanceTransform((mask_np > 0).astype(np.uint8), cv2.DIST_L2, 5)
537:             dist_outside = cv2.distanceTransform((mask_np == 0).astype(np.uint8), cv2.DIST_L2, 5)
538:             sdf_np = dist_inside.astype(np.float32) - dist_outside.astype(np.float32)
539:             sdf_np = np.tanh(sdf_np / 10.0).astype(np.float32)
540:             sdf[b, 0] = torch.from_numpy(sdf_np)
541:         except Exception:
542:             # Fallback if cv2 fails
543:             sdf[b, 0] = torch.zeros_like(mask[b].float())
544:
545:     return sdf
546:
547: def _polygon_loss(self, predictions: dict, targets: dict) -> torch.Tensor:
548:     """Polygon/DVX loss (preserved from original)"""
549:     pred_polys = predictions.get("polygons")

```

```

550:         tgt_polys = targets.get("polygons")
551:         valid_mask = targets.get("valid_mask")
552:
553:         if pred_polys is None or tgt_polys is None:
554:             return torch.tensor(0.0, device=pred_polys.device if pred_polys is not None else self.device)
555:
556:         pred_polys = pred_polys.float()
557:         tgt_polys = tgt_polys.float().to(pred_polys.device)
558:
559:         point_loss = self.mse_loss(pred_polys, tgt_polys)
560:
561:         pred_valid = predictions.get("polygon_validity")
562:         if pred_valid is None or valid_mask is None:
563:             validity_loss = torch.tensor(0.0, device=pred_polys.device)
564:         else:
565:             pred_valid = pred_valid.float().to(pred_polys.device)
566:             valid_mask_f = valid_mask.float().to(pred_polys.device)
567:             validity_loss = self.mse_loss(pred_valid, valid_mask_f)
568:
569:         smoothness_loss = self._polygon_smoothness(pred_polys)
570:         rect_loss = self._rectilinearity_loss(pred_polys)
571:
572:         return point_loss + 0.1 * validity_loss + 0.05 * smoothness_loss + 0.1 * rect_loss
573:
574:     def _polygon_smoothness(self, polygons: torch.Tensor) -> torch.Tensor:
575:         """Polygon smoothness loss (preserved)"""
576:         if polygons is None or polygons.numel() == 0:
577:             return torch.tensor(0.0, device=polygons.device if polygons is not None else self.device)
578:
579:         p1 = polygons
580:         p2 = torch.roll(polygons, -1, dims=2)
581:         p3 = torch.roll(polygons, -2, dims=2)
582:         curvature = torch.norm(p1 - 2.0 * p2 + p3, dim=-1)
583:         return curvature.mean()
584:
585:     def _rectilinearity_loss(self, polygons: torch.Tensor) -> torch.Tensor:
586:         """Encourage axis-aligned structure (preserved)"""
587:         if polygons is None or polygons.numel() == 0:
588:             return torch.tensor(0.0, device=polygons.device if polygons is not None else self.device)
589:
590:         edges = torch.roll(polygons, -1, dims=2) - polygons
591:         edge_norms = torch.norm(edges, dim=-1, keepdim=True)
592:         edges_normalized = edges / (edge_norms + 1e-6)
593:
594:         edge1 = edges_normalized
595:         edge2 = torch.roll(edges_normalized, -1, dims=2)
596:
597:         cos_angles = (edge1 * edge2).sum(dim=-1)
598:         cos2 = cos_angles ** 2
599:         perp_penalty = cos2
600:         parallel_penalty = (cos2 - 1.0) ** 2
601:         angle_penalty = torch.minimum(perp_penalty, parallel_penalty)
602:         return angle_penalty.mean()
603:
604:     def _voxel_iou_loss(self, pred_voxels: torch.Tensor, target_voxels: torch.Tensor) -> torch.Tensor:
605:         """3D voxel IoU loss (preserved)"""
606:         pred_prob = torch.sigmoid(torch.clamp(pred_voxels, -10.0, 10.0)) # FIX: safe sigmoid range
607:         target = target_voxels.float().to(pred_prob.device)
608:
609:         intersection = (pred_prob * target).view(pred_prob.shape[0], -1).sum(dim=1)
610:         union = (pred_prob.view(pred_prob.shape[0], -1).sum(dim=1) +
611:                 target.view(target.shape[0], -1).sum(dim=1) - intersection)
612:
613:         iou = (intersection + 1e-6) / (union + 1e-6)
614:         return (1.0 - iou).mean()
615:
616:     def _topology_loss(self, segmentation_logits: torch.Tensor) -> torch.Tensor:
617:         """Traditional topology loss (preserved)"""
618:         seg_soft = F.softmax(segmentation_logits, dim=1)
619:         C = seg_soft.shape[1]
620:         device = seg_soft.device
621:
622:         walls = seg_soft[:, 1] if C > 1 else torch.zeros_like(seg_soft[:, 0])

```

```

623:     doors = seg_soft[:, 2] if C > 2 else torch.zeros_like(walls)
624:     windows = seg_soft[:, 3] if C > 3 else torch.zeros_like(walls)
625:
626:     door_wall_overlap = doors * walls
627:     window_wall_overlap = windows * walls
628:
629:     door_penalty = torch.maximum(doors - door_wall_overlap, torch.zeros_like(doors))
630:     window_penalty = torch.maximum(windows - window_wall_overlap, torch.zeros_like(windows))
631:
632:     connectivity_loss = self._connectivity_loss(walls)
633:
634:     return door_penalty.mean() + window_penalty.mean() + 0.1 * connectivity_loss
635:
636: def _connectivity_loss(self, wall_prob: torch.Tensor) -> torch.Tensor:
637:     """Connectivity loss for walls (preserved)"""
638:     if wall_prob is None or wall_prob.numel() == 0:
639:         return torch.tensor(0.0, device=wall_prob.device if wall_prob is not None else self.device)
640:
641:     kernel = torch.ones((1, 1, 3, 3), device=wall_prob.device, dtype=wall_prob.dtype) / 9.0
642:     neighbors = F.conv2d(wall_prob.unsqueeze(1), kernel, padding=1).squeeze(1)
643:
644:     isolation_penalty = wall_prob * torch.exp(-neighbors)
645:     return isolation_penalty.mean()
646:
647:
648: class LossScheduler:
649:     """Manages curriculum-based loss weight scheduling"""
650:
651:     def __init__(self, config):
652:         self.config = config
653:         self.loss_schedules = config.loss_schedule
654:
655:     def get_scheduled_weights(self, current_stage: int, current_epoch: int,
656:                               stage_epoch: int, total_stage_epochs: int,
657:                               base_weights: Dict[str, float]) -> Dict[str, float]:
658:         """
659:         Calculate loss weights based on curriculum schedule
660:
661:         Args:
662:             current_stage: Current training stage (1, 2, 3)
663:             current_epoch: Global epoch count
664:             stage_epoch: Epoch within current stage
665:             total_stage_epochs: Total epochs planned for current stage
666:             base_weights: Base weight configuration
667:         """
668:         scheduled_weights = base_weights.copy()
669:
670:         for loss_name, schedule_type in self.loss_schedules.items():
671:             if loss_name not in scheduled_weights:
672:                 continue
673:
674:             base_weight = scheduled_weights[loss_name]
675:
676:             if schedule_type == "static":
677:                 # Keep original weight
678:                 continue
679:
680:             elif schedule_type == "progressive":
681:                 # Gradually increase throughout training
682:                 if loss_name == "topology":
683:                     start_weight = self.config.topology_start_weight
684:                     end_weight = self.config.topology_end_weight
685:                     ramp_epochs = self.config.topology_ramp_epochs
686:                     progress = min(current_epoch / ramp_epochs, 1.0)
687:                     scheduled_weights[loss_name] = start_weight + progress * (end_weight - start_weight)
688:
689:             elif schedule_type == "linear_ramp":
690:                 # Linear increase within current stage
691:                 progress = stage_epoch / max(total_stage_epochs, 1)
692:                 scheduled_weights[loss_name] = base_weight * progress
693:
694:             elif schedule_type == "exponential":
695:                 # Exponential increase

```

```

696:         progress = stage_epoch / max(total_stage_epochs, 1)
697:         scheduled_weights[loss_name] = base_weight * (progress ** 2)
698:
699:     elif schedule_type == "early_decay":
700:         # Decay after Stage 1 (for SDF loss)
701:         if current_stage > 1:
702:             scheduled_weights[loss_name] = base_weight * 0.3
703:
704:     elif schedule_type == "staged_ramp":
705:         # Ramp up during specific stage (polygon in Stage 2)
706:         if current_stage == 2:
707:             progress = stage_epoch / max(total_stage_epochs, 1)
708:             scheduled_weights[loss_name] = base_weight * progress
709:         elif current_stage < 2:
710:             scheduled_weights[loss_name] = 0.0
711:
712:     elif schedule_type == "late_ramp":
713:         # Ramp up in Stage 3 (voxel loss)
714:         if current_stage == 3:
715:             progress = stage_epoch / max(total_stage_epochs, 1)
716:             scheduled_weights[loss_name] = base_weight * progress
717:         elif current_stage < 3:
718:             scheduled_weights[loss_name] = 0.0
719:
720:     elif schedule_type == "mid_ramp":
721:         # Activate mid-training (latent consistency)
722:         if current_stage >= 2:
723:             if current_stage == 2:
724:                 progress = min(stage_epoch / (total_stage_epochs * 0.5), 1.0)
725:                 scheduled_weights[loss_name] = base_weight * progress
726:             else: # Stage 3
727:                 scheduled_weights[loss_name] = base_weight
728:         else:
729:             scheduled_weights[loss_name] = 0.0
730:
731:     elif schedule_type == "delayed_ramp":
732:         # FIX: gentler ramp for graph constraints
733:         if current_epoch >= self.config.graph_start_epoch:
734:             epochs_since_start = current_epoch - self.config.graph_start_epoch
735:             ramp_duration = 50 # FIX: slower ramp (was 20)
736:             progress = min(epochs_since_start / ramp_duration, 1.0)
737:             scheduled_weights[loss_name] = self.config.graph_end_weight * progress
738:         else:
739:             scheduled_weights[loss_name] = 0.0
740:
741:     return scheduled_weights

```

File: training\trainer.py

```

=====
1: """
2: Dynamic Multi-stage training system with adaptive curriculum learning
3: Implements novel training strategies: dynamic stage transitions, topology-aware scheduling,
4: multi-objective optimization, and cross-modal consistency learning
5: """
6:
7: import torch
8: import torch.nn.utils
9:
10: # training/trainer.py - Fixed AMP imports
11: from torch.amp import autocast, GradScaler
12: import time
13: import numpy as np
14: import random
15: from pathlib import Path
16: from tqdm import tqdm
17: from typing import Dict, List, Optional, Tuple
18: from collections import deque
19:
20: from .losses import ResearchGradeLoss, LossScheduler
21: from config import DEFAULT_TRAINING_CONFIG, DEFAULT_LOSS_CONFIG, StageTransitionCriteria
22:

```



```

23:
24: class CurriculumState:
25:     """Tracks curriculum learning state and metrics"""
26:
27:     def __init__(self, config):
28:         self.config = config
29:
30:         # Loss history for plateau detection
31:         self.loss_history = {
32:             "stage1": deque(maxlen=config.plateau_detection_window * 2),
33:             "stage2": deque(maxlen=config.plateau_detection_window * 2),
34:             "stage3": deque(maxlen=config.plateau_detection_window * 2),
35:         }
36:
37:         # Component loss tracking
38:         self.component_losses = {
39:             "segmentation": deque(maxlen=20),
40:             "dice": deque(maxlen=20),
41:             "polygon": deque(maxlen=20),
42:             "voxel": deque(maxlen=20),
43:             "topology": deque(maxlen=20),
44:             "latent_consistency": deque(maxlen=20),
45:             "graph": deque(maxlen=20),
46:         }
47:
48:         # Gradient magnitude tracking for dynamic weighting
49:         self.gradient_norms = {
50:             name: deque(maxlen=config.gradient_norm_window)
51:             for name in self.component_losses.keys()
52:         }
53:
54:         # Stage transition tracking
55:         self.epochs_without_improvement = 0
56:         self.best_val_loss = float("inf")
57:         self.stage_transition_epochs = []
58:
59:         # Dynamic weights history
60:         self.weight_history = []
61:
62:     def update_loss_history(self, stage: str, val_loss: float):
63:         """Update validation loss history for plateau detection"""
64:         if stage in self.loss_history:
65:             self.loss_history[stage].append(val_loss)
66:
67:         # Update improvement tracking
68:         if val_loss < self.best_val_loss:
69:             self.best_val_loss = val_loss
70:             self.epochs_without_improvement = 0
71:         else:
72:             self.epochs_without_improvement += 1
73:
74:     def update_component_losses(self, loss_components: Dict[str, float]):
75:         """Update individual loss component history"""
76:         for name, loss_val in loss_components.items():
77:             if name in self.component_losses:
78:                 self.component_losses[name].append(loss_val)
79:
80:     def should_transition(self, current_stage: int) -> bool:
81:         """Check if should transition to next stage"""
82:         if current_stage == 1:
83:             val_losses = list(self.loss_history["stage1"])
84:             return StageTransitionCriteria.should_transition_from_stage1(
85:                 [], val_losses, self.config
86:             )
87:         elif current_stage == 2:
88:             polygon_losses = list(self.component_losses["polygon"])
89:             return StageTransitionCriteria.should_transition_from_stage2(
90:                 polygon_losses, self.config
91:             )
92:
93:         return False
94:
95:

```

```

96: class AdaptiveMultiStageTrainer:
97:     """
98:     Advanced multi-stage trainer with dynamic curriculum learning:
99:     - Adaptive stage transitioning based on performance plateaus
100:    - Topology-aware loss scheduling
101:    - Multi-objective optimization with dynamic weighting
102:    - Cross-modal latent consistency learning
103:    - Graph-based topology constraints
104:    """
105:
106:    # Class constant for rolling checkpoint path
107:    ROLLING_CHECKPOINT = "latest_checkpoint.pth"
108:
109:    def __init__(self, model, train_loader, val_loader, device=None, config=None):
110:        if config is None:
111:            config = DEFAULT_TRAINING_CONFIG
112:
113:        self.model = model.to(device or config.device)
114:        self.train_loader = train_loader
115:        self.val_loader = val_loader
116:        self.device = device or config.device
117:        self.config = config
118:
119:        # Initialize curriculum state
120:        self.curriculum_state = CurriculumState(config.curriculum)
121:        self.loss_scheduler = LossScheduler(config.curriculum)
122:
123:        # Training state tracking for resume functionality
124:        self.current_stage = 1
125:        self.current_epoch = 0
126:        self.global_epoch = 0
127:        self.stage_epoch = 0
128:        self.stage_start_time = None
129:        self.epoch_times = []
130:
131:        # Add AMP and optimization settings - Updated for new PyTorch API
132:        self.use_amp = getattr(config, "use_mixed_precision", True)
133:        self.scaler = GradScaler("cuda", enabled=self.use_amp)
134:        self.accumulation_steps = getattr(config, "accumulation_steps", 1)
135:        self.dvx_step_freq = getattr(config, "dvx_step_freq", 1)
136:        self.voxel_size_stage = getattr(config, "voxel_size_stage", None)
137:        self.image_size_stage = getattr(config, "image_size_stage", None)
138:        self._step = 0
139:
140:        # Enhanced optimizers with better hyperparameters
141:        self.optimizer_2d = torch.optim.AdamW(
142:            list(self.model.encoder.parameters())
143:            + list(self.model.seg_head.parameters())
144:            + list(self.model.attr_head.parameters())
145:            + list(self.model.sdf_head.parameters()),
146:            lr=config.stage1_lr,
147:            weight_decay=config.stage1_weight_decay,
148:            betas=(0.9, 0.999),
149:        )
150:
151:        self.optimizer_dvx = torch.optim.AdamW(
152:            self.model.dvx.parameters(),
153:            lr=config.stage2_lr,
154:            weight_decay=config.stage2_weight_decay,
155:            betas=(0.9, 0.999),
156:        )
157:
158:        self.optimizer_full = torch.optim.AdamW(
159:            self.model.parameters(),
160:            lr=config.stage3_lr,
161:            weight_decay=config.stage3_weight_decay,
162:            betas=(0.9, 0.999),
163:        )
164:
165:        # Enhanced learning rate schedulers
166:        if config.use_cosine_restarts:
167:            self.scheduler_2d = torch.optim.lr_scheduler.CosineAnnealingWarmRestarts(
168:                self.optimizer_2d, T_0=20, T_mult=1

```

```

169:         )
170:         self.scheduler_dvx = torch.optim.lr_scheduler.CosineAnnealingWarmRestarts(
171:             self.optimizer_dvx, T_0=15, T_mult=1
172:         )
173:         self.scheduler_full = torch.optim.lr_scheduler.CosineAnnealingWarmRestarts(
174:             self.optimizer_full, T_0=30, T_mult=1
175:         )
176:     else:
177:         self.scheduler_2d = torch.optim.lr_scheduler.CosineAnnealingLR(
178:             self.optimizer_2d, T_max=config.max_stage1_epochs
179:         )
180:         self.scheduler_dvx = torch.optim.lr_scheduler.CosineAnnealingLR(
181:             self.optimizer_dvx, T_max=config.max_stage2_epochs
182:         )
183:         self.scheduler_full = torch.optim.lr_scheduler.CosineAnnealingLR(
184:             self.optimizer_full, T_max=config.max_stage3_epochs
185:         )
186:
187:     # Enhanced loss function with dynamic weighting
188:     base_loss_kwargs = {
189:         k: v
190:         for k, v in DEFAULT_LOSS_CONFIG.__dict__.items()
191:         if k != "enable_dynamic_weighting"
192:     }
193:     self.loss_fn = ResearchGradeLoss(
194:         **base_loss_kwargs,
195:         enable_dynamic_weighting=bool(config.curriculum.use_gradnorm),
196:         gradnorm_alpha=float(config.curriculum.gradnorm_alpha),
197:         device=self.device,
198:     )
199:
200:     self.history = {
201:         "stage1": {"train_loss": [], "val_loss": [], "component_losses": []},
202:         "stage2": {"train_loss": [], "val_loss": [], "component_losses": []},
203:         "stage3": {"train_loss": [], "val_loss": [], "component_losses": []},
204:         "stage_transitions": [],
205:         "dynamic_weights": [],
206:         "curriculum_events": [],
207:     }
208:
209:     def _get_eta_string(self, epoch, total_epochs):
210:         """Calculate and format ETA string"""
211:         if len(self.epoch_times) == 0:
212:             return "ETA: calculating..."
213:
214:         avg_epoch_time = sum(self.epoch_times) / len(self.epoch_times)
215:         remaining_epochs = total_epochs - epoch - 1
216:         eta_seconds = avg_epoch_time * remaining_epochs
217:
218:         if eta_seconds < 60:
219:             return f"ETA: {int(eta_seconds)}s"
220:         elif eta_seconds < 3600:
221:             return f"ETA: {int(eta_seconds // 60)}m {int(eta_seconds % 60)}s"
222:         else:
223:             hours = int(eta_seconds // 3600)
224:             minutes = int((eta_seconds % 3600) // 60)
225:             return f"ETA: {hours}h {minutes}m"
226:
227:     def _get_shared_parameters(self):
228:         """Get shared parameters for GradNorm weighting"""
229:         # Return encoder parameters as shared across tasks
230:         return list(self.model.encoder.parameters())
231:
232:     def _update_loss_weights_for_curriculum(
233:         self, current_stage: int, stage_epoch: int, total_stage_epochs: int
234:     ):
235:         """Update loss weights based on curriculum schedule"""
236:         base_weights = {
237:             "seg": self.loss_fn.initial_weights["seg"],
238:             "dice": self.loss_fn.initial_weights["dice"],
239:             "sdf": self.loss_fn.initial_weights["sdf"],
240:             "attr": self.loss_fn.initial_weights["attr"],
241:             "polygon": self.loss_fn.initial_weights["polygon"],

```

```

242:         "voxel": self.loss_fn.initial_weights["voxel"],
243:         "topology": self.loss_fn.initial_weights["topology"],
244:         "latent_consistency": self.loss_fn.initial_weights["latent_consistency"],
245:         "graph": self.loss_fn.initial_weights["graph"],
246:     }
247:
248:     scheduled_weights = self.loss_scheduler.get_scheduled_weights(
249:         current_stage,
250:         self.global_epoch,
251:         stage_epoch,
252:         total_stage_epochs,
253:         base_weights,
254:     )
255:
256:     self.loss_fn.update_loss_weights(scheduled_weights)
257:
258:     # Log weight changes
259:     self.history["dynamic_weights"].append(
260:         {
261:             "epoch": self.global_epoch,
262:             "stage": current_stage,
263:             "weights": scheduled_weights.copy(),
264:         }
265:     )
266:
267: def _train_epoch(self, mode="stage1"):
268:     """Enhanced training epoch with AMP, gradient accumulation, and DVX gating"""
269:     self.model.train()
270:     total_loss = 0
271:     component_loss_sums = {}
272:
273:     # Select appropriate optimizer based on mode
274:     if mode == "stage1":
275:         optimizer = self.optimizer_2d
276:     elif mode == "stage2":
277:         optimizer = self.optimizer_dvx
278:     else: # stage3
279:         optimizer = self.optimizer_full
280:
281:     # Progress bar for training batches
282:     train_pbar = tqdm(
283:         self.train_loader, desc=f"Training {mode.upper()}", leave=False, ncols=120
284:     )
285:
286:     batch_count = 0
287:     epoch_start_time = time.time()
288:
289:     for batch_idx, batch in enumerate(train_pbar):
290:         self._step += 1
291:         batch = {
292:             k: v.to(self.device, non_blocking=True) if torch.is_tensor(v) else v
293:             for k, v in batch.items()
294:         }
295:
296:         # Gate heavy DVX/extrusion: only run full forward every dvx_step_freq steps
297:         run_full_geometric = (self.dvx_step_freq <= 1) or (
298:             (self._step % self.dvx_step_freq) == 0
299:         )
300:
301:         # First-batch profiling (optional timing helper)
302:         if batch_idx == 0 and self.global_epoch == 0:
303:             torch.cuda.synchronize()
304:             t0 = time.time()
305:             with autocast("cuda", enabled=self.use_amp):
306:                 out = self.model(batch["image"], run_full_geometric=True)
307:                 # Prepare targets for loss computation
308:                 targets = self._prepare_targets(batch, mode)
309:                 shared_params = (
310:                     self._get_shared_parameters()
311:                     if self.config.curriculum.use_gradnorm
312:                     else None
313:                 )
314:                 l, _ = self.loss_fn(

```

```

315:         out, targets, shared_params, run_full_geometric=True
316:     )
317:     torch.cuda.synchronize()
318:     print(f"First-batch forward+loss time: {time.time() - t0:.3f}s")
319:
320:     with autocast("cuda", enabled=self.use_amp):
321:         # Forward pass with geometric gating
322:         predictions = self.model(
323:             batch["image"], run_full_geometric=run_full_geometric
324:         )
325:
326:         # Add latent embeddings if model supports it
327:         if hasattr(self.model, "get_latent_embeddings"):
328:             latent_2d, latent_3d = self.model.get_latent_embeddings(
329:                 batch["image"]
330:             )
331:             predictions["latent_2d_embedding"] = latent_2d
332:             predictions["latent_3d_embedding"] = latent_3d
333:
334:         # Prepare targets based on training mode
335:         targets = self._prepare_targets(batch, mode)
336:
337:         # Get shared parameters for dynamic weighting
338:         shared_params = (
339:             self._get_shared_parameters()
340:             if self.config.curriculum.use_gradnorm
341:             else None
342:         )
343:
344:         # Compute loss with dynamic weighting and geometric gating
345:         loss, loss_components = self.loss_fn(
346:             predictions,
347:             targets,
348:             shared_params,
349:             run_full_geometric=run_full_geometric,
350:         )
351:
352:         # Scale loss for gradient accumulation
353:         loss = loss / self.accumulation_steps
354:
355:     # Scale and backward pass
356:     self.scaler.scale(loss).backward()
357:
358:     # Gradient accumulation step
359:     if ((batch_idx + 1) % self.accumulation_steps) == 0:
360:         # Unscale and clip gradients
361:         self.scaler.unscale_(optimizer)
362:
363:         # Apply gradient clipping
364:         torch.nn.utils.clip_grad_norm_(
365:             self.model.parameters(), self.config.grad_clip_norm
366:         )
367:
368:         # Optimizer step with scaler
369:         self.scaler.step(optimizer)
370:         self.scaler.update()
371:         optimizer.zero_grad()
372:
373:     current_loss = loss.item() * self.accumulation_steps
374:     total_loss += current_loss
375:
376:     # Track component losses
377:     for name, component_loss in loss_components.items():
378:         if name != "total":
379:             loss_val = (
380:                 component_loss.item()
381:                 if torch.is_tensor(component_loss)
382:                 else component_loss
383:             )
384:             if name not in component_loss_sums:
385:                 component_loss_sums[name] = 0
386:             component_loss_sums[name] += loss_val
387:

```

```

388:         batch_count += 1
389:
390:         # Occasional lightweight logging
391:         if (batch_idx + 1) % 50 == 0:
392:             elapsed = time.time() - epoch_start_time
393:             avg_time_per_batch = elapsed / (batch_idx + 1)
394:             current_weights = {
395:                 k: f"{v:.3f}" for k, v in self.loss_fn.weights.items() if v > 0.001
396:             }
397:             print(
398:                 f"[Epoch {self.global_epoch}] Batch {batch_idx+1}/{len(self.train_loader)} | "
399:                 f"avg batch {avg_time_per_batch:.3f}s | loss {total_loss/batch_count:.4f}"
400:             )
401:
402:         # Update progress bar
403:         current_weights = {
404:             k: f"{v:.3f}" for k, v in self.loss_fn.weights.items() if v > 0.001
405:         }
406:         train_pbar.set_postfix(
407:             {
408:                 "loss": f"{current_loss:.4f}",
409:                 "weights": str(current_weights)[:50] + "...",
410:                 if len(str(current_weights)) > 50
411:                 else str(current_weights),
412:             }
413:         )
414:
415:         # Final epoch timing
416:         epoch_time = time.time() - epoch_start_time
417:         avg_loss = total_loss / batch_count
418:         print(
419:             f"Epoch {self.global_epoch} finished in {epoch_time/60:.2f} min. avg loss: {avg_loss:.4f}"
420:         )
421:
422:         # Average component losses
423:         avg_component_losses = {
424:             name: loss_sum / batch_count
425:             for name, loss_sum in component_loss_sums.items()
426:         }
427:
428:         return avg_loss, avg_component_losses
429:
430:     def _prepare_targets(self, batch, mode):
431:         """Prepare targets based on training mode"""
432:         if mode == "stage1":
433:             return {"mask": batch["mask"], "attributes": batch["attributes"]}
434:         elif mode == "stage2":
435:             return {
436:                 "polygons_gt": {
437:                     "polygons": batch["polygons_gt"]["polygons"].to(self.device),
438:                     "valid_mask": batch["polygons_gt"]["valid_mask"].to(self.device),
439:                 }
440:             }
441:         else: # stage3
442:             return {
443:                 "mask": batch["mask"],
444:                 "attributes": batch["attributes"],
445:                 "voxels_gt": batch["voxels_gt"],
446:                 "polygons_gt": {
447:                     "polygons": batch["polygons_gt"]["polygons"].to(self.device),
448:                     "valid_mask": batch["polygons_gt"]["valid_mask"].to(self.device),
449:                 },
450:             }
451:
452:     def _validate(self, mode="stage1"):
453:         """Enhanced validation with detailed metrics and AMP support"""
454:         self.model.eval()
455:         total_loss = 0
456:         component_loss_sums = {}
457:
458:         val_pbar = tqdm(
459:             self.val_loader, desc=f"Validating {mode.upper()}", leave=False, ncols=120
460:         )

```

```

461:
462:     batch_count = 0
463:     with torch.no_grad():
464:         for batch in val_pbar:
465:             batch = {
466:                 k: v.to(self.device, non_blocking=True) if torch.is_tensor(v) else v
467:                 for k, v in batch.items()
468:             }
469:
470:             with autocast("cuda", enabled=self.use_amp):
471:                 # Always run full geometric computation during validation
472:                 predictions = self.model(batch["image"], run_full_geometric=True)
473:
474:                 # Add latent embeddings if available
475:                 if hasattr(self.model, "get_latent_embeddings"):
476:                     latent_2d, latent_3d = self.model.get_latent_embeddings(
477:                         batch["image"]
478:                     )
479:                     predictions["latent_2d_embedding"] = latent_2d
480:                     predictions["latent_3d_embedding"] = latent_3d
481:
482:                 targets = self._prepare_targets(batch, mode)
483:
484:                 loss, loss_components = self.loss_fn(
485:                     predictions, targets, run_full_geometric=True
486:                 )
487:
488:                 current_loss = loss.item()
489:                 total_loss += current_loss
490:
491:                 # Track component losses
492:                 for name, component_loss in loss_components.items():
493:                     if name != "total":
494:                         loss_val = (
495:                             component_loss.item()
496:                             if torch.is_tensor(component_loss)
497:                             else component_loss
498:                         )
499:                         if name not in component_loss_sums:
500:                             component_loss_sums[name] = 0
501:                         component_loss_sums[name] += loss_val
502:
503:                 batch_count += 1
504:                 val_pbar.set_postfix({"loss": f"{current_loss:.4f}"})
505:
506:     avg_component_losses = {
507:         name: loss_sum / batch_count
508:         for name, loss_sum in component_loss_sums.items()
509:     }
510:
511:     return total_loss / batch_count, avg_component_losses
512:
513: def train_stage_adaptive(self, stage: int, max_epochs: int, min_epochs: int):
514:     """
515:     Train a stage with adaptive termination based on curriculum learning
516:
517:     Args:
518:         stage: Stage number (1, 2, 3)
519:         max_epochs: Maximum epochs for this stage
520:         min_epochs: Minimum epochs before considering transition
521:     """
522:     print("=" * 60)
523:     print(f"STAGE {stage}: Adaptive Training with Dynamic Curriculum")
524:     print("=" * 60)
525:
526:     self.current_stage = stage
527:     self.stage_start_time = time.time()
528:
529:     # Only reset if not resuming
530:     if not hasattr(self, "epoch_times") or self.epoch_times is None:
531:         self.epoch_times = []
532:
533:     start_epoch = int(self.stage_epoch or 0)

```



```

607:         "reason": "performance_plateau",
608:     }
609: )
610:
611:     self.history["curriculum_events"].append(
612:         {
613:             "type": "stage_transition",
614:             "stage": stage,
615:             "epoch": self.global_epoch,
616:             "details": f"Converged after {epoch+1} epochs",
617:         }
618:     )
619:     break
620:
621:     # Save rolling checkpoint
622:     if (epoch + 1) % self.config.checkpoint_freq == 0:
623:         self._save_rolling_checkpoint()
624:
625:     print(f"\nStage {stage} completed after {epoch+1} epochs")
626:
627: def _configure_stage_parameters(self, stage: int):
628:     """Configure which parameters require gradients for each stage"""
629:     # First freeze everything
630:     for param in self.model.parameters():
631:         param.requires_grad = False
632:
633:     if stage == 1:
634:         # Stage 1: Segmentation + Attributes (2D only)
635:         for param in self.model.encoder.parameters():
636:             param.requires_grad = True
637:         for param in self.model.seg_head.parameters():
638:             param.requires_grad = True
639:         for param in self.model.attr_head.parameters():
640:             param.requires_grad = True
641:         for param in self.model.sdf_head.parameters():
642:             param.requires_grad = True
643:
644:     elif stage == 2:
645:         # Stage 2: DVX training (polygon fitting) - keep encoder frozen initially
646:         for param in self.model.dvx.parameters():
647:             param.requires_grad = True
648:         # Optionally unfreeze encoder in later epochs
649:         if self.stage_epoch > 10:
650:             for param in self.model.encoder.parameters():
651:                 param.requires_grad = True
652:
653:     else: # stage == 3
654:         # Stage 3: End-to-end fine-tuning (all parameters)
655:         for param in self.model.parameters():
656:             param.requires_grad = True
657:
658: def _display_epoch_results(
659:     self,
660:     epoch: int,
661:     total_epochs: int,
662:     train_loss: float,
663:     val_loss: float,
664:     train_components: Dict,
665:     val_components: Dict,
666:     epoch_time: float,
667: ):
668:     """Display comprehensive epoch results with curriculum information"""
669:     eta_str = self._get_eta_string(epoch, total_epochs)
670:
671:     print(f"Train Loss: {train_loss:.4f}, Val Loss: {val_loss:.4f}")
672:     print(f"Epoch time: {epoch_time:.1f}s, {eta_str}")
673:
674:     # Show significant component losses
675:     significant_components = {
676:         k: v
677:         for k, v in val_components.items()
678:         if v > 0.01
679:         and k

```

```

680:         in [
681:             "seg",
682:             "dice",
683:             "polygon",
684:             "voxel",
685:             "topology",
686:             "latent_consistency",
687:             "graph",
688:         ]
689:     }
690:     if significant_components:
691:         comp_str = ", ".join(
692:             [f"{k}: {v:.3f}" for k, v in significant_components.items()]
693:         )
694:         print(f"Components: {comp_str}")
695:
696:     # Show current loss weights for active components
697:     active_weights = {k: v for k, v in self.loss_fn.weights.items() if v > 0.001}
698:     if active_weights:
699:         weight_str = ", ".join([f"{k}: {v:.3f}" for k, v in active_weights.items()])
700:         print(f"Weights: {weight_str}")
701:
702:     # Show curriculum status
703:     plateau_epochs = self.curriculum_state.epochs_without_improvement
704:     if plateau_epochs > 0:
705:         print(f"Plateau: {plateau_epochs} epochs without improvement")
706:
707:     def _save_rolling_checkpoint(self):
708:         """Enhanced checkpoint saving with curriculum state, RNG state, and scaler state"""
709:         checkpoint = {
710:             "model_state_dict": self.model.state_dict(),
711:             "optimizer_2d_state_dict": self.optimizer_2d.state_dict(),
712:             "optimizer_dvx_state_dict": self.optimizer_dvx.state_dict(),
713:             "optimizer_full_state_dict": self.optimizer_full.state_dict(),
714:             "scheduler_2d_state_dict": self.scheduler_2d.state_dict(),
715:             "scheduler_dvx_state_dict": self.scheduler_dvx.state_dict(),
716:             "scheduler_full_state_dict": self.scheduler_full.state_dict(),
717:             "scaler_state_dict": self.scaler.state_dict(), # Add AMP scaler state
718:             "loss_fn_state": {
719:                 "weights": self.loss_fn.weights,
720:                 "initial_weights": self.loss_fn.initial_weights,
721:             },
722:             "history": self.history,
723:             "config": self.config,
724:             "current_stage": self.current_stage,
725:             "current_epoch": self.current_epoch,
726:             "global_epoch": self.global_epoch,
727:             "stage_epoch": self.stage_epoch,
728:             "epoch_times": self.epoch_times,
729:             "step_counter": self._step, # Save step counter for DVX gating
730:             "curriculum_state": {
731:                 "loss_history": dict(self.curriculum_state.loss_history),
732:                 "component_losses": dict(self.curriculum_state.component_losses),
733:                 "epochs_without_improvement": self.curriculum_state.epochs_without_improvement,
734:                 "best_val_loss": self.curriculum_state.best_val_loss,
735:                 "stage_transition_epochs": self.curriculum_state.stage_transition_epochs,
736:             },
737:             "rng_state": {
738:                 "torch": torch.get_rng_state(),
739:                 "cuda": torch.cuda.get_rng_state_all()
740:                 if torch.cuda.is_available()
741:                 else None,
742:                 "numpy": np.random.get_state(),
743:                 "python": random.getstate(),
744:             },
745:         }
746:
747:         checkpoint_path = self.ROLLING_CHECKPOINT
748:         torch.save(checkpoint, checkpoint_path)
749:         print(f"Rolling checkpoint saved: {checkpoint_path}")
750:
751:     def load_checkpoint(self, filename):
752:         """Enhanced checkpoint loading with curriculum state restoration and device handling"""

```

```

753:         checkpoint = torch.load(filename, map_location=self.device)
754:
755:         self.model.load_state_dict(checkpoint["model_state_dict"])
756:         self.optimizer_2d.load_state_dict(checkpoint["optimizer_2d_state_dict"])
757:         self.optimizer_dvx.load_state_dict(checkpoint["optimizer_dvx_state_dict"])
758:         self.optimizer_full.load_state_dict(checkpoint["optimizer_full_state_dict"])
759:
760:         # Load scaler state for AMP
761:         if "scaler_state_dict" in checkpoint:
762:             self.scaler.load_state_dict(checkpoint["scaler_state_dict"])
763:
764:         # Safer scheduler loading
765:         for sched_key, sched_obj in [
766:             ("scheduler_2d_state_dict", self.scheduler_2d),
767:             ("scheduler_dvx_state_dict", self.scheduler_dvx),
768:             ("scheduler_full_state_dict", self.scheduler_full),
769:         ]:
770:             if sched_key in checkpoint:
771:                 sched_obj.load_state_dict(checkpoint[sched_key])
772:
773:         # Load loss weights with proper device handling
774:         if "loss_fn_state" in checkpoint:
775:             loaded_weights = checkpoint["loss_fn_state"]["weights"]
776:             if isinstance(loaded_weights, dict):
777:                 self.loss_fn.weights = {
778:                     k: (v.to(self.device) if torch.is_tensor(v) else v)
779:                     for k, v in loaded_weights.items()
780:                 }
781:             else:
782:                 self.loss_fn.weights = loaded_weights
783:             self.loss_fn.initial_weights = checkpoint["loss_fn_state"][
784:                 "initial_weights"
785:             ]
786:
787:         if "history" in checkpoint:
788:             self.history = checkpoint["history"]
789:
790:         # Restore training state
791:         if "current_stage" in checkpoint:
792:             self.current_stage = checkpoint["current_stage"]
793:         if "current_epoch" in checkpoint:
794:             self.current_epoch = checkpoint["current_epoch"]
795:         if "global_epoch" in checkpoint:
796:             self.global_epoch = checkpoint["global_epoch"]
797:         if "stage_epoch" in checkpoint:
798:             self.stage_epoch = checkpoint["stage_epoch"]
799:         if "epoch_times" in checkpoint:
800:             self.epoch_times = checkpoint["epoch_times"]
801:         if "step_counter" in checkpoint:
802:             self._step = checkpoint["step_counter"]
803:
804:         # Restore curriculum state
805:         if "curriculum_state" in checkpoint:
806:             cs = checkpoint["curriculum_state"]
807:             for key, history in cs["loss_history"].items():
808:                 self.curriculum_state.loss_history[key] = deque(
809:                     history, maxlen=self.config.curriculum.plateau_detection_window * 2
810:                 )
811:             for key, history in cs["component_losses"].items():
812:                 self.curriculum_state.component_losses[key] = deque(history, maxlen=20)
813:             self.curriculum_state.epochs_without_improvement = cs.get(
814:                 "epochs_without_improvement", 0
815:             )
816:             self.curriculum_state.best_val_loss = cs.get("best_val_loss", float("inf"))
817:             self.curriculum_state.stage_transition_epochs = cs.get(
818:                 "stage_transition_epochs", []
819:             )
820:
821:         # Restore RNG states
822:         if "rng_state" in checkpoint:
823:             rs = checkpoint["rng_state"]
824:
825:         # --- Torch RNG (CPU) ---

```

```

826:         try:
827:             torch_state = rs.get("torch", None)
828:             if torch_state is not None:
829:                 # If it's already a torch tensor with uint8 dtype, use directly
830:                 if torch.is_tensor(torch_state) and torch_state.dtype == torch.uint8:
831:                     torch.set_rng_state(torch_state)
832:                 else:
833:                     # Convert lists / numpy arrays / other tensors to uint8 torch tensor
834:                     torch.set_rng_state(torch.tensor(torch_state, dtype=torch.uint8))
835:             except Exception as e:
836:                 print(f"Warning: could not restore torch RNG state ({e}), skipping.")
837:
838:         # --- CUDA RNG (all devices) ---
839:         try:
840:             cuda_state = rs.get("cuda", None)
841:             if cuda_state is not None and torch.cuda.is_available():
842:                 # cuda_state might be a list of states (one per device)
843:                 cuda_tensors = []
844:                 for s in cuda_state:
845:                     if torch.is_tensor(s) and s.dtype == torch.uint8:
846:                         cuda_tensors.append(s)
847:                     else:
848:                         cuda_tensors.append(torch.tensor(s, dtype=torch.uint8))
849:                 torch.cuda.set_rng_state_all(cuda_tensors)
850:             except Exception as e:
851:                 print(f"Warning: could not restore CUDA RNG state ({e}), skipping.")
852:
853:         # --- numpy RNG ---
854:         try:
855:             if "numpy" in rs and rs["numpy"] is not None:
856:                 np.random.set_state(rs["numpy"])
857:             except Exception as e:
858:                 print(f"Warning: could not restore numpy RNG state ({e}), skipping.")
859:
860:         # --- python random RNG ---
861:         try:
862:             if "python" in rs and rs["python"] is not None:
863:                 random.setstate(rs["python"])
864:             except Exception as e:
865:                 print(f"Warning: could not restore python RNG state ({e}), skipping.")
866:
867:         # Restore DataLoader sampler states if available
868:         if "dataloader_state" in checkpoint:
869:             dl_state = checkpoint["dataloader_state"]
870:             if dl_state["train_sampler_state"] is not None and hasattr(
871:                 self.train_loader.sampler, "__dict__"
872:             ):
873:                 try:
874:                     self.train_loader.sampler.__dict__.update(
875:                         dl_state["train_sampler_state"]
876:                     )
877:                 except Exception:
878:                     print("Warning: Could not restore train_loader sampler state")
879:             if dl_state["val_sampler_state"] is not None and hasattr(
880:                 self.val_loader.sampler, "__dict__"
881:             ):
882:                 try:
883:                     self.val_loader.sampler.__dict__.update(
884:                         dl_state["val_sampler_state"]
885:                     )
886:                 except Exception:
887:                     print("Warning: Could not restore val_loader sampler state")
888:
889:         print(f"Checkpoint loaded: {filename}")
890:         print(
891:             f"Resuming from Stage {self.current_stage}, Global Epoch {self.global_epoch}"
892:         )
893:         print(
894:             f"Curriculum state restored with {self.curriculum_state.epochs_without_improvement} epochs with"
895:         )
896:
897:     def train_all_stages(self):
898:         """

```

```

899:         Run complete adaptive multi-stage training pipeline
900:
901:         This is the main entry point that orchestrates the dynamic curriculum learning
902:         """
903:         if Path(self.ROLLING_CHECKPOINT).exists():
904:             print(f"Found existing checkpoint: {self.ROLLING_CHECKPOINT}")
905:             print("Resuming adaptive training from checkpoint...")
906:             self.load_checkpoint(self.ROLLING_CHECKPOINT)
907:         else:
908:             print("Starting fresh adaptive training pipeline...")
909:             self.current_stage = 1
910:             self.current_epoch = 0
911:             self.global_epoch = 0
912:
913:         print("\n" + "=" * 80)
914:         print("ADAPTIVE MULTI-STAGE TRAINING WITH DYNAMIC CURRICULUM")
915:         print("Novel Training Strategies:")
916:         print("? Adaptive Stage Transitioning (Dynamic Curriculum)")
917:         print("? Topology-aware Loss Scheduling")
918:         print("? Multi-objective Optimization with Dynamic Weighting")
919:         print("? Cross-modal Latent Consistency Learning")
920:         print("? Graph-based Topology Constraints")
921:         print("=" * 80)
922:
923:         # Stage 1: Adaptive 2D training
924:         if self.current_stage <= 1:
925:             print("\n? STAGE 1: Adaptive 2D Segmentation + Attributes Training")
926:             self.train_stage_adaptive(
927:                 stage=1,
928:                 max_epochs=self.config.max_stagel_epochs,
929:                 min_epochs=self.config.min_stagel_epochs,
930:             )
931:             self.current_stage = 2
932:             self.stage_epoch = 0
933:             print("\nStage 1 completed. Transitioning to Stage 2...")
934:
935:         # Stage 2: Adaptive DVX training
936:         if self.current_stage <= 2:
937:             print("\n? STAGE 2: Adaptive DVX Polygon Fitting Training")
938:             self.train_stage_adaptive(
939:                 stage=2,
940:                 max_epochs=self.config.max_stage2_epochs,
941:                 min_epochs=self.config.min_stage2_epochs,
942:             )
943:             self.current_stage = 3
944:             self.stage_epoch = 0
945:             print("\nStage 2 completed. Transitioning to Stage 3...")
946:
947:         # Stage 3: Adaptive end-to-end fine-tuning
948:         if self.current_stage <= 3:
949:             print("\n? STAGE 3: Adaptive End-to-End Fine-tuning with Full Loss Suite")
950:             self.train_stage_adaptive(
951:                 stage=3,
952:                 max_epochs=self.config.max_stage3_epochs,
953:                 min_epochs=self.config.min_stage3_epochs,
954:             )
955:             print("\nStage 3 completed!")
956:
957:         print("\n" + "=" * 80)
958:         print("? ALL ADAPTIVE TRAINING STAGES COMPLETED!")
959:         print("=" * 80)
960:
961:         # Generate training report
962:         self._generate_training_report()
963:
964:         # Save final model
965:         self._save_checkpoint("final_adaptive_model.pth")
966:
967:         # Clean up rolling checkpoint
968:         if Path(self.ROLLING_CHECKPOINT).exists():
969:             Path(self.ROLLING_CHECKPOINT).unlink()
970:             print(f"Cleaned up rolling checkpoint: {self.ROLLING_CHECKPOINT}")
971:

```

```

972:         return self.history
973:
974:     def _generate_training_report(self):
975:         """Generate comprehensive training report with curriculum insights"""
976:         print("\n" + "=" * 60)
977:         print("ADAPTIVE TRAINING REPORT")
978:         print("=" * 60)
979:
980:         # Stage transition summary
981:         if self.history["stage_transitions"]:
982:             print("\n? Stage Transitions:")
983:             for transition in self.history["stage_transitions"]:
984:                 print(
985:                     f" ? Stage {transition['from_stage']} ? {transition['from_stage']+1}: "
986:                     f"Epoch {transition['epoch']} (Global: {transition['global_epoch']})"
987:                 )
988:                 print(f"     Reason: {transition['reason']}")
989:
990:         # Dynamic weight evolution
991:         if self.history["dynamic_weights"]:
992:             print(
993:                 f"\n?? Dynamic Weight Updates: {len(self.history['dynamic_weights'])} updates"
994:             )
995:             final_weights = self.history["dynamic_weights"][-1]["weights"]
996:             print(" Final loss weights:")
997:             for name, weight in final_weights.items():
998:                 if weight > 0.001:
999:                     print(f"     {name}: {weight:.3f}")
1000:
1001:         # Curriculum events
1002:         if self.history["curriculum_events"]:
1003:             print(
1004:                 f"\n? Curriculum Events: {len(self.history['curriculum_events'])} events"
1005:             )
1006:             for event in self.history["curriculum_events"][-5:]: # Show last 5 events
1007:                 print(
1008:                     f" ? {event['type']} at global epoch {event['epoch']}: {event['details']}"
1009:                 )
1010:
1011:         # Performance summary
1012:         print("\n? Final Performance:")
1013:         for stage_name, data in self.history.items():
1014:             if isinstance(data, dict) and "val_loss" in data and data["val_loss"]:
1015:                 final_loss = data["val_loss"][-1]
1016:                 best_loss = min(data["val_loss"])
1017:                 print(
1018:                     f" ? {stage_name.upper()}: Final={final_loss:.4f}, Best={best_loss:.4f}"
1019:                 )
1020:
1021:         print("\n? Training completed with novel adaptive curriculum strategies!")
1022:         print("=" * 60)
1023:
1024:     def _save_checkpoint(self, filename):
1025:         """Save final training checkpoint"""
1026:         checkpoint = {
1027:             "model_state_dict": self.model.state_dict(),
1028:             "optimizer_2d_state_dict": self.optimizer_2d.state_dict(),
1029:             "optimizer_dvx_state_dict": self.optimizer_dvx.state_dict(),
1030:             "optimizer_full_state_dict": self.optimizer_full.state_dict(),
1031:             "scheduler_2d_state_dict": self.scheduler_2d.state_dict(),
1032:             "scheduler_dvx_state_dict": self.scheduler_dvx.state_dict(),
1033:             "scheduler_full_state_dict": self.scheduler_full.state_dict(),
1034:             "scaler_state_dict": self.scaler.state_dict(),
1035:             "loss_fn_state": {
1036:                 "weights": self.loss_fn.weights,
1037:                 "initial_weights": self.loss_fn.initial_weights,
1038:             },
1039:             "history": self.history,
1040:             "config": self.config,
1041:             "final_stage": self.current_stage,
1042:             "total_epochs": self.global_epoch,
1043:             "training_complete": True,
1044:             "curriculum_summary": {

```

```

1045:         "stage_transitions": len(self.history["stage_transitions"]),
1046:         "weight_updates": len(self.history["dynamic_weights"]),
1047:         "curriculum_events": len(self.history["curriculum_events"]),
1048:     },
1049: }
1050: torch.save(checkpoint, filename)
1051: print(f"Final model saved: {filename}")
1052:
1053:
1054: # Legacy compatibility class
1055: class MultiStageTrainer(AdaptiveMultiStageTrainer):
1056:     """
1057:     Legacy wrapper for backward compatibility
1058:     Redirects to the new adaptive trainer
1059:     """
1060:
1061:     def __init__(self, *args, **kwargs):
1062:         super().__init__(*args, **kwargs)
1063:         print("Note: Using enhanced AdaptiveMultiStageTrainer with dynamic curriculum")
1064:
1065:     def train_stage1(self, epochs=None):
1066:         """Legacy method - redirects to adaptive training"""
1067:         max_epochs = epochs or self.config.max_stage1_epochs
1068:         min_epochs = self.config.min_stage1_epochs
1069:         return self.train_stage_adaptive(1, max_epochs, min_epochs)
1070:
1071:     def train_stage2(self, epochs=None):
1072:         """Legacy method - redirects to adaptive training"""
1073:         max_epochs = epochs or self.config.max_stage2_epochs
1074:         min_epochs = self.config.min_stage2_epochs
1075:         return self.train_stage_adaptive(2, max_epochs, min_epochs)
1076:
1077:     def train_stage3(self, epochs=None):
1078:         """Legacy method - redirects to adaptive training"""
1079:         max_epochs = epochs or self.config.max_stage3_epochs
1080:         min_epochs = self.config.min_stage3_epochs
1081:         return self.train_stage_adaptive(3, max_epochs, min_epochs)

```

File: utils\visualization.py

```

=====
1: """
2: Visualization and utility functions
3: """
4:
5: import matplotlib.pyplot as plt
6: import numpy as np
7: import cv2
8: import torch
9: from pathlib import Path
10: from evaluation.metrics import compute_iou
11:
12:
13: def plot_training_history(history, save_path="training_history.png"):
14:     """Plot training curves for all stages"""
15:     fig, axes = plt.subplots(1, 3, figsize=(15, 5))
16:
17:     for idx, (stage, data) in enumerate(history.items()):
18:         if isinstance(data, dict) and "train_loss" in data and data["train_loss"]: # Only plot if stage was trained
19:             axes[idx].plot(data["train_loss"], label="Train", linewidth=2)
20:             axes[idx].plot(data["val_loss"], label="Validation", linewidth=2)
21:             axes[idx].set_title(f"{stage.upper()} Training")
22:             axes[idx].set_xlabel("Epoch")
23:             axes[idx].set_ylabel("Loss")
24:             axes[idx].legend()
25:             axes[idx].grid(True, alpha=0.3)
26:
27:     plt.tight_layout()
28:     plt.savefig(save_path, dpi=300, bbox_inches="tight")
29:     plt.show()
30:
31:

```

```

32: def plot_curriculum_analysis(history, save_path="curriculum_analysis.png"):
33:     """Plot curriculum learning analysis including stage transitions and adaptive behavior"""
34:     fig, axes = plt.subplots(2, 2, figsize=(15, 10))
35:
36:     # Plot 1: Stage transition timeline
37:     if "stage_transitions" in history and history["stage_transitions"]:
38:         transitions = history["stage_transitions"]
39:
40:         # Extract transition epochs and reasons
41:         transition_epochs = [t["epoch"] for t in transitions]
42:         transition_stages = [t["from_stage"] + " ? " + t["to_stage"] for t in transitions]
43:         transition_reasons = [t.get("reason", "threshold") for t in transitions]
44:
45:         # Create timeline
46:         y_positions = range(len(transition_epochs))
47:         colors = ['red' if 'patience' in reason else 'green' for reason in transition_reasons]
48:
49:         axes[0, 0].barh(y_positions, transition_epochs, color=colors, alpha=0.7)
50:         axes[0, 0].set_yticks(y_positions)
51:         axes[0, 0].set_yticklabels(transition_stages)
52:         axes[0, 0].set_xlabel("Epoch")
53:         axes[0, 0].set_title("Stage Transition Timeline")
54:         axes[0, 0].grid(True, alpha=0.3)
55:
56:         # Add legend
57:         axes[0, 0].legend(['Patience-based', 'Threshold-based'], loc='lower right')
58:     else:
59:         axes[0, 0].text(0.5, 0.5, "No stage transitions recorded",
60:                        ha='center', va='center', transform=axes[0, 0].transAxes)
61:         axes[0, 0].set_title("Stage Transition Timeline")
62:
63:     # Plot 2: Loss component evolution
64:     if "dynamic_weights" in history and history["dynamic_weights"]:
65:         weight_data = history["dynamic_weights"]
66:         epochs = [entry["epoch"] for entry in weight_data]
67:
68:         # Plot each loss component weight
69:         weight_names = list(weight_data[0]["weights"].keys()) if weight_data else []
70:         for weight_name in weight_names[:5]: # Limit to top 5 for readability
71:             weights = [entry["weights"].get(weight_name, 0) for entry in weight_data]
72:             if any(w > 0.001 for w in weights): # Only plot significant weights
73:                 axes[0, 1].plot(epochs, weights, label=weight_name, linewidth=2, marker='o', markersize=3)
74:
75:         axes[0, 1].set_xlabel("Global Epoch")
76:         axes[0, 1].set_ylabel("Loss Weight")
77:         axes[0, 1].set_title("Dynamic Loss Weight Evolution")
78:         axes[0, 1].legend()
79:         axes[0, 1].grid(True, alpha=0.3)
80:     else:
81:         axes[0, 1].text(0.5, 0.5, "No dynamic weights recorded",
82:                        ha='center', va='center', transform=axes[0, 1].transAxes)
83:         axes[0, 1].set_title("Dynamic Loss Weight Evolution")
84:
85:     # Plot 3: Curriculum progress indicators
86:     if "curriculum_events" in history and history["curriculum_events"]:
87:         events = history["curriculum_events"]
88:         event_types = {}
89:
90:         for event in events:
91:             event_type = event.get("type", "unknown")
92:             if event_type not in event_types:
93:                 event_types[event_type] = []
94:             event_types[event_type].append(event["epoch"])
95:
96:         # Plot event timeline
97:         y_offset = 0
98:         for event_type, epochs in event_types.items():
99:             axes[1, 0].scatter(epochs, [y_offset] * len(epochs),
100:                              label=event_type, s=50, alpha=0.7)
101:             y_offset += 1
102:
103:         axes[1, 0].set_xlabel("Epoch")
104:         axes[1, 0].set_ylabel("Event Type")

```



```

105:         axes[1, 0].set_title("Curriculum Learning Events")
106:         axes[1, 0].legend()
107:         axes[1, 0].grid(True, alpha=0.3)
108:     else:
109:         axes[1, 0].text(0.5, 0.5, "No curriculum events recorded",
110:                         ha='center', va='center', transform=axes[1, 0].transAxes)
111:         axes[1, 0].set_title("Curriculum Learning Events")
112:
113:     # Plot 4: Stage performance comparison
114:     stage_names = ["stage1", "stage2", "stage3"]
115:     stage_performance = {}
116:
117:     for stage_name in stage_names:
118:         if stage_name in history and isinstance(history[stage_name], dict):
119:             stage_data = history[stage_name]
120:             if "val_loss" in stage_data and stage_data["val_loss"]:
121:                 stage_performance[stage_name] = {
122:                     "final_loss": stage_data["val_loss"][-1],
123:                     "best_loss": min(stage_data["val_loss"]),
124:                     "epochs": len(stage_data["val_loss"])
125:                 }
126:
127:     if stage_performance:
128:         stages = list(stage_performance.keys())
129:         final_losses = [stage_performance[s]["final_loss"] for s in stages]
130:         best_losses = [stage_performance[s]["best_loss"] for s in stages]
131:
132:         x = np.arange(len(stages))
133:         width = 0.35
134:
135:         axes[1, 1].bar(x - width/2, final_losses, width, label='Final Loss', alpha=0.8)
136:         axes[1, 1].bar(x + width/2, best_losses, width, label='Best Loss', alpha=0.8)
137:
138:         axes[1, 1].set_xlabel("Training Stage")
139:         axes[1, 1].set_ylabel("Validation Loss")
140:         axes[1, 1].set_title("Stage Performance Comparison")
141:         axes[1, 1].set_xticks(x)
142:         axes[1, 1].set_xticklabels([s.upper() for s in stages])
143:         axes[1, 1].legend()
144:         axes[1, 1].grid(True, alpha=0.3)
145:
146:         # Add epoch count annotations
147:         for i, stage in enumerate(stages):
148:             epochs = stage_performance[stage]["epochs"]
149:             axes[1, 1].text(i, max(final_losses) * 0.9, f'{epochs} epochs',
150:                             ha='center', va='bottom', fontsize=9)
151:     else:
152:         axes[1, 1].text(0.5, 0.5, "No stage performance data",
153:                         ha='center', va='center', transform=axes[1, 1].transAxes)
154:         axes[1, 1].set_title("Stage Performance Comparison")
155:
156:     plt.tight_layout()
157:     plt.savefig(save_path, dpi=300, bbox_inches="tight")
158:     plt.close()
159:
160:     print(f"Curriculum analysis saved to {save_path}")
161:
162:
163: def visualize_predictions(image, predictions, targets=None, save_path=None):
164:     """Visualize model predictions"""
165:     fig, axes = plt.subplots(2, 3, figsize=(15, 10))
166:
167:     # Original image
168:     if len(image.shape) == 4:
169:         img_np = image[0].permute(1, 2, 0).cpu().numpy()
170:     else:
171:         img_np = image.permute(1, 2, 0).cpu().numpy()
172:
173:     axes[0, 0].imshow(img_np)
174:     axes[0, 0].set_title("Input Image")
175:     axes[0, 0].axis('off')
176:
177:     # Predicted segmentation

```

```

178:     if "segmentation" in predictions:
179:         seg_pred = torch.argmax(predictions["segmentation"], dim=1)[0].cpu().numpy()
180:         axes[0, 1].imshow(seg_pred, cmap='tab10')
181:         axes[0, 1].set_title("Predicted Segmentation")
182:         axes[0, 1].axis('off')
183:
184:     # Ground truth segmentation (if available)
185:     if targets and "mask" in targets:
186:         gt_mask = targets["mask"][0].cpu().numpy()
187:         axes[0, 2].imshow(gt_mask, cmap='tab10')
188:         axes[0, 2].set_title("Ground Truth Segmentation")
189:         axes[0, 2].axis('off')
190:
191:     # SDF prediction
192:     if "sdf" in predictions:
193:         sdf_pred = predictions["sdf"][0, 0].cpu().numpy()
194:         im = axes[1, 0].imshow(sdf_pred, cmap='RdBu', vmin=-1, vmax=1)
195:         axes[1, 0].set_title("Predicted SDF")
196:         axes[1, 0].axis('off')
197:         plt.colorbar(im, ax=axes[1, 0])
198:
199:     # Polygon visualization
200:     if "polygons" in predictions:
201:         poly_vis = visualize_polygons(
202:             predictions["polygons"][0],
203:             predictions["polygon_validity"][0],
204:             image_size=(256, 256)
205:         )
206:         axes[1, 1].imshow(poly_vis)
207:         axes[1, 1].set_title("Predicted Polygons")
208:         axes[1, 1].axis('off')
209:
210:     # 3D voxel slice
211:     if "voxels_pred" in predictions:
212:         voxels = torch.sigmoid(predictions["voxels_pred"][0]).cpu().numpy()
213:         # Show middle slice
214:         mid_slice = voxels[voxels.shape[0]//2]
215:         axes[1, 2].imshow(mid_slice, cmap='viridis')
216:         axes[1, 2].set_title("3D Voxels (Mid Slice)")
217:         axes[1, 2].axis('off')
218:
219:     plt.tight_layout()
220:
221:     if save_path:
222:         plt.savefig(save_path, dpi=300, bbox_inches="tight")
223:
224:     plt.show()
225:
226:
227: def visualize_polygons(polygons, validity, image_size=(256, 256), threshold=0.5):
228:     """Visualize predicted polygons"""
229:     vis_img = np.zeros((*image_size, 3), dtype=np.uint8)
230:
231:     for poly_idx, (polygon, valid_score) in enumerate(zip(polygons, validity)):
232:         if valid_score > threshold:
233:             # Convert to image coordinates
234:             points = polygon.cpu().numpy() * np.array(image_size)
235:
236:             # Remove zero-padded points
237:             valid_points = points[points.sum(axis=1) > 0]
238:
239:             if len(valid_points) >= 3:
240:                 points_int = valid_points.astype(np.int32)
241:
242:                 # Different colors for different polygons
243:                 color = plt.cm.tab10(poly_idx % 10)[:3]
244:                 color = tuple(int(c * 255) for c in color)
245:
246:                 cv2.polylines(vis_img, [points_int], True, color, 2)
247:
248:                 # Add polygon index
249:                 center = points_int.mean(axis=0).astype(int)
250:                 cv2.putText(vis_img, str(poly_idx), tuple(center),

```

```

251:                 cv2.FONT_HERSHEY_SIMPLEX, 0.5, (255, 255, 255), 1)
252:
253:     return vis_img
254:
255:
256: def save_model_outputs(predictions, output_dir, sample_id):
257:     """Save all model outputs for detailed analysis"""
258:     output_dir = Path(output_dir)
259:     output_dir.mkdir(exist_ok=True)
260:
261:     sample_dir = output_dir / sample_id
262:     sample_dir.mkdir(exist_ok=True)
263:
264:     # Save segmentation
265:     if "segmentation" in predictions:
266:         seg_pred = torch.argmax(predictions["segmentation"], dim=1)[0].cpu().numpy()
267:         cv2.imwrite(str(sample_dir / "segmentation.png"), seg_pred * 50)
268:
269:     # Save SDF
270:     if "sdf" in predictions:
271:         sdf_pred = predictions["sdf"][0, 0].cpu().numpy()
272:         sdf_normalized = ((sdf_pred + 1) * 127.5).astype(np.uint8)
273:         cv2.imwrite(str(sample_dir / "sdf.png"), sdf_normalized)
274:
275:     # Save attributes
276:     if "attributes" in predictions:
277:         attrs = predictions["attributes"][0].cpu().numpy()
278:         np.save(sample_dir / "attributes.npy", attrs)
279:
280:     # Save polygons
281:     if "polygons" in predictions:
282:         polygons = predictions["polygons"][0].cpu().numpy()
283:         validity = predictions["polygon_validity"][0].cpu().numpy()
284:
285:         np.save(sample_dir / "polygons.npy", polygons)
286:         np.save(sample_dir / "polygon_validity.npy", validity)
287:
288:     # Save voxels
289:     if "voxels_pred" in predictions:
290:         voxels = torch.sigmoid(predictions["voxels_pred"][0]).cpu().numpy()
291:         np.save(sample_dir / "voxels.npy", voxels)
292:
293:
294: def create_comparison_grid(input_images, predictions, targets=None, num_samples=4):
295:     """Create a comparison grid showing inputs, predictions, and targets"""
296:     fig, axes = plt.subplots(num_samples, 4, figsize=(16, 4 * num_samples))
297:
298:     for i in range(min(num_samples, len(input_images))):
299:         # Input image
300:         img = input_images[i].permute(1, 2, 0).cpu().numpy()
301:         axes[i, 0].imshow(img)
302:         axes[i, 0].set_title(f"Sample {i+1}: Input")
303:         axes[i, 0].axis('off')
304:
305:         # Predicted segmentation
306:         seg_pred = torch.argmax(predictions["segmentation"][i], dim=0).cpu().numpy()
307:         axes[i, 1].imshow(seg_pred, cmap='tab10')
308:         axes[i, 1].set_title("Predicted Seg")
309:         axes[i, 1].axis('off')
310:
311:         # Ground truth segmentation (if available)
312:         if targets and "mask" in targets:
313:             gt_mask = targets["mask"][i].cpu().numpy()
314:             axes[i, 2].imshow(gt_mask, cmap='tab10')
315:             axes[i, 2].set_title("GT Segmentation")
316:         else:
317:             axes[i, 2].text(0.5, 0.5, "No GT", ha='center', va='center',
318:                             transform=axes[i, 2].transAxes)
319:             axes[i, 2].set_title("GT Segmentation")
320:             axes[i, 2].axis('off')
321:
322:         # Polygon overlay
323:         poly_vis = visualize_polygons(

```

```

324:         predictions["polygons"][i],
325:         predictions["polygon_validity"][i]
326:     )
327:     axes[i, 3].imshow(poly_vis)
328:     axes[i, 3].set_title("Predicted Polygons")
329:     axes[i, 3].axis('off')
330:
331: plt.tight_layout()
332: return fig
333:
334:
335: def analyze_failure_cases(predictions, targets, threshold_iou=0.5):
336:     """Analyze failure cases for debugging"""
337:     failure_indices = []
338:
339:     for i, (pred_seg, gt_mask) in enumerate(zip(predictions["segmentation"], targets["mask"])):
340:         seg_pred = torch.argmax(pred_seg, dim=0)
341:         iou = compute_iou(seg_pred, gt_mask)
342:
343:         if iou < threshold_iou:
344:             failure_indices.append({
345:                 "index": i,
346:                 "iou": iou,
347:                 "pred_classes": torch.unique(seg_pred).tolist(),
348:                 "gt_classes": torch.unique(gt_mask).tolist()
349:             })
350:
351:     return failure_indices
352:
353:
354: class ProgressiveVisualization:
355:     """Track and visualize training progress"""
356:
357:     def __init__(self, save_dir="./training_progress"):
358:         self.save_dir = Path(save_dir)
359:         self.save_dir.mkdir(exist_ok=True)
360:
361:     def log_epoch_results(self, epoch, stage, predictions, targets, sample_image):
362:         """Log results for a specific epoch"""
363:         epoch_dir = self.save_dir / f"{stage}_epoch_{epoch}"
364:         epoch_dir.mkdir(exist_ok=True)
365:
366:         # Save prediction visualization
367:         fig = plt.figure(figsize=(12, 8))
368:         visualize_predictions(sample_image, predictions, targets)
369:         plt.savefig(epoch_dir / "predictions.png", dpi=150, bbox_inches="tight")
370:         plt.close()
371:
372:         # Save individual outputs
373:         save_model_outputs(predictions, epoch_dir, "sample")
374:
375:     def create_training_animation(self, stage, metric_name="total_loss"):
376:         """Create animated GIF showing training progress"""
377:         # This would create an animation of training progress
378:         # Implementation depends on having saved epoch results
379:         pass
380:
381:
382: def compute_architectural_metrics(predictions, image_size=(256, 256)):
383:     """Compute architecture-specific metrics"""
384:     metrics = {}
385:
386:     if "segmentation" in predictions:
387:         seg_pred = torch.argmax(predictions["segmentation"], dim=1)[0]
388:
389:         # Room count
390:         room_mask = (seg_pred == 0).cpu().numpy().astype(np.uint8)
391:         contours, _ = cv2.findContours(room_mask, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
392:         room_count = len([c for c in contours if cv2.contourArea(c) > 100])
393:         metrics["room_count"] = room_count
394:
395:         # Wall connectivity
396:         wall_mask = (seg_pred == 1).cpu().numpy().astype(np.uint8)

```

```

397:         wall_components = cv2.connectedComponents(wall_mask)[0] - 1 # Subtract background
398:         metrics["wall_components"] = max(0, wall_components)
399:
400:         # Door and window counts
401:         door_pixels = (seg_pred == 2).sum().item()
402:         window_pixels = (seg_pred == 3).sum().item()
403:         metrics["door_pixels"] = door_pixels
404:         metrics["window_pixels"] = window_pixels
405:
406:         if "polygons" in predictions:
407:             validity = predictions["polygon_validity"][0]
408:             valid_polygons = (validity > 0.5).sum().item()
409:             metrics["valid_polygon_count"] = valid_polygons
410:
411:             # Average polygon area
412:             polygons = predictions["polygons"][0]
413:             areas = []
414:             for poly_idx, (polygon, valid) in enumerate(zip(polygons, validity)):
415:                 if valid > 0.5:
416:                     # Compute polygon area using shoelace formula
417:                     points = polygon.cpu().numpy() * np.array(image_size)
418:                     valid_points = points[points.sum(axis=1) > 0]
419:                     if len(valid_points) >= 3:
420:                         area = compute_polygon_area(valid_points)
421:                         areas.append(area)
422:
423:             metrics["avg_polygon_area"] = np.mean(areas) if areas else 0.0
424:
425:         return metrics
426:
427:
428: def compute_polygon_area(points):
429:     """Compute polygon area using shoelace formula"""
430:     if len(points) < 3:
431:         return 0.0
432:
433:     x = points[:, 0]
434:     y = points[:, 1]
435:
436:     # Shoelace formula
437:     area = 0.5 * abs(sum(x[i] * y[i+1] - x[i+1] * y[i] for i in range(-1, len(x)-1)))
438:     return area
439:
440:
441: def create_model_summary_report(model, sample_input, save_path="model_summary.txt"):
442:     """Create detailed model summary report"""
443:     with open(save_path, "w") as f:
444:         f.write("Neural-Geometric 3D Model Generator - Model Summary\n")
445:         f.write("=" * 60 + "\n\n")
446:
447:         # Model architecture
448:         f.write("MODEL ARCHITECTURE:\n")
449:         f.write("-" * 20 + "\n")
450:
451:         total_params = sum(p.numel() for p in model.parameters())
452:         trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
453:
454:         f.write(f"Total parameters: {total_params:,}\n")
455:         f.write(f"Trainable parameters: {trainable_params:,}\n")
456:         f.write(f"Model size: {total_params * 4 / 1024 / 1024:.2f} MB\n\n")
457:
458:         # Component breakdown
459:         f.write("COMPONENT PARAMETERS:\n")
460:         f.write("-" * 25 + "\n")
461:
462:         encoder_params = sum(p.numel() for p in model.encoder.parameters())
463:         seg_params = sum(p.numel() for p in model.seg_head.parameters())
464:         attr_params = sum(p.numel() for p in model.attr_head.parameters())
465:         sdf_params = sum(p.numel() for p in model.sdf_head.parameters())
466:         dvx_params = sum(p.numel() for p in model.dvx.parameters())
467:         ext_params = sum(p.numel() for p in model.extrusion.parameters())
468:
469:         f.write(f"Encoder: {encoder_params:,} ({encoder_params/total_params*100:.1f}%)\n")

```

```

470:         f.write(f"Segmentation Head: {seg_params:,} ({seg_params/total_params*100:.1f}%)\\n")
471:         f.write(f"Attribute Head: {attr_params:,} ({attr_params/total_params*100:.1f}%)\\n")
472:         f.write(f"SDF Head: {sdf_params:,} ({sdf_params/total_params*100:.1f}%)\\n")
473:         f.write(f"DVX Module: {dvx_params:,} ({dvx_params/total_params*100:.1f}%)\\n")
474:         f.write(f"Extrusion Module: {ext_params:,} ({ext_params/total_params*100:.1f}%)\\n\\n")
475:
476:         # Forward pass analysis
477:         f.write("FORWARD PASS ANALYSIS:\\n")
478:         f.write("-" * 25 + "\\n")
479:
480:         model.eval()
481:         with torch.no_grad():
482:             predictions = model(sample_input)
483:
484:             for key, value in predictions.items():
485:                 if torch.is_tensor(value):
486:                     f.write(f"{key}: {list(value.shape)} - {value.dtype}\\n")
487:                 else:
488:                     f.write(f"{key}: {type(value)}\\n")
489:
490:         print(f"Model summary saved to {save_path}")
491:
492:
493: def debug_gradient_flow(model, loss):
494:     """Debug gradient flow through the model"""
495:     print("Gradient Flow Analysis:")
496:     print("-" * 30)
497:
498:     total_norm = 0
499:     component_norms = {}
500:
501:     for name, param in model.named_parameters():
502:         if param.grad is not None:
503:             param_norm = param.grad.norm().item()
504:             total_norm += param_norm ** 2
505:
506:             # Group by component
507:             component = name.split('.')[0]
508:             if component not in component_norms:
509:                 component_norms[component] = 0
510:             component_norms[component] += param_norm ** 2
511:
512:     total_norm = total_norm ** 0.5
513:
514:     print(f"Total gradient norm: {total_norm:.4f}")
515:     print("Component gradient norms:")
516:
517:     for component, norm in component_norms.items():
518:         norm = norm ** 0.5
519:         print(f"  {component}: {norm:.4f} ({norm/total_norm*100:.1f}%)")
520:
521:
522: def create_3d_visualization(voxels, output_path="3d_preview.png"):
523:     """Create 3D visualization of voxel prediction"""
524:     try:
525:         import matplotlib.pyplot as plt
526:         from mpl_toolkits.mplot3d import Axes3D
527:
528:         # Convert to binary
529:         if isinstance(voxels, torch.Tensor):
530:             voxels = voxels.cpu().numpy()
531:
532:         binary_voxels = voxels > 0.5
533:
534:         # Get occupied voxel coordinates
535:         occupied = np.where(binary_voxels)
536:
537:         if len(occupied[0]) == 0:
538:             print("No occupied voxels to visualize")
539:             return
540:
541:         # Create 3D plot
542:         fig = plt.figure(figsize=(10, 8))

```

```

543:         ax = fig.add_subplot(111, projection='3d')
544:
545:         # Plot occupied voxels
546:         ax.scatter(occupied[0], occupied[1], occupied[2],
547:                   c=occupied[2], cmap='viridis', s=1, alpha=0.6)
548:
549:         ax.set_xlabel('X')
550:         ax.set_ylabel('Y')
551:         ax.set_zlabel('Z')
552:         ax.set_title('3D Voxel Occupancy')
553:
554:         plt.savefig(output_path, dpi=150, bbox_inches="tight")
555:         plt.close()
556:
557:         print(f"3D visualization saved to {output_path}")
558:
559:     except ImportError:
560:         print("3D visualization requires matplotlib with 3D support")

```

■ File: vortex.py

```

=====
1: import cv2
2: import numpy as np
3: from pathlib import Path
4: import json
5:
6: # Base dataset path
7: data_root = Path("./data/floorplans")
8:
9: def is_valid_mask(mask_file):
10:     m = cv2.imread(str(mask_file), 0)
11:     return m is not None and np.sum(m) > 0
12:
13: def is_valid_voxel(voxel_file):
14:     try:
15:         data = np.load(str(voxel_file))
16:         key = "voxels" # Use the correct key in your npz files
17:         if key not in data.files:
18:             print(f"?? Key '{key}' not found in {voxel_file}, available keys: {data.files}")
19:             return False
20:         v = data[key]
21:         return not np.isnan(v).any()
22:     except Exception as e:
23:         print(f"?? Error reading {voxel_file}: {e}")
24:         return False
25:
26: def is_valid_polygon(polygon_file):
27:     try:
28:         with open(polygon_file, "r") as f:
29:             json.load(f)
30:         return True
31:     except Exception as e:
32:         print(f"?? Invalid polygon {polygon_file}: {e}")
33:         return False
34:
35: def check_split(split="train"):
36:     split_dir = data_root / split
37:     mask_files = list(split_dir.rglob("mask.png"))
38:     voxel_files = list(split_dir.rglob("voxel_GT.npz"))
39:     polygon_files = list(split_dir.rglob("polygon.json"))
40:
41:     print(f"\n? Checking split: {split}")
42:
43:     # Check mask files
44:     total_masks = len(mask_files)
45:     bad_masks = 0
46:     for f in mask_files:
47:         if not is_valid_mask(f):
48:             bad_masks += 1
49:             print(f"?? Invalid mask: {f}")
50:

```

```

51:     # Check voxel files
52:     total_voxels = len(voxel_files)
53:     bad_voxels = 0
54:     for f in voxel_files:
55:         if not is_valid_voxel(f):
56:             bad_voxels += 1
57:             print(f"?? Invalid voxel: {f}")
58:
59:     # Check polygon files
60:     total_polygons = len(polygon_files)
61:     bad_polygons = 0
62:     for f in polygon_files:
63:         if not is_valid_polygon(f):
64:             bad_polygons += 1
65:
66:     # Summary
67:     print(f"\n? Summary for split: {split}")
68:     print(f"Total mask files checked: {total_masks}, Invalid: {bad_masks}")
69:     print(f"Total voxel files checked: {total_voxels}, Invalid: {bad_voxels}")
70:     print(f"Total polygon files checked: {total_polygons}, Invalid: {bad_polygons}")
71:
72: def main():
73:     for split in ["train", "val", "test"]:
74:         check_split(split)
75:
76: if __name__ == "__main__":
77:     main()

```
